

Public signup for this instance is disabled. Our Jira Guidelines page explains how to get an account.



Hadoop Common / HADOOP-2768

# DFSIO write performance benchmark shows a regression

#### Details

Type: ☐ Bug
Priority: ☐ Blocker

Affects Version/s: 0.16.0

Component/s: None Labels: None

Status: CLOSED
Resolution: Fixed

Fix Version/s: 0.16.0

## Description

We are seeing a performance regression in DFSIO write benchmark, possibly due to Hadoop-1707.

Here is the data:

- Aggregate cluster throughput went from 0.30G to 0.14G with 39 concurrent clients, on 20 nodes
- Aggregate cluster throughput went from 1.15G to 0.85G with 39 concurrent clients, on 100 nodes

#### Attachments

perf.patch	1 kB	01/Feb/08 22:04
perf.patch	2 kB	01/Feb/08 21:58
perf1.patch	3 kB	03/Feb/08 08:57
perf2.patch	3 kB	04/Feb/08 05:09

#### Activity

Dhruba Borthakur added a comment - 01/Feb/08 21:58

Datanodes write to local file using a bufferedotputstream

Mukund Madhugiri added a comment - 01/Feb/08 23:12

With the patch, I see a performance improvement: Aggregate cluster throughput with the patch is at 1.145G on 100 nodes

Raghu Angadi added a comment - 01/Feb/08 23:26

+1.

Yes, this will increase buffer copies. But we intend get rid of copies without changing effective write size for sockets for 0.17 (HADOOP 2154). I went through the same experience with HADOOP 1134, initially didn't have BufferedOutputStream and added it after the benchmarks.

- ▼ Hadoop QA added a comment 02/Feb/08 02:17
  - -1 overall. Here are the results of testing the latest attachment http://issues.apache.org/jira/secure/attachment/12374587/perf.patch against trunk revision 616796.

@author +1. The patch does not contain any @author tags.

javadoc +1. The javadoc tool did not generate any warning messages.

javac +1. The applied patch does not generate any new javac compiler warnings.

release audit +1. The applied patch does not generate any new release audit warnings.

findbugs +1. The patch does not introduce any new Findbugs warnings.

1

core tests -1. The patch failed core unit tests.

contrib tests +1. The patch passed contrib unit tests.

Test results: http://hudson.zones.apache.org/hudson/job/Hadoop-Patch/1725/testReport/

Findbugs warnings: http://hudson.zones.apache.org/hudson/job/Hadoop-Patch/1725/artifact/trunk/build/test/findbugs/newPatchFindbugsWarnings.html

Checkstyle results: http://hudson.zones.apache.org/hudson/job/Hadoop-Patch/1725/artifact/trunk/build/test/checkstyle-errors.html

Console output: http://hudson.zones.apache.org/hudson/job/Hadoop-Patch/1725/console

This message is automatically generated.

#### Hadoop QA added a comment - 02/Feb/08 05:43

-1 overall. Here are the results of testing the latest attachment http://issues.apache.org/jira/secure/attachment/12374587/perf.patch against trunk revision 616796.

@author +1. The patch does not contain any @author tags.

javadoc +1. The javadoc tool did not generate any warning messages.

javac +1. The applied patch does not generate any new javac compiler warnings.

release audit +1. The applied patch does not generate any new release audit warnings.

findbugs +1. The patch does not introduce any new Findbugs warnings.

core tests -1. The patch failed core unit tests.

contrib tests +1. The patch passed contrib unit tests.

Test results: http://hudson.zones.apache.org/hudson/job/Hadoop-Patch/1727/testReport/

Findbugs warnings: http://hudson.zones.apache.org/hudson/job/Hadoop-Patch/1727/artifact/trunk/build/test/findbugs/newPatchFindbugsWarnings.html

Checkstyle results: http://hudson.zones.apache.org/hudson/job/Hadoop-Patch/1727/artifact/trunk/build/test/checkstyle-errors.html

Console output: http://hudson.zones.apache.org/hudson/job/Hadoop-Patch/1727/console

This message is automatically generated.

## Nigel Daley added a comment - 02/Feb/08 20:19

In both core test runs above, TestBalancer failed with this exception:

[junit] 2008-02-02 00:51:29,097 INFO dfs.DataNode (DataNode.java:writeBlock(1169)) - writeBlock

blk\_-4006373334377701738 received exception java.io.IOException: setBlockPosition trying to set position to 3 which is not a multiple of bytesPerChecksum 10

[junit] 2008-02-02 00:51:29,098 ERROR dfs.DataNode (DataNode.java:run(961)) - 127.0.0.1:49578:DataXceiver:

java.io.IOException: setBlockPosition trying to set position to 3 which is not a multiple of bytesPerChecksum 10

 $[junit]\ at\ org. apache. hadoop. dfs. DataNode \$BlockReceiver. set BlockPosition (DataNode. java: 2301)$ 

[junit] at org.apache.hadoop.dfs.DataNode\$BlockReceiver.receivePacket(DataNode.java:2126)

[junit] at org.apache.hadoop.dfs.DataNode\$BlockReceiver.receiveBlock(DataNode.java:2230)

[junit] at org.apache.hadoop.dfs.DataNode\$DataXceiver.writeBlock(DataNode.java:1150)

[junit] at org.apache.hadoop.dfs.DataNode\$DataXceiver.run(DataNode.java:938)

[junit] at java.lang.Thread.run(Thread.java:595)

#### Dhruba Borthakur added a comment - 03/Feb/08 08:55

Failure in unit tests TestBalancer

## Dhruba Borthakur added a comment - 03/Feb/08 08:57

Since the datanode started using a bufferoutputstream, the verification that the offset of a packet in the block matches the current channel offset needed to be modified.

#### ▼ ○ Hadoop QA added a comment - 03/Feb/08 10:34

-1 overall. Here are the results of testing the latest attachment http://issues.apache.org/jira/secure/attachment/12374633/perf1.patch

against trunk revision 616796.

@author +1. The patch does not contain any @author tags.

javadoc +1. The javadoc tool did not generate any warning messages.

javac +1. The applied patch does not generate any new javac compiler warnings.

release audit +1. The applied patch does not generate any new release audit warnings.

findbugs -1. The patch appears to introduce 1 new Findbugs warnings.

core tests +1. The patch passed core unit tests.

contrib tests +1. The patch passed contrib unit tests.

Test results: http://hudson.zones.apache.org/hudson/job/Hadoop-Patch/1733/testReport/

Findbugs warnings: http://hudson.zones.apache.org/hudson/job/Hadoop-Patch/1733/artifact/trunk/build/test/findbugs/newPatchFindbugsWarnings.html

Checkstyle results: http://hudson.zones.apache.org/hudson/job/Hadoop-Patch/1733/artifact/trunk/build/test/checkstyle-errors.html

Console output: http://hudson.zones.apache.org/hudson/job/Hadoop-Patch/1733/console

This message is automatically generated.

Dhruba Borthakur added a comment - 04/Feb/08 05:09

Findbugs warnings

Dhruba Borthakur added a comment - 04/Feb/08 05:09

Fix findbugs warnings

- ✓ Hadoop QA added a comment 04/Feb/08 07:25
  - -1 overall. Here are the results of testing the latest attachment http://issues.apache.org/jira/secure/attachment/12374659/perf2.patch against trunk revision 616796.

@author +1. The patch does not contain any @author tags.

javadoc +1. The javadoc tool did not generate any warning messages.

javac +1. The applied patch does not generate any new javac compiler warnings.

release audit +1. The applied patch does not generate any new release audit warnings.

findbugs +1. The patch does not introduce any new Findbugs warnings.

core tests -1. The patch failed core unit tests.

contrib tests +1. The patch passed contrib unit tests.

Test results: http://hudson.zones.apache.org/hudson/job/Hadoop-Patch/1735/testReport/

Findbugs warnings: http://hudson.zones.apache.org/hudson/job/Hadoop-

Patch/1735/artifact/trunk/build/test/findbugs/newPatchFindbugsWarnings.html

 $\label{lem:checkstyle} \textbf{Checkstyle results: } http://hudson.zones.apache.org/hudson/job/Hadoop-Patch/1735/artifact/trunk/build/test/checkstyle-errors.html.$ 

Console output: http://hudson.zones.apache.org/hudson/job/Hadoop-Patch/1735/console

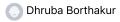
This message is automatically generated.

✓ Nigel Daley added a comment - 04/Feb/08 18:00

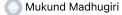
I just committed this. Thanks Dhruba!

## ✓ People

Assignee:



Reporter:



Votes:

• Vote for this issue

Watchers:

O Start watching this issue

## Dates

Created:

01/Feb/08 18:59

Updated:

08/Jul/09 16:42

Resolved:

04/Feb/08 18:00