



Hadoop Common / HADOOP-3382

Memory leak when files are not cleanly closed

Details

Type:	Bug	Status:	CLOSED
Priority:	Blocker	Resolution:	Fixed
Affects Version/s:	0.15.0	Fix Version/s:	0.17.0
Component/s:	None		
Labels:	None		
Hadoop Flags:	Reviewed		
Release Note:	Fixed a memory leak associated with 'abandoned' files (i.e. not cleanly closed). This held up significant amounts of memory depending on activity and how long NameNode has been running.		

Description

FSNamesystem.internalReleaseCreate() in invoked on files that are open for writing but not cleanly closed. e.g. when client invokes abandonFileInProgress() or when lease expires. It deletes the last block if it has a length of zero. The block is deleted from the file INode but not from blocksMap. Then leaves a reference to such file until NameNode is restarted. When this happens [HADOOP-3381](#) multiplies amount of memory leak.

Attachments

HADOOP-3382.patch	1 kB	14/May/08 18:16
HADOOP-3382.patch	0.7 kB	14/May/08 07:44
HADOOP-3382.patch	0.7 kB	14/May/08 07:42
memleak.txt	2 kB	13/May/08 21:59

Activity



[Raghu Angadi](#) created issue - 13/May/08 18:40

[Tsz-wo Sze](#) added a comment - 13/May/08 20:53

abandonFileInProgress() in ClientProtocol is deprecated. We should remove it as well.

[Raghu Angadi](#) added a comment - 13/May/08 21:59

memleak.txt attached.

Koji confirmed that this leak exists and it could lead to many other INode objects to leak, using a simple one node cluster and manually interrupting writing one of the files.

[Raghu Angadi](#) made changes - 13/May/08 21:59

Field	Original Value	New Value
Attachment		memleak.txt [12381989]

[Robert Chansler](#) added a comment - 13/May/08 23:17

0.17 won't be good enough for the big grids without this.

[Robert Chansler](#) made changes - 13/May/08 23:17

Fix Version/s	0.17.0 [12312913]
Description	{{FSNamesystem.internalReleaseCreate()}} in invoked on files that are open for writing but not


invoked on files that are open for writing but not cleanly closed. e.g. when client invokes `{{abandonFileInProgress()}}` or when lease expires. It deletes the last block if it has a length of zero. The block is deleted from the file `INode` but not from `{{blocksMap}}`. Then leaves a reference to such file until `NameNode` is restarted. When this happens [HADOOP-3384](#) multiplies amount of memory leak.

cleanly closed. e.g. when client invokes `{{abandonFileInProgress()}}` or when lease expires. It deletes the last block if it has a length of zero. The block is deleted from the file `INode` but not from `{{blocksMap}}`. Then leaves a reference to such file until `NameNode` is restarted. When this happens [HADOOP-3384](#) multiplies amount of memory leak.

Priority

Major [3]

Blocker [1]

▼  [Raghu Angadi](#) added a comment - 14/May/08 07:42 - edited


The patch that is tested manually is attached. As my comment in the patch describes, we do different things when a block is removed in different context. What is done a block is removed from `NameNode` should be in one place. This patch only fixes the observed leak.

I was thinking of writing unit test using `abandonFileInProgress()` but it is already deprecated. Since it is fairly simple patch and tested manually, it is probably ok not to have a unit test.

 [Raghu Angadi](#) made changes - 14/May/08 07:42


Attachment

[HADOOP-3382.patch](#) [12382015]

 [Raghu Angadi](#) made changes - 14/May/08 07:44

Attachment


[HADOOP-3382.patch](#) [12382016]

 [Raghu Angadi](#) made changes - 14/May/08 07:45

Status

Open [1]

Patch Available [10002]

▼  [Hadoop QA](#) added a comment - 14/May/08 10:53

-1 overall. Here are the results of testing the latest attachment <http://issues.apache.org/jira/secure/attachment/12382016/HADOOP-3382.patch> against trunk revision 656153.

+1 @author. The patch does not contain any @author tags.

-1 tests included. The patch doesn't appear to include any new or modified tests. Please justify why no tests are needed for this patch.

+1 javadoc. The javadoc tool did not generate any warning messages.

+1 javac. The applied patch does not increase the total number of javac compiler warnings.

+1 findbugs. The patch does not introduce any new Findbugs warnings.

+1 release audit. The applied patch does not increase the total number of release audit warnings.

+1 core tests. The patch passed core unit tests.

+1 contrib tests. The patch passed contrib unit tests.


Test results: <http://hudson.zones.apache.org/hudson/job/Hadoop-Patch/2463/testReport/>

Findbugs warnings: <http://hudson.zones.apache.org/hudson/job/Hadoop-Patch/2463/artifact/trunk/build/test/findbugs/newPatchFindbugsWarnings.html>

Checkstyle results: <http://hudson.zones.apache.org/hudson/job/Hadoop-Patch/2463/artifact/trunk/build/test/checkstyle-errors.html>

Console output: <http://hudson.zones.apache.org/hudson/job/Hadoop-Patch/2463/console>

This message is automatically generated.

▼  [Raghu Angadi](#) added a comment - 14/May/08 16:01

Thanks to Koji for helping the investigation and verification of the fix.

▼  [Tsz-wo Sze](#) added a comment - 14/May/08 17:18

I guess [HADOOP-3382](#).patch is for 0.16. Should we also remove all datanodes from `BlockInfo`?

- ▼

Tsz-wo Sze

added a comment - 14/May/08 17:51

created [HADOOP-3390](#) to remove ClientProtocol.abandonFileInProgress()
- Raghu Angadi

added a comment - 14/May/08 17:57 - edited

> Should we also remove all datanodes from BlockInfo?
Possibly. Please see my comment in the code and in jira above. There is no explicit and/or consistent policy regd what needs to be cleaned up. I didn't want to fix that in this jira.

Pretty much in all these cases datanodes would not completed the block. I will add cleaning up the datanodes part.
- Raghu Angadi

added a comment - 14/May/08 18:16

Updated patch adds block to invalidated blocks for any datanodes that have this block. Ideally processBlockReport should handle this. When it iterates, if a block does not have any file associated with it, it should be removed.
- Raghu Angadi

made changes - 14/May/08 18:16

Attachment

HADOOP-3382.patch [12382065]
- ▼

Tsz-wo Sze

added a comment - 14/May/08 18:21

+1 patch looks good
- Tsz-wo Sze

made changes - 14/May/08 18:21

Hadoop Flags

[Reviewed]
- Raghu Angadi

made changes - 14/May/08 18:29

Status

Patch Available [10002]

Open [1]
- Raghu Angadi

made changes - 14/May/08 18:29

Status

Open [1]

Patch Available [10002]
- ▼

Hadoop QA

added a comment - 14/May/08 21:45

-1 overall. Here are the results of testing the latest attachment
<http://issues.apache.org/jira/secure/attachment/12382065/HADOOP-3382.patch>
against trunk revision 656270.

+1 @author. The patch does not contain any @author tags.

-1 tests included. The patch doesn't appear to include any new or modified tests.
Please justify why no tests are needed for this patch.

+1 javadoc. The javadoc tool did not generate any warning messages.

+1 javac. The applied patch does not increase the total number of javac compiler warnings.

+1 findbugs. The patch does not introduce any new Findbugs warnings.

+1 release audit. The applied patch does not increase the total number of release audit warnings.

-1 core tests. The patch failed core unit tests.

+1 contrib tests. The patch passed contrib unit tests.

Test results: <http://hudson.zones.apache.org/hudson/job/Hadoop-Patch/2469/testReport/>
Findbugs warnings: <http://hudson.zones.apache.org/hudson/job/Hadoop-Patch/2469/artifact/trunk/build/test/findbugs/newPatchFindbugsWarnings.html>
Checkstyle results: <http://hudson.zones.apache.org/hudson/job/Hadoop-Patch/2469/artifact/trunk/build/test/checkstyle-errors.html>
Console output: <http://hudson.zones.apache.org/hudson/job/Hadoop-Patch/2469/console>


This message is automatically generated.
- ▼

Raghu Angadi

added a comment - 14/May/08 23:00

TestMiniMRBingUp failes. This is very simple test and does not even use DFS. The failure seems to another case of [HADOOP-3354](#) (thanks Lohit). I think this patch can be committed.

👤 Raghu Angadi made changes - 14/May/08 23:26		
Status	Patch Available [10002]	Open [1]
👤 Raghu Angadi made changes - 14/May/08 23:26		
Status	Open [1]	Patch Available [10002]
▼ 👤 Raghu Angadi added a comment - 14/May/08 23:48		
I just committed this.		
👤 Raghu Angadi made changes - 14/May/08 23:48		
Release Note	Fixed a memory leak associated with 'abandoned' files (i.e. not cleanly closed). This held up significant amounts of memory depending on activity and how long NameNode has been running.	
Affects Version/s	0.15.0 [12312565]	
Affects Version/s	0.16.0 [12312740]	
▼ 🗳️ Hadoop QA added a comment - 15/May/08 02:11		
<p>-1 overall. Here are the results of testing the latest attachment http://issues.apache.org/jira/secure/attachment/12382065/HADOOP-3382.patch against trunk revision 656480.</p> <p>+1 @author. The patch does not contain any @author tags.</p> <p>-1 tests included. The patch doesn't appear to include any new or modified tests. Please justify why no tests are needed for this patch.</p> <p>-1 patch. The patch command could not apply the patch.</p> <p>Console output: http://hudson.zones.apache.org/hudson/job/Hadoop-Patch/2474/console</p> <p>This message is automatically generated.</p>		
▼ 🗳️ Hudson added a comment - 15/May/08 12:43		
Integrated in Hadoop-trunk #492 (See http://hudson.zones.apache.org/hudson/job/Hadoop-trunk/492/)		
👤 Raghu Angadi made changes - 15/May/08 18:33		
Resolution	Fixed [1]	
Status	Patch Available [10002]	Resolved [5]
👤 Nigel Daley made changes - 21/May/08 20:06		
Status	Resolved [5]	Closed [6]
👤 Owen O'Malley made changes - 08/Jul/09 16:43		
Component/s	dfs [12310710]	
Transition		
👤 Raghu Angadi made transition - 14/May/08 23:26		
PATCH AVAILABLE ➡️ OPEN		
15h 40m		
2		
👤 Raghu Angadi made transition - 14/May/08 23:26		
OPEN ➡️ PATCH AVAILABLE		
13h 4m		
3		
👤 Raghu Angadi made transition - 15/May/08 18:33		
PATCH AVAILABLE ➡️ RESOLVED		
19h 7m		
1		

 Nigel Daley made transition - 21/May/08 20:06


RESOLVED → **CLOSED**

6d 1h 32m


1

▼ **People**

Assignee:

 Raghu Angadi

Reporter:

 Raghu Angadi

Votes:

 0 Vote for this issue

Watchers:

 0 Start watching this issue

▼ **Dates**

Created:

13/May/08 18:40

Updated:

08/Jul/09 16:43

Resolved:

15/May/08 18:33