



log DNS addresses on s3a init

Details

Type:	Sub-task	Status:	RESOLVED
Priority:	Minor	Resolution:	Fixed
Affects Version/s:	2.8.0	Fix Version/s:	3.3.0
Component/s:	fs/s3		
Labels:	None		
Target Version/s:	3.4.0		

Description

~~HADOOP-13874~~ has shown that network problems can kill perf, and that it's v. hard to track down, even if you turn up the logging in `hadoop.fs.s3a` and `com.amazon` layers to debug.

we could maybe improve things by printing out the `IPAddress` of the s3 endpoint, as that could help with the network tracing. Printing from within `hadoop` shows the one given to S3a, not a different one returned by any load balancer.

Issue Links

relates to

- [HADOOP-13871](#) `ITestS3AInputStreamPerformance.testTimeToOpenAndReadWholeFileBlocks` performance... RESOLVED
- [HADOOP-14551](#) S3A init hangs if you try to connect while the system is offline RESOLVED

links to

- [GitHub Pull Request #1955](#)

Activity

[Andres Perez](#) added a comment - 08/Dec/16 19:00 - edited

Would adding this to `S3AFileSystem#initialize` be enough?

```
...
    bucket = name.getHost();

    if(LOG.isDebugEnabled()) {
        LOG.debug("Bucket endpoint: "
            + InetAddress.getByName(bucket).toString());
    }
...
```

[Steve Loughran](#) added a comment - 05/Sep/17 08:58

yes, that'd work...we wouldn't be logging it except on debug, so eliminating rdns on the main path


but, we should use `NetUtils.createSocketAddrForHost()` for its handling of problems, even if that doesn't quite match what's going on in URL handling.

[Steve Loughran](#) added a comment - 05/Sep/17 09:04

Actually, [HADOOP-14551](#) highlights that the AWS SDK doesn't seem to handle DNS lookup failure (and this isn't an rDNS, is it?), so an early check for hostname lookup and fail fast if wrong could actually be better.

[Steve Loughran](#) added a comment - 12/Feb/18 12:56

+list https proxy settings being picked up in httpclient

▼  Steve Loughran added a comment - 07/Apr/20 13:06

do in bucket existence check or async -DNS delays shouldn't hold up init

▼  Hudson added a comment - 17/Apr/20 13:44

SUCCESS: Integrated in Jenkins build Hadoop-trunk-Commit #18154 (See <https://builds.apache.org/job/Hadoop-trunk-Commit/18154/>)

~~HADOOP-13873~~. log DNS addresses on s3a initialization. (stevel: rev 56350664a76b1ea8e1a942a251880ae3fab12f0c)

- (edit) `hadoop-tools/hadoop-aws/src/main/java/org/apache/hadoop/fs/s3a/S3AFileSystem.java`
- (edit) `hadoop-tools/hadoop-aws/src/main/java/org/apache/hadoop/fs/s3a/Constants.java`
- (edit) `hadoop-tools/hadoop-aws/src/main/java/org/apache/hadoop/fs/s3a/impl/NetworkBinding.java`

▼ People

Assignee:

 Mukund Thakur

Reporter:

 Steve Loughran

Votes:

 0 Vote for this issue

Watchers:

 4 Start watching this issue

▼ Dates

Created:

08/Dec/16 11:40

Updated:

17/Apr/20 17:06

Resolved:

17/Apr/20 13:21