



Hadoop HDFS / HDFS-6735

A minor optimization to avoid pread() be blocked by read() inside the same DFSInputStream

Details

Type:	Improvement	Status:	CLOSED
Priority:	Major	Resolution:	Fixed
Affects Version/s:	3.0.0-alpha1	Fix Version/s:	2.7.0
Component/s:	hdfs-client		
Labels:	None		
Hadoop Flags:	Reviewed		

Description

In current DFSInputStream impl, there're a couple of coarser-grained locks in read/pread path, and it has become a HBase read latency pain point so far. In [HDFS-6698](#), i made a minor patch against the first encountered lock, around getFileLength, in deed, after reading code and testing, it shows still other locks we could improve.

In this jira, i'll make a patch against other locks, and a simple test case to show the issue and the improved result.

This is important for HBase application, since in current HFile read path, we issue all read()/pread() requests in the same DFSInputStream for one HFile. (Multi streams solution is another story i had a plan to do, but probably will take more time than i expected)

Attachments

HDFS-6735.txt	9 kB	23/Jul/14 03:52
HDFS-6735-v2.txt	13 kB	23/Jul/14 06:59
HDFS-6735-v3.txt	11 kB	09/Nov/14 07:29
HDFS-6735-v4.txt	22 kB	23/Nov/14 21:00
HDFS-6735-v5.txt	22 kB	24/Nov/14 05:08
HDFS-6735-v6.txt	23 kB	25/Nov/14 21:10
HDFS-6735-v7.txt	22 kB	27/Nov/14 17:21
HDFS-6735-v8.txt	26 kB	28/Nov/14 20:18

Issue Links

is related to

HADOOP-15557 CryptoInputStream can't handle concurrent access; inconsistent with HDFS	OPEN
---	-------------

relates to

HDFS-7494 Checking of closed in DFSInputStream#pread() should be protected by synchronization	CLOSED
HDFS-7495 Remove updatePosition argument from DFSInputStream#getBlockAt()	CLOSED
HADOOP-9361 Strictly define the expected behavior of filesystem APIs and write tests to verify compliance	CLOSED

Sub-Tasks

1.	try to optimize DFSInputStream.getFileLength()	PATCH AVAIL...	Liang Xie
2.	Documenting DFSCClient#DFSInputStream expectations reading and preading in concurrent context	CLOSED	Michael Stack

Activity

 [Liang Xie](#) created issue - 23/Jul/14 03:13 [Liang Xie](#) made changes - 23/Jul/14 03:52

Field	Original Value	New Value
Attachment		HDFS-6735.txt [12657271]

[Liang Xie](#) made changes - 23/Jul/14 03:52

Status	Open [1]	Patch Available [10002]
--------	------------	---------------------------

[Liang Xie](#) made changes - 23/Jul/14 04:03

Description	<p>In current DFSInputStream impl, there're a couple of coarser-grained locks in read/pread path, and it has became a HBase read latency pain point so far. In HDFS-6698, i made a minor patch against the first encountered lock, around getFileLength, in deed, after reading code and testing, it shows still other locks we could improve.</p> <p>In this jira, i'll make a patch against other locks, and a simple test case to show the issue and the improved result.</p>	<p>In current DFSInputStream impl, there're a couple of coarser-grained locks in read/pread path, and it has became a HBase read latency pain point so far. In HDFS-6698, i made a minor patch against the first encountered lock, around getFileLength, in deed, after reading code and testing, it shows still other locks we could improve.</p> <p>In this jira, i'll make a patch against other locks, and a simple test case to show the issue and the improved result.</p>
-------------	--	--

This is important for HBase application, since in current HFile read path, we issue all read()/pread() requests in the same DFSInputStream for one HFile. (Multi streams solution is another story i had a plan to do, but probably will take more time than i expected)

[Liang Xie](#) made changes - 23/Jul/14 06:59

Attachment	HDFS-6735-v2.txt [12657297]
------------	---

[Steve Loughran](#) made changes - 01/Aug/14 10:08

Link	This issue relates to HADOOP-9361 [HADOOP-9361]
------	---

[Lars Hofhansl](#) made changes - 09/Nov/14 07:29

Attachment	HDFS-6735-v3.txt [12680470]
------------	---

[Lars Hofhansl](#) made changes - 23/Nov/14 20:57

Attachment	HDFS-6735-v4.txt [12683238]
------------	---

[Lars Hofhansl](#) made changes - 23/Nov/14 20:58

Attachment	HDFS-6735-v4.txt [12683238]
------------	---

[Lars Hofhansl](#) made changes - 23/Nov/14 21:00

Attachment	HDFS-6735-v4.txt [12683239]
------------	---

[Lars Hofhansl](#) made changes - 24/Nov/14 05:08

Attachment	HDFS-6735-v5.txt [12683270]
------------	---


[Lars Hofhansl](#) made changes - 25/Nov/14 21:10

Attachment	HDFS-6735-v6.txt [12683657]
------------	---

[Lars Hofhansl](#) made changes - 27/Nov/14 00:30

Attachment


HDFS-6735-v6.txt [12683956]

 [Lars Hofhansl](#) made changes - 27/Nov/14 02:24

Assignee


Liang Xie [xieliang007]

Lars Hofhansl [lhofhansl]


 [Lars Hofhansl](#) made changes - 27/Nov/14 17:06

Attachment


HDFS-6735-v6.txt [12684091]

 [Lars Hofhansl](#) made changes - 27/Nov/14 17:12


Attachment

[HDFS-6735](#)-v6.txt [12683956] [Lars Hofhansl](#) made changes - 27/Nov/14 17:12


Attachment

[HDFS-6735](#)-v6.txt [12684091] [Lars Hofhansl](#) made changes - 27/Nov/14 17:21

Attachment

[HDFS-6735-v7.txt](#) [12684096] [Lars Hofhansl](#) made changes - 28/Nov/14 20:18

Attachment

[HDFS-6735-v8.txt](#) [12684213] [Michael Stack](#) made changes - 03/Dec/14 05:01

Fix Version/s

2.7.0 [12327584]

Hadoop Flags

Reviewed [10343]


Resolution

Fixed [1]


Status

Patch Available [10002]

Resolved [5]

 [Colin McCabe](#) made changes - 12/Dec/14 21:29

Link

This issue relates to [HDFS-7494](#) [[HDFS-7494](#)] [Colin McCabe](#) made changes - 13/Dec/14 00:44


Link

This issue relates to [HDFS-7495](#) [[HDFS-7495](#)] [Vinod Kumar Vavilapalli](#) made changes - 24/Apr/15 23:22

Status

Resolved [5]

Closed [6]

 [Andrew Wang](#) made changes - 12/May/16 18:13

Affects Version/s

3.0.0-alpha1 [12335732]

Affects Version/s


3.0.0 [12320356]

 [Weiwei Yang](#) made changes - 24/Oct/17 12:42

Link

This issue is related to [HDFS-12701](#) [[HDFS-12701](#)] [Weiwei Yang](#) made changes - 24/Oct/17 13:00


Link

This issue is related to [HDFS-12701](#) [[HDFS-12701](#)] [Steve Loughran](#) made changes - 25/Jun/18 14:15

Link

This issue is related to [HADOOP-15557](#) [[HADOOP-15557](#)]**People**

Assignee:

 [Lars Hofhansl](#)

Reporter:



Liang Xie

Votes:

0 Vote for this issue

Watchers:

17 Start watching this issue

▼ Dates

Created:

23/Jul/14 03:13

Updated:

25/Jun/18 14:15

Resolved:

03/Dec/14 05:01