

Public signup for this instance is **disabled**. Go to our [Self serve sign up page](#) to request an account.



Hadoop HDFS / HDFS-7684

The host:port settings of the daemons should be trimmed before use

Details

Type:	Improvement	Status:	CLOSED
Priority:	Major	Resolution:	Fixed
Affects Version/s:	2.4.1, 2.5.1	Fix Version/s:	2.7.0
Component/s:	None		
Labels:	None		
Target Version/s:	2.7.0		
Hadoop Flags:	Reviewed		

Description

With the following setting,

```
<property>
<name>dfs.namenode.secondary.http-address</name>
<value>myhostname:50090 </value>
</property>
```

The secondary NameNode could not be started

```
$ hadoop-daemon.sh start secondarynamenode
starting secondarynamenode, logging to /home/hadoop/hadoop-2.4.1/logs/hadoop-hadoop-secondarynamenode-xxx.out
/home/hadoop/hadoop-2.4.1/bin/hdfs
Exception in thread "main" java.lang.IllegalArgumentException: Does not contain a valid host:port authority: myhostname:50090
at org.apache.hadoop.net.NetUtils.createSocketAddr(NetUtils.java:196)
at org.apache.hadoop.net.NetUtils.createSocketAddr(NetUtils.java:163)
at org.apache.hadoop.net.NetUtils.createSocketAddr(NetUtils.java:152)
at org.apache.hadoop.hdfs.server.namenode.SecondaryNameNode.getHttpAddress(SecondaryNameNode.java:203)
at org.apache.hadoop.hdfs.server.namenode.SecondaryNameNode.initialize(SecondaryNameNode.java:214)
at org.apache.hadoop.hdfs.server.namenode.SecondaryNameNode.<init>(SecondaryNameNode.java:192)
at org.apache.hadoop.hdfs.server.namenode.SecondaryNameNode.main(SecondaryNameNode.java:651)
```

We were really confused and misled by the log message: we thought about the DNS problems (changed to IP address but no success) and the network problem (tried to test the connections with no success...)

It turned out to be that the setting is not trimmed and the additional space character in the end of the setting caused the problem... OMG!!!...

Searching on the Internet, we find we are really not alone. So many users encountered similar trim problems! The following lists a few:

<http://solaimurugan.blogspot.com/2013/10/hadoop-multi-node-cluster-configuration.html>
<http://stackoverflow.com/questions/11263664/error-while-starting-the-hadoop-using-strat-all-sh>
<https://issues.apache.org/jira/browse/HDFS-2799>
<https://issues.apache.org/jira/browse/HBASE-6973>

Attachments

HDFS.7684.001.patch	17 kB	06/Feb/15 19:20
HDFS.7684.002.patch	18 kB	07/Feb/15 01:19
HDFS-7684.003.patch	18 kB	11/Feb/15 01:25
HDFS-7684.004.patch	18 kB	12/Feb/15 21:43

Issue Links

relates to

↑ [HADOOP-9869](#) Configuration.getSocketAddr()/getEnum() should use getTrimmed()

≡ CLOSED

↑ [HADOOP-11589](#) NetUtils.createSocketAddr should trim the input URI

≡ CLOSED

Activity



▼ [Xiaoyu Yao](#) added a comment - 30/Jan/15 21:09

Thanks [tianyinyin](#) for reporting this. The one that you hit can be fixed by changing the conf.get to conf.getTrimmed.

```
final String httpsAddrString = conf.get(
    DFSConfigKeys.DFS_NAMENODE_SECONDARY_HTTPS_ADDRESS_KEY,
    DFSConfigKeys.DFS_NAMENODE_SECONDARY_HTTPS_ADDRESS_DEFAULT);
InetSocketAddress httpsAddr = NetUtils.createSocketAddr(httpsAddrString);
```

Searched the call of NetUtils.createSocketAddr() in HDFS code, I found many other places with similar untrimmed host:port issues. For example in DataNodeManager#DataNodeManager() below. I think we should fix them as well with this JIRA.

```
this.defaultXferPort = NetUtils.createSocketAddr(
    conf.get(DFSConfigKeys.DFS_DATANODE_ADDRESS_KEY,
    DFSConfigKeys.DFS_DATANODE_ADDRESS_DEFAULT)).getPort();
this.defaultInfoPort = NetUtils.createSocketAddr(
    conf.get(DFSConfigKeys.DFS_DATANODE_HTTP_ADDRESS_KEY,
    DFSConfigKeys.DFS_DATANODE_HTTP_ADDRESS_DEFAULT)).getPort();
this.defaultInfoSecurePort = NetUtils.createSocketAddr(
    conf.get(DFSConfigKeys.DFS_DATANODE_HTTPS_ADDRESS_KEY,
    DFSConfigKeys.DFS_DATANODE_HTTPS_ADDRESS_DEFAULT)).getPort();
this.defaultIpcPort = NetUtils.createSocketAddr(
    conf.get(DFSConfigKeys.DFS_DATANODE_IPC_ADDRESS_KEY,
    DFSConfigKeys.DFS_DATANODE_IPC_ADDRESS_DEFAULT)).getPort();
```

▼ [Tianyinyin Xu](#) added a comment - 30/Jan/15 21:17

Yes, exactly. It seems that Hadoop has a bunch of such trimming issues that bothered a number of users...

Thanks, Xiaoyu!

~t

▼ [Anu Engineer](#) added a comment - 06/Feb/15 19:20

Trims the Configuration values before reading them in.

▼ [Hadoop QA](#) added a comment - 06/Feb/15 22:43

-1 overall. Here are the results of testing the latest attachment
<http://issues.apache.org/jira/secure/attachment/12697096/HDFS.7684.001.patch>
 against trunk revision eaab959.

+1 @author. The patch does not contain any @author tags.

+1 tests included. The patch appears to include 2 new or modified test files.

+1 javac. The applied patch does not increase the total number of javac compiler warnings.

+1 javadoc. There were no new javadoc warning messages.

+1 eclipse:eclipse. The patch built with eclipse:eclipse.

+1 findbugs. The patch does not introduce any new Findbugs (version 2.0.3) warnings.

-1 release audit. The applied patch generated 1 release audit warnings.

-1 core tests. The patch failed these unit tests in hadoop-hdfs-project/hadoop-hdfs:


org.apache.hadoop.hdfs.server.balancer.TestBalancer
 org.apache.hadoop.hdfs.TestLeaseRecovery2

Test results: <https://builds.apache.org/job/PreCommit-HDFS-Build/9464/testReport/>

Release audit warnings: <https://builds.apache.org/job/PreCommit-HDFS-Build/9464/artifact/patchprocess/patchReleaseAuditProblems.txt>

Console output: <https://builds.apache.org/job/PreCommit-HDFS-Build/9464/console>

This message is automatically generated.


▼  Akira Ajsaka added a comment - 06/Feb/15 23:32

Hi anu, thank you for creating the patch. I'm +1 for trimming the values.

-1 release audit. The applied patch generated 1 release audit warnings.

Would you add a license header to TestMalformedURLs.java?


In addition, would you please change the indent to 2 spaces instead of 4 in TestMalformedURLs.java?

▼  Anu Engineer added a comment - 07/Feb/15 01:19

update :

1) Added Lic. to test code file

2) Fixed the Tabs side

▼  Hadoop QA added a comment - 07/Feb/15 04:58

-1 overall. Here are the results of testing the latest attachment
<http://issues.apache.org/jira/secure/attachment/12697202/HDFS.7684.002.patch>
against trunk revision 8de80ff.

+1 @author. The patch does not contain any @author tags.

+1 tests included. The patch appears to include 2 new or modified test files.

+1 javac. The applied patch does not increase the total number of javac compiler warnings.

+1 javadoc. There were no new javadoc warning messages.

+1 eclipse:eclipse. The patch built with eclipse:eclipse.

+1 findbugs. The patch does not introduce any new Findbugs (version 2.0.3) warnings.

+1 release audit. The applied patch does not increase the total number of release audit warnings.

-1 core tests. The patch failed these unit tests in hadoop-hdfs-project/hadoop-hdfs:

org.apache.hadoop.cli.TestHDFSCLI

Test results: <https://builds.apache.org/job/PreCommit-HDFS-Build/9480//testReport/>

Console output: <https://builds.apache.org/job/PreCommit-HDFS-Build/9480//console>


This message is automatically generated.

▼  Anu Engineer added a comment - 09/Feb/15 18:00

The failure does not seem to be related to this change set.

Thanks

Anu

▼  Akira Ajsaka added a comment - 09/Feb/15 22:41

2) Fixed the Tabs side


I didn't see the indent was fixed to 2 whitespaces. Would you please update the patch?

Also, would you remove some empty lines from the test? One empty line is sufficient.

```
+import org.apache.hadoop.hdfs.DFSConfigKeys;  
+  
+  
+  
+  
+  
+public class TestMalformedURLs {
```

▼  Anu Engineer added a comment - 11/Feb/15 01:25

Hopefully this really fixes the spaces issue. Sorry about missing that earlier

▼  Hadoop QA added a comment - 11/Feb/15 07:52

-1 overall. Here are the results of testing the latest attachment

<http://issues.apache.org/jira/secure/attachment/12697953/HDFS-7684.003.patch>
against trunk revision 7c6b654.

+1 @author. The patch does not contain any @author tags.

+1 tests included. The patch appears to include 2 new or modified test files.

+1 javac. The applied patch does not increase the total number of javac compiler warnings.

+1 javadoc. There were no new javadoc warning messages.

+1 eclipse:eclipse. The patch built with eclipse:eclipse.

-1 findbugs. The patch appears to introduce 1 new Findbugs (version 2.0.3) warnings.

+1 release audit. The applied patch does not increase the total number of release audit warnings.


+1 core tests. The patch passed unit tests in hadoop-hdfs-project/hadoop-hdfs.

Test results: <https://builds.apache.org/job/PreCommit-HDFS-Build/9529//testReport/>

Findbugs warnings: <https://builds.apache.org/job/PreCommit-HDFS-Build/9529//artifact/patchprocess/newPatchFindbugsWarningshadoop-hdfs.html>


Console output: <https://builds.apache.org/job/PreCommit-HDFS-Build/9529//console>

This message is automatically generated.

▼  Chris Nauroth added a comment - 11/Feb/15 18:08

The Findbugs warning is likely unrelated, fixed last night in [HDFS-7754](#). I submitted another Jenkins run:

<https://builds.apache.org/job/PreCommit-HDFS-Build/9542/>

▼  Hadoop QA added a comment - 11/Feb/15 20:55

-1 overall. Here are the results of testing the latest attachment

<http://issues.apache.org/jira/secure/attachment/12697953/HDFS-7684.003.patch>
against trunk revision 22441ab.

+1 @author. The patch does not contain any @author tags.

+1 tests included. The patch appears to include 2 new or modified test files.

+1 javac. The applied patch does not increase the total number of javac compiler warnings.

+1 javadoc. There were no new javadoc warning messages.

+1 eclipse:eclipse. The patch built with eclipse:eclipse.

+1 findbugs. The patch does not introduce any new Findbugs (version 2.0.3) warnings.

+1 release audit. The applied patch does not increase the total number of release audit warnings.


-1 core tests. The patch failed these unit tests in hadoop-hdfs-project/hadoop-hdfs:

org.apache.hadoop.hdfs.server.namenode.ha.TestRetryCacheWithHA


Test results: <https://builds.apache.org/job/PreCommit-HDFS-Build/9542//testReport/>

Console output: <https://builds.apache.org/job/PreCommit-HDFS-Build/9542//console>


This message is automatically generated.

▼  Akira Ajisaka added a comment - 12/Feb/15 19:11

Looking around the code, I'm thinking it would make more sense to trim the value of the parameters in `NetUtils.createSocketAddr`. What do you think, [cnauroth](#) and [anu](#)?


▼  Akira Ajisaka added a comment - 12/Feb/15 19:18

I think if the parameters were trimmed in `NetUtils`, the problem in [HADOOP-9869](#) would be fixed also.

▼  [Chris Nauroth](#) added a comment - 12/Feb/15 19:57

Trimming in `NetUtils` might be a valuable additional thing to do in its own independent patch, but I also think trimming when getting out of the `Configuration` is helpful. This can make the code more robust if any logic needs to be applied to the configuration value before it gets passed to `NetUtils`. One example is `NameNode#getServiceAddress`, which needs to check for empty string.


Anu, thank you for the patch. I have just a few minor formatting nitpicks in `TestMalformedURLs`. In `testTryStartingCluster`, there is a typo "configuration". Then, in `tearDown`, the `cluster.shutdown()` line needs to be indented by 2 spaces instead of 4. I'll be +1 after those changes.

▼  [Akira Ajisaka](#) added a comment - 12/Feb/15 20:12

I also think trimming when getting out of the `Configuration` is helpful. This can make the code more robust if any logic needs to be applied to the configuration value before it gets passed to `NetUtils`.


I agree. Now I have one minor comment: Would you simply use `assertNotEquals` instead of `assertTrue` in the test?

```
+ assertTrue(config.get(DFSConfigKeys.DFS_NAMENODE_HTTP_ADDRESS_KEY)  
+             != config.getTrimmed(DFSConfigKeys.DFS_NAMENODE_HTTP_ADDRESS_KEY));
```


▼  [Akira Ajisaka](#) added a comment - 12/Feb/15 20:54

Trimming in `NetUtils` might be a valuable additional thing to do in its own independent patch


Filed [HADOOP-11589](#).

▼  [Anu Engineer](#) added a comment - 12/Feb/15 21:44

Addressed comments from [cnauroth](#) and [ajisakaa](#)

▼  [Chris Nauroth](#) added a comment - 12/Feb/15 22:03

+1 for patch v004, pending Jenkins run. Akira, thank you for the help with the code review. Good catch on the assert.

▼  [Hadoop QA](#) added a comment - 13/Feb/15 01:22

+1 overall. Here are the results of testing the latest attachment <http://issues.apache.org/jira/secure/attachment/12698533/HDFS-7684.004.patch> against trunk revision 6f5290b.

+1 @author. The patch does not contain any @author tags.

+1 tests included. The patch appears to include 2 new or modified test files.

+1 javac. The applied patch does not increase the total number of javac compiler warnings.

+1 javadoc. There were no new javadoc warning messages.

+1 eclipse:eclipse. The patch built with eclipse:eclipse.

+1 findbugs. The patch does not introduce any new Findbugs (version 2.0.3) warnings.


+1 release audit. The applied patch does not increase the total number of release audit warnings.

+1 core tests. The patch passed unit tests in hadoop-hdfs-project/hadoop-hdfs.

Test results: <https://builds.apache.org/job/PreCommit-HDFS-Build/9564//testReport/>

Console output: <https://builds.apache.org/job/PreCommit-HDFS-Build/9564//console>

This message is automatically generated.

▼  [Akira Ajisaka](#) added a comment - 13/Feb/15 01:29

+1. Thanks [anu](#) for the contribution, and thanks [cnauroth](#) for the review. Committing this.


▼  [Akira Ajisaka](#) added a comment - 13/Feb/15 01:43

Committed this to trunk and branch-2. Thanks again, [anu](#) and [cnauroth](#).

▼  Hudson added a comment - 13/Feb/15 01:53

FAILURE: Integrated in Hadoop-trunk-Commit #7098 (See <https://builds.apache.org/job/Hadoop-trunk-Commit/7098/>) [HDFS-7684](#). The host:port settings of the daemons should be trimmed before use. Contributed by Anu Engineer. (aajisaka: rev 46b6d23e8fbed4c2ba537dd752116c173805bca7)

- hadoop-hdfs-project/hadoop-hdfs/CHANGES.txt
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/SecondaryNameNode.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NameNode.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/datanode/DataNode.java
- hadoop-hdfs-project/hadoop-hdfs/src/test/java/org/apache/hadoop/hdfs/server/namenode/TestMalformedURLs.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/blockmanagement/DatanodeManager.java
- hadoop-hdfs-project/hadoop-hdfs/src/test/resources/hdfs-site.malformed.xml
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/datanode/web/DatanodeHttpServer.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NameNodeHttpServer.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/ImageServlet.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/BackupNode.java

▼  Hudson added a comment - 13/Feb/15 10:41

FAILURE: Integrated in Hadoop-Yarn-trunk-Java8 #103 (See <https://builds.apache.org/job/Hadoop-Yarn-trunk-Java8/103/>) [HDFS-7684](#). The host:port settings of the daemons should be trimmed before use. Contributed by Anu Engineer. (aajisaka: rev 46b6d23e8fbed4c2ba537dd752116c173805bca7)

- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NameNodeHttpServer.java
- hadoop-hdfs-project/hadoop-hdfs/src/test/java/org/apache/hadoop/hdfs/server/namenode/TestMalformedURLs.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/BackupNode.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/ImageServlet.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NameNode.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/datanode/DataNode.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/blockmanagement/DatanodeManager.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/datanode/web/DatanodeHttpServer.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/SecondaryNameNode.java
- hadoop-hdfs-project/hadoop-hdfs/src/test/resources/hdfs-site.malformed.xml
- hadoop-hdfs-project/hadoop-hdfs/CHANGES.txt

▼  Hudson added a comment - 13/Feb/15 10:44

FAILURE: Integrated in Hadoop-Yarn-trunk #837 (See <https://builds.apache.org/job/Hadoop-Yarn-trunk/837/>) [HDFS-7684](#). The host:port settings of the daemons should be trimmed before use. Contributed by Anu Engineer. (aajisaka: rev 46b6d23e8fbed4c2ba537dd752116c173805bca7)

- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/SecondaryNameNode.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/ImageServlet.java
- hadoop-hdfs-project/hadoop-hdfs/CHANGES.txt
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/datanode/DataNode.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NameNode.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/BackupNode.java
- hadoop-hdfs-project/hadoop-hdfs/src/test/resources/hdfs-site.malformed.xml
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/blockmanagement/DatanodeManager.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/datanode/web/DatanodeHttpServer.java
- hadoop-hdfs-project/hadoop-hdfs/src/test/java/org/apache/hadoop/hdfs/server/namenode/TestMalformedURLs.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NameNodeHttpServer.java

▼  Hudson added a comment - 13/Feb/15 14:26

SUCCESS: Integrated in Hadoop-Hdfs-trunk #2035 (See <https://builds.apache.org/job/Hadoop-Hdfs-trunk/2035/>) [HDFS-7684](#). The host:port settings of the daemons should be trimmed before use. Contributed by Anu Engineer. (aajisaka: rev 46b6d23e8fbed4c2ba537dd752116c173805bca7)

- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/BackupNode.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NameNode.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/ImageServlet.java

- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NameNodeHttpServer.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/datanode/web/DatanodeHttpServer.java
- hadoop-hdfs-project/hadoop-hdfs/src/test/resources/hdfs-site.malformed.xml
- hadoop-hdfs-project/hadoop-hdfs/CHANGES.txt
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/SecondaryNameNode.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/datanode/DataNode.java
- hadoop-hdfs-project/hadoop-hdfs/src/test/java/org/apache/hadoop/hdfs/server/namenode/TestMalformedURLs.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/blockmanagement/DatanodeManager.java

▼  Hudson added a comment - 13/Feb/15 15:13

FAILURE: Integrated in Hadoop-Mapreduce-trunk-Java8 #104 (See <https://builds.apache.org/job/Hadoop-Mapreduce-trunk-Java8/104/>)

[HDFS-7684](#). The host:port settings of the daemons should be trimmed before use. Contributed by Anu Engineer. (aajisaka: rev 46b6d23e8fbed4c2ba537dd752116c173805bca7)


- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/blockmanagement/DatanodeManager.java
- hadoop-hdfs-project/hadoop-hdfs/src/test/resources/hdfs-site.malformed.xml
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NameNodeHttpServer.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/ImageServlet.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/SecondaryNameNode.java
- hadoop-hdfs-project/hadoop-hdfs/src/test/java/org/apache/hadoop/hdfs/server/namenode/TestMalformedURLs.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/datanode/DataNode.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/datanode/web/DatanodeHttpServer.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NameNode.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/BackupNode.java
- hadoop-hdfs-project/hadoop-hdfs/CHANGES.txt

▼  Hudson added a comment - 13/Feb/15 15:25

FAILURE: Integrated in Hadoop-Mapreduce-trunk #2054 (See <https://builds.apache.org/job/Hadoop-Mapreduce-trunk/2054/>)

[HDFS-7684](#). The host:port settings of the daemons should be trimmed before use. Contributed by Anu Engineer. (aajisaka: rev 46b6d23e8fbed4c2ba537dd752116c173805bca7)

- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/BackupNode.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NameNode.java
- hadoop-hdfs-project/hadoop-hdfs/src/test/resources/hdfs-site.malformed.xml
- hadoop-hdfs-project/hadoop-hdfs/src/test/java/org/apache/hadoop/hdfs/server/namenode/TestMalformedURLs.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/blockmanagement/DatanodeManager.java
- hadoop-hdfs-project/hadoop-hdfs/CHANGES.txt
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/ImageServlet.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/datanode/web/DatanodeHttpServer.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NameNodeHttpServer.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/SecondaryNameNode.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/datanode/DataNode.java

▼  Chris Nauroth added a comment - 13/Feb/15 18:08

aajisaka, thank you for the commit!

▼  Hudson added a comment - 17/Feb/15 08:12

SUCCESS: Integrated in Hadoop-Hdfs-trunk-Java8 #97 (See <https://builds.apache.org/job/Hadoop-Hdfs-trunk-Java8/97/>)


[HDFS-7684](#). The host:port settings of the daemons should be trimmed before use. Contributed by Anu Engineer. (aajisaka: rev 46b6d23e8fbed4c2ba537dd752116c173805bca7)

- hadoop-hdfs-project/hadoop-hdfs/CHANGES.txt
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NameNodeHttpServer.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/datanode/web/DatanodeHttpServer.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/SecondaryNameNode.java
- hadoop-hdfs-project/hadoop-hdfs/src/test/resources/hdfs-site.malformed.xml
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/BackupNode.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/datanode/DataNode.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/blockmanagement/DatanodeManager.java
- hadoop-hdfs-project/hadoop-hdfs/src/test/java/org/apache/hadoop/hdfs/server/namenode/TestMalformedURLs.java

- `hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/ImageServlet.java`
- `hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NameNode.java`

▼ People

Assignee:

 Anu Engineer

Reporter:

 Tianyin Xu

Votes:

 0 Vote for this issue

Watchers:

 6 Start watching this issue

▼ Dates

Created:

27/Jan/15 10:05

Updated:

10/Apr/15 20:29

Resolved:

13/Feb/15 01:43