

Public signup for this instance is **disabled**. Go to our [Self serve sign up page](#) to request an account.



Hadoop HDFS / HDFS-8025

# Addendum fix for HDFS-3087 Decomissioning on NN restart can complete without blocks being replicated

Details

Type:	Bug	Status:	<b>CLOSED</b>
Priority:	Major	Resolution:	Fixed
Affects Version/s:	None	Fix Version/s:	2.7.0
Component/s:	None		
Labels:	None		

Description

Per discussion with [andrew.wang](#) on [HDFS-7441](#), we should include [HDFS-3087](#) and enhance the unit test.

Attachments

<a href="#">HDFS-8025.patch</a>	4 kB	30/Mar/15 23:25
<a href="#">HDFS-8025-2.patch</a>	4 kB	03/Apr/15 04:35

Activity

- [Ming Ma](#) added a comment - 30/Mar/15 23:25
- The reason unit test "TestDecommission#testDecommissionWithNamenodeRestart" can pass because blockReport was processed earlier than DecommissonManager thread. Thus the test case doesn't verify the scenario.

The patch added checkBlockReportReceived back. It also improves unit tests in the following ways. a) increase the blockreport interval and make it more likely for DecommissonManager thread to process the node before blockReport is processed. b) Check if the blocks are properly replicated right after node is transitioned to Decommissioned. There is no need to wait for the block replication given nodes should only be transitioned to Decommissioned after block are properly replicated.
- [Hadoop QA](#) added a comment - 31/Mar/15 03:36
- 1 overall**. Here are the results of testing the latest attachment <http://issues.apache.org/jira/secure/attachment/12708258/HDFS-8025.patch> against trunk revision 1a495fb.

**+1 @author**. The patch does not contain any @author tags.

**+1 tests included**. The patch appears to include 1 new or modified test files.

**+1 javac**. The applied patch does not increase the total number of javac compiler warnings.

**+1 javadoc**. There were no new javadoc warning messages.

**+1 eclipse:eclipse**. The patch built with eclipse:eclipse.

**+1 findbugs**. The patch does not introduce any new Findbugs (version 2.0.3) warnings.

**+1 release audit**. The applied patch does not increase the total number of release audit warnings.

**-1 core tests**. The patch failed these unit tests in hadoop-hdfs-project/hadoop-hdfs:

org.apache.hadoop.hdfs.TestEncryptionZonesWithKMS


The following test timeouts occurred in hadoop-hdfs-project/hadoop-hdfs:

org.apache.hadoop.hdfs.server.blockmanagement.TestDatanodeManager

Test results: <https://builds.apache.org/job/PreCommit-HDFS-Build/10118//testReport/>

Console output: <https://builds.apache.org/job/PreCommit-HDFS-Build/10118//console>

This message is automatically generated.


▼  Andrew Wang added a comment - 01/Apr/15 20:01

Hey Ming, thanks for working on this. Did you consider putting the new check in `BlockManager#isNodeHealthyForDecommission`? This is a better home for this check (since this does relate to node health), and I think it should be approximately as efficient since a node will have 0 blocks before its FBR.

It would also be good to add some logging for this case, since a perpetual complaint about decom is the lack of visibility as to why a node is stuck.

▼  Ming Ma added a comment - 03/Apr/15 04:35

Thanks Andrew. Here is the updated patch with your suggestions.

▼  Hadoop QA added a comment - 03/Apr/15 08:15

-1 overall. Here are the results of testing the latest attachment <http://issues.apache.org/jira/secure/attachment/12709165/HDFS-8025-2.patch> against trunk revision 72f6bd4.

+1 @author. The patch does not contain any @author tags.

+1 tests included. The patch appears to include 1 new or modified test files.

+1 javac. The applied patch does not increase the total number of javac compiler warnings.

+1 javadoc. There were no new javadoc warning messages.

+1 eclipse:eclipse. The patch built with eclipse:eclipse.

+1 findbugs. The patch does not introduce any new Findbugs (version 2.0.3) warnings.


+1 release audit. The applied patch does not increase the total number of release audit warnings.

-1 core tests. The patch failed these unit tests in hadoop-hdfs-project/hadoop-hdfs:


org.apache.hadoop.hdfs.server.blockmanagement.TestDatanodeManager  
org.apache.hadoop.hdfs.server.namenode.ha.TestRetryCacheWithHA

Test results: <https://builds.apache.org/job/PreCommit-HDFS-Build/10170/testReport/>  
Console output: <https://builds.apache.org/job/PreCommit-HDFS-Build/10170/console>

This message is automatically generated.

▼  Andrew Wang added a comment - 08/Apr/15 23:06

+1 LGTM, will commit shortly

▼  Andrew Wang added a comment - 08/Apr/15 23:10

Committed to trunk and branch-2, thanks Ming.

▼  Hudson added a comment - 08/Apr/15 23:30

FAILURE: Integrated in Hadoop-trunk-Commit #7540 (See <https://builds.apache.org/job/Hadoop-trunk-Commit/7540/>)  
~~HDFS-8025~~. Addendum fix for ~~HDFS-3087~~ Decommissioning on NN restart can complete without blocks being replicated. Contributed by Ming Ma. (wang: rev 5a540c3d3107199f4632e2ad7ee8ff913b107a04)

- hadoop-hdfs-project/hadoop-hdfs/CHANGES.txt
- hadoop-hdfs-project/hadoop-hdfs/src/test/java/org/apache/hadoop/hdfs/TestDecommission.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/blockmanagement/BlockManager.java

▼  Hudson added a comment - 09/Apr/15 10:41

FAILURE: Integrated in Hadoop-Yarn-trunk-Java8 #158 (See <https://builds.apache.org/job/Hadoop-Yarn-trunk-Java8/158/>)  
~~HDFS-8025~~. Addendum fix for ~~HDFS-3087~~ Decommissioning on NN restart can complete without blocks being replicated. Contributed by Ming Ma. (wang: rev 5a540c3d3107199f4632e2ad7ee8ff913b107a04)

- hadoop-hdfs-project/hadoop-hdfs/src/test/java/org/apache/hadoop/hdfs/TestDecommission.java
- hadoop-hdfs-project/hadoop-hdfs/CHANGES.txt

- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/blockmanagement/BlockManager.java

▼  Hudson added a comment - 09/Apr/15 11:34

FAILURE: Integrated in Hadoop-Hdfs-trunk #2090 (See <https://builds.apache.org/job/Hadoop-Hdfs-trunk/2090/>)  
[HDFS-8025](#). Addendum fix for [HDFS-3087](#) Decommissioning on NN restart can complete without blocks being replicated. Contributed by Ming Ma. (wang: rev 5a540c3d3107199f4632e2ad7ee8ff913b107a04)

- hadoop-hdfs-project/hadoop-hdfs/src/test/java/org/apache/hadoop/hdfs/TestDecommission.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/blockmanagement/BlockManager.java
- hadoop-hdfs-project/hadoop-hdfs/CHANGES.txt

▼  Hudson added a comment - 09/Apr/15 11:35


FAILURE: Integrated in Hadoop-Hdfs-trunk-Java8 #149 (See <https://builds.apache.org/job/Hadoop-Hdfs-trunk-Java8/149/>)  
[HDFS-8025](#). Addendum fix for [HDFS-3087](#) Decommissioning on NN restart can complete without blocks being replicated. Contributed by Ming Ma. (wang: rev 5a540c3d3107199f4632e2ad7ee8ff913b107a04)

- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/blockmanagement/BlockManager.java
- hadoop-hdfs-project/hadoop-hdfs/src/test/java/org/apache/hadoop/hdfs/TestDecommission.java
- hadoop-hdfs-project/hadoop-hdfs/CHANGES.txt

▼  Hudson added a comment - 09/Apr/15 11:56

SUCCESS: Integrated in Hadoop-Yarn-trunk #892 (See <https://builds.apache.org/job/Hadoop-Yarn-trunk/892/>)  
[HDFS-8025](#). Addendum fix for [HDFS-3087](#) Decommissioning on NN restart can complete without blocks being replicated. Contributed by Ming Ma. (wang: rev 5a540c3d3107199f4632e2ad7ee8ff913b107a04)

- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/blockmanagement/BlockManager.java
- hadoop-hdfs-project/hadoop-hdfs/src/test/java/org/apache/hadoop/hdfs/TestDecommission.java
- hadoop-hdfs-project/hadoop-hdfs/CHANGES.txt

▼  Kihwal Lee added a comment - 09/Apr/15 15:03

I put this in branch-2.7.

▼  Hudson added a comment - 09/Apr/15 18:14

FAILURE: Integrated in Hadoop-Mapreduce-trunk-Java8 #159 (See <https://builds.apache.org/job/Hadoop-Mapreduce-trunk-Java8/159/>)  
[HDFS-8025](#). Addendum fix for [HDFS-3087](#) Decommissioning on NN restart can complete without blocks being replicated. Contributed by Ming Ma. (wang: rev 5a540c3d3107199f4632e2ad7ee8ff913b107a04)

- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/blockmanagement/BlockManager.java
- hadoop-hdfs-project/hadoop-hdfs/CHANGES.txt
- hadoop-hdfs-project/hadoop-hdfs/src/test/java/org/apache/hadoop/hdfs/TestDecommission.java

▼  Hudson added a comment - 09/Apr/15 18:32

FAILURE: Integrated in Hadoop-Mapreduce-trunk #2108 (See <https://builds.apache.org/job/Hadoop-Mapreduce-trunk/2108/>)  
[HDFS-8025](#). Addendum fix for [HDFS-3087](#) Decommissioning on NN restart can complete without blocks being replicated. Contributed by Ming Ma. (wang: rev 5a540c3d3107199f4632e2ad7ee8ff913b107a04)

- hadoop-hdfs-project/hadoop-hdfs/src/test/java/org/apache/hadoop/hdfs/TestDecommission.java
- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/blockmanagement/BlockManager.java
- hadoop-hdfs-project/hadoop-hdfs/CHANGES.txt

▼  Ming Ma added a comment - 13/Apr/15 15:14

Thanks Andrew and Kihwal.

#### ▼ People

Assignee:

 Ming Ma

Reporter:

 Ming Ma

Votes:

0

Vote for this issue

Watchers:

5

Start watching this issue

▼

Dates

---

Created:

30/Mar/15 22:06

Updated:

24/Apr/15 23:22

Resolved:

08/Apr/15 23:10