Hadoop HDFS / HDFS-9906

Remove spammy log spew when a datanode is restarted

Details

Type:	⬆ Improvement	Status:	**CLOSED**
Priority:	🔺 Major	Resolution:	Fixed
Affects Version/s:	2.7.2	Fix Version/s:	2.8.0, 2.7.3, 3.0.0-alpha1
Component/s:	namenode		
Labels:	None		
Hadoop Flags:	Reviewed		

Description

```
WARN BlockStateChange: BLOCK* addStoredBlock: Redundant addStoredBlock request received for blk_1109897077_36157149 on node
192.168.1.1:50010 size 268435456
```

This happens waaaay too much to add any useful information. We should either move this to a different level or only warn once per machine.

Attachments

📄 HDFS-9906.patch	1 kB	05/Mar/16 02:10

Activity

⬆

Haohui Mai added a comment - 04/Mar/16 21:34

The information is useful for debugging but I agree that we don't need to print all of them. We can consider doing a rate limit on the logs. For example, print only 5 logs / minutes at most.

Arpit Agarwal added a comment - 04/Mar/16 22:15

Rate limiting is overkill. +1 on just making it debug.

Brahma Reddy Battula added a comment - 05/Mar/16 02:10

Me too +1 on making it debug..

Arpit Agarwal added a comment - 05/Mar/16 02:18

+1 pending Jenkins. Thanks for the quick patch brahmareddy.

I'll hold off committing until next week in case there are any objections.

Hadoop QA added a comment - 05/Mar/16 05:05

❌ **-1 overall**

Vote	Subsystem	Runtime	Comment
0	reexec	12m 56s	Docker mode activated.
+1	@author	0m 0s	The patch does not contain any @author tags.
-1	test4tests	0m 0s	The patch doesn't appear to include any new or modified tests. Please justify why no new tests are needed for this patch. Also please list what manual steps were performed to verify this patch.
+1	mvninstall	7m 42s	trunk passed

+1	compile	0m 57s	trunk passed with JDK v1.8.0_74
+1	compile	0m 43s	trunk passed with JDK v1.7.0_95
+1	checkstyle	0m 21s	trunk passed
+1	mvnsite	0m 53s	trunk passed
+1	mvneclipse	0m 12s	trunk passed
+1	findbugs	1m 54s	trunk passed
+1	javadoc	1m 19s	trunk passed with JDK v1.8.0_74
+1	javadoc	1m 49s	trunk passed with JDK v1.7.0_95
+1	mvninstall	1m 3s	the patch passed
+1	compile	1m 16s	the patch passed with JDK v1.8.0_74
+1	javac	1m 15s	the patch passed
+1	compile	0m 57s	the patch passed with JDK v1.7.0_95
+1	javac	0m 57s	the patch passed
+1	checkstyle	0m 21s	the patch passed
+1	mvnsite	0m 54s	the patch passed
+1	mvneclipse	0m 11s	the patch passed
+1	whitespace	0m 0s	Patch has no whitespace issues.
+1	findbugs	2m 5s	the patch passed
+1	javadoc	1m 9s	the patch passed with JDK v1.8.0_74
+1	javadoc	1m 48s	the patch passed with JDK v1.7.0_95
-1	unit	69m 4s	hadoop-hdfs in the patch failed with JDK v1.8.0_74.
-1	unit	60m 3s	hadoop-hdfs in the patch failed with JDK v1.7.0_95.
+1	asflicense	0m 25s	Patch does not generate ASF License warnings.
		170m 28s	

Reason	Tests
JDK v1.8.0_74 Failed junit tests	hadoop.hdfs.server.namenode.TestEditLog
	hadoop.hdfs.server.blockmanagement.TestRBWBlockInvalidation
	hadoop.hdfs.TestFileAppend
	hadoop.metrics2.sink.TestRollingFileSystemSinkWithHdfs
	hadoop.hdfs.TestBlockStoragePolicy
JDK v1.7.0_95 Failed junit tests	hadoop.hdfs.server.blockmanagement.TestBlockManager
	hadoop.hdfs.server.namenode.TestEditLog
	hadoop.hdfs.TestDataTransferKeepalive

Subsystem	Report/Notes
Docker	Image:yetus/hadoop:0ca8df7
JIRA Patch URL	https://issues.apache.org/jira/secure/attachment/12791590/HDFS-9906.patch
JIRA Issue	HDFS-9906
Optional Tests	asflicense compile javac javadoc mvninstall mvnsite unit findbugs checkstyle
uname	Linux 0d16c4758f02 3.13.0-36-lowlatency #63-Ubuntu SMP PREEMPT Wed Sep 3 21:56:12 UTC 2014 x86_64 x86_64 x86_64 GNU/Linux
Build tool	maven

Personality	/testptch/hadoop/patchprocess/precommit/personality/provided.sh
git revision	trunk / 2759689
Default Java	1.7.0_95
Multi-JDK versions	/usr/lib/jvm/java-8-oracle:1.8.0_74 /usr/lib/jvm/java-7-openjdk-amd64:1.7.0_95
findbugs	v3.0.0
unit	https://builds.apache.org/job/PreCommit-HDFS-Build/14724/artifact/patchprocess/patch-unit-hadoop-hdfs-project_hadoop-hdfs-jdk1.8.0_74.txt
unit	https://builds.apache.org/job/PreCommit-HDFS-Build/14724/artifact/patchprocess/patch-unit-hadoop-hdfs-project_hadoop-hdfs-jdk1.7.0_95.txt
unit test logs	https://builds.apache.org/job/PreCommit-HDFS-Build/14724/artifact/patchprocess/patch-unit-hadoop-hdfs-project_hadoop-hdfs-jdk1.8.0_74.txt https://builds.apache.org/job/PreCommit-HDFS-Build/14724/artifact/patchprocess/patch-unit-hadoop-hdfs-project_hadoop-hdfs-jdk1.7.0_95.txt
JDK v1.7.0_95 Test Results	https://builds.apache.org/job/PreCommit-HDFS-Build/14724/testReport/
modules	C: hadoop-hdfs-project/hadoop-hdfs U: hadoop-hdfs-project/hadoop-hdfs
Console output	https://builds.apache.org/job/PreCommit-HDFS-Build/14724/console
Powered by	Apache Yetus 0.3.0-SNAPSHOT http://yetus.apache.org

This message was automatically generated.

Chris Nauroth added a comment - 05/Mar/16 05:26

I am also +1 for the patch to downtune this logging to debug level. brahmareddy, thank you for the patch.

Elliott Neil Clark added a comment - 05/Mar/16 23:36

+1

Arpit Agarwal added a comment - 07/Mar/16 20:22

I've committed this for 2.8.0. Thank you for the contribution brahmareddy, and thanks cnauroth and eclark for the reviews.

Hudson added a comment - 07/Mar/16 20:45

FAILURE: Integrated in Hadoop-trunk-Commit #9438 (See https://builds.apache.org/job/Hadoop-trunk-Commit/9438/)
HDFS-9906. Remove spammy log spew when a datanode is restarted. (arp: rev 724d2299cd2516d90c030f6e20d814cceb439228)

- hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/blockmanagement/BlockManager.java

---

Xiaoyu Yao added a comment - 15/Jun/16 06:30

Cherrypicked to branch-2.7.

Akira Ajisaka added a comment - 15/Jun/16 06:49

Updated fix version. Hi xyao, now branch-2.7.3 is cut, so if you want to include this in 2.7.3, you need to cherry-pick this to branch-2.7.3 as well.

Xiaoyu Yao added a comment - 15/Jun/16 17:03

Thanks ajisakaa, cherry-pick to 2.7.3 and update fix version to 2.7.3. cc: vinodkv.

Vinod Kumar Vavilapalli added a comment - 25/Aug/16 23:00

Closing the JIRA as part of 2.7.3 release.

---

darkz added a comment - 29/Sep/18 11:39 - edited

you can use

```
hadoop daemonlog –setlevel 192.168.120.180:50070 BlockStateChange ERROR
```

to change the log level online without the namenode restart

▾ People

Assignee:

🔘 Brahma Reddy Battula

Reporter:

🔘 Elliott Neil Clark

Votes:

0 Vote for this issue

Watchers:

15 Start watching this issue

▾ Dates

Created:

04/Mar/16 17:57

Updated:

29/Sep/18 11:39

Resolved:

07/Mar/16 20:22