

Public signup for this instance is **disabled**. Go to our [Self serve sign up page](#) to request an account.



Hadoop HDFS / HDFS-14340

Lower the log level when can't get postOpAttr

Details

Type:	Improvement	Status:	RESOLVED
Priority:	Minor	Resolution:	Fixed
Affects Version/s:	None	Fix Version/s:	3.3.0, 3.2.1, 3.1.3
Component/s:	nfs		
Labels:	easyfix		
Hadoop Flags:	Reviewed		

Description

I think should lower the log level when can't get postOpAttr in *hadoop-2.8.5-src/hadoop-hdfs-project/hadoop-hdfs-nfs/src/main/java/org/apache/hadoop/hdfs/nfs/nfs3/RpcProgramNfs3.java*.

```
*****The first code snippet*****
//the problematic log level ERROR, at line 1044
try {
    dirWcc = Nfs3Utils.createWccData(Nfs3Utils.getWccAttr(preOpDirAttr),
    dfsClient, dirFilePath, iug);
} catch (IOException e1) {
    LOG.error("Can't get postOpDirAttr for dirFileId: "
    + dirHandle.getFileId(), e1);
}
*****The second code snippet*****
//other practice in similar code snippets, line number is 475, the log assigned with INFO level
try {
    wccData = Nfs3Utils.createWccData(Nfs3Utils.getWccAttr(preOpAttr), dfsClient, fileIdPath, iug);
} catch (IOException e1) {
    LOG.info("Can't get postOpAttr for fileIdPath: " + fileIdPath, e1);
}
*****The third code snippet*****
//other practice in similar code snippets, line number is 1405, the log assigned with INFO level
try {
    fromDirWcc = Nfs3Utils.createWccData(
    Nfs3Utils.getWccAttr(fromPreOpAttr), dfsClient, fromDirFilePath, iug);
    toDirWcc = Nfs3Utils.createWccData(Nfs3Utils.getWccAttr(toPreOpAttr),
    dfsClient, toDirFilePath, iug);
}
```

Therefore, I think the logging practices should be consistent in similar contexts. When the code catches *IOException* for *getWccAttr()* method, it more likely prints a log message with **INFO** level, a lower level.

Attachments

HDFS-14340.trunk.patch	0.9 kB	07/Mar/19 13:46
------------------------	--------	-----------------

Activity

star added a comment - 07/Mar/19 01:35

According to the original code, I think:

1. There's a warn log before that exception.
2. The code still work well even with exception thrown out.

```
//代码占位符
} catch (IOException e) {
    LOG.warn("Exception ", e);
    WccData wccData = null;
    try {
```

```
wccData = Nfs3Utils.createWccData(Nfs3Utils.getWccAttr(preOpAttr),
    dfsClient, fileIdPath, iug);
} catch (IOException e1) {
    LOG.info("Can't get postOpAttr for fileIdPath: " + fileIdPath, e1);
}

int status = mapErrorStatus(e);
return new SETATTR3Response(status, wccData);
}
```

Anuhan Torgonshar added a comment - 07/Mar/19 02:18 - edited

Hi, starphin, yes, the code still works well in all of three situations.

In first code snippet, after catching an exception, it still runs following code, which is the same as second and third snippet :

```
int status = mapErrorStatus(e);
return new CREATE3Response(status, fileHandle, postOpObjAttr, dirWcc);
```

Thus, it seems that the first log statement should be assigned to a lower level, INFO, too.

star added a comment - 07/Mar/19 02:58

Checking more code, indeed there are inconsistency in those code——write() with info log level, create with error log info, while others with warn and info log level.

Anuhan Torgonshar added a comment - 07/Mar/19 03:19 - edited

starphin, I agree. In this situation, from the static text in logging statement, we can know that the `getWccAttr()` method triggers the exception, so i think such getting attribute failure can be logged with a lower level.

Thanks for the replies.

star added a comment - 07/Mar/19 03:27

OneisAll, you can make a patch for the issue.

Anuhan Torgonshar added a comment - 07/Mar/19 07:34

Lower the level of log statement which logged an IOException triggered by `getWccAttr()`.

Hadoop QA added a comment - 07/Mar/19 07:40

-1 overall

Vote	Subsystem	Runtime	Comment
0	reexec	0m 0s	Docker mode activated.
-1	patch	0m 6s	HDFS-14340 does not apply to trunk. Rebase required? Wrong Branch? See https://wiki.apache.org/hadoop/HowToContribute for help.

Subsystem	Report/Notes
JIRA Issue	HDFS-14340
JIRA Patch URL	https://issues.apache.org/jira/secure/attachment/12961517/HDFS-14340.patch
Console output	https://builds.apache.org/job/PreCommit-HDFS-Build/26422/console
Powered by	Apache Yetus 0.8.0 http://yetus.apache.org

This message was automatically generated.


Hadoop QA added a comment - 07/Mar/19 09:01

-1 overall

Vote	Subsystem	Runtime	Comment
0	reexec	0m 0s	Docker mode activated.
-1	patch	0m 19s	HDFS-14340 does not apply to master. Rebase required? Wrong Branch? See https://wiki.apache.org/hadoop/HowToContribute for help.

Subsystem	Report/Notes
JIRA Issue	HDFS-14340
JIRA Patch URL	https://issues.apache.org/jira/secure/attachment/12961522/HDFS-14340.master.001.patch
Console output	https://builds.apache.org/job/PreCommit-HDFS-Build/26423/console
Powered by	Apache Yetus 0.8.0 http://yetus.apache.org

This message was automatically generated.

▼  Hadoop QA added a comment - 07/Mar/19 12:00

 -1 overall

Vote	Subsystem	Runtime	Comment
0	reexec	0m 0s	Docker mode activated.
-1	patch	0m 7s	HDFS-14340 does not apply to master. Rebase required? Wrong Branch? See https://wiki.apache.org/hadoop/HowToContribute for help.

Subsystem	Report/Notes
JIRA Issue	HDFS-14340
JIRA Patch URL	https://issues.apache.org/jira/secure/attachment/12961535/HDFS-14340.master.001.patch
Console output	https://builds.apache.org/job/PreCommit-HDFS-Build/26424/console
Powered by	Apache Yetus 0.8.0 http://yetus.apache.org


This message was automatically generated.

▼  star added a comment - 07/Mar/19 13:22

1. switch to branch 'trunk'
2. edit your code and execute cmd 'git diff > ../~~HDFS-14340~~.trunk.patch'

▼  Anuhan Torgonshar added a comment - 07/Mar/19 14:04

starphin,thanks! I uploaded the patch file again.

▼  Hadoop QA added a comment - 07/Mar/19 14:34

 -1 overall

Vote	Subsystem	Runtime	Comment
0	reexec	0m 16s	Docker mode activated.
Prechecks			
+1	@author	0m 0s	The patch does not contain any @author tags.


-1	test4tests	0m 0s	The patch doesn't appear to include any new or modified tests. Please justify why no new tests are needed for this patch. Also please list what manual steps were performed to verify this patch.
trunk Compile Tests			
+1	mvninstall	16m 0s	trunk passed
+1	compile	0m 19s	trunk passed
+1	checkstyle	0m 12s	trunk passed
+1	mvnsite	0m 21s	trunk passed
+1	shadedclient	9m 56s	branch has no errors when building and testing our client artifacts.
+1	findbugs	0m 29s	trunk passed
+1	javadoc	0m 14s	trunk passed
Patch Compile Tests			
+1	mvninstall	0m 18s	the patch passed
+1	compile	0m 16s	the patch passed
+1	javac	0m 16s	the patch passed
+1	checkstyle	0m 9s	the patch passed
+1	mvnsite	0m 17s	the patch passed
+1	whitespace	0m 0s	The patch has no whitespace issues.
+1	shadedclient	10m 43s	patch has no errors when building and testing our client artifacts.
+1	findbugs	0m 33s	the patch passed
+1	javadoc	0m 13s	the patch passed
Other Tests			
+1	unit	2m 5s	hadoop-hdfs-nfs in the patch passed.
+1	asflicense	0m 20s	The patch does not generate ASF License warnings.
		43m 27s	

Subsystem	Report/Notes
Docker	Client=17.05.0-ce Server=17.05.0-ce Image:yetus/hadoop:8f97d6f
JIRA Issue	HDFS-14340
JIRA Patch URL	https://issues.apache.org/jira/secure/attachment/12961555/HDFS-14340.trunk.patch
Optional Tests	dupname asflicense compile javac javadoc mvninstall mvnsite unit shadedclient findbugs checkstyle
uname	Linux 95b7e84bf391 4.4.0-139-generic #165-Ubuntu SMP Wed Oct 24 10:58:50 UTC 2018 x86_64 x86_64 x86_64 GNU/Linux
Build tool	maven
Personality	/testptch/patchprocess/precommit/personality/provided.sh
git revision	trunk / 475011b
maven	version: Apache Maven 3.3.9
Default Java	1.8.0_191
findbugs	v3.1.0-RC1
Test Results	https://builds.apache.org/job/PreCommit-HDFS-Build/26425/testReport/
Max. process+thread count	723 (vs. ulimit of 10000)
modules	C: hadoop-hdfs-project/hadoop-hdfs-nfs U: hadoop-hdfs-project/hadoop-hdfs-nfs
Console output	https://builds.apache.org/job/PreCommit-HDFS-Build/26425/console
Powered by	Apache Yetus 0.8.0 http://yetus.apache.org


This message was automatically generated.

▼  Anuhan Torgonshar added a comment - 08/Mar/19 01:26

The modification fix a log level in catch clause, which is can not verified in test code directly.

▼  Wei-Chiu Chuang added a comment - 18/Jun/19 05:46

+1 Thanks starphin for the review. Committing it.

▼  Wei-Chiu Chuang added a comment - 18/Jun/19 05:47


Pushed to trunk, branch-3.2 branch-3.1. Thanks

▼  Hudson added a comment - 18/Jun/19 05:47


FAILURE: Integrated in Jenkins build Hadoop-trunk-Commit #16770 (See <https://builds.apache.org/job/Hadoop-trunk-Commit/16770/>)

[HDFS-14340](#). Lower the log level when can't get postOpAttr. Contributed (weichiu: rev 098c325a78dce2697ec94a22a4a3981fa5a3849e)

- (edit) hadoop-hdfs-project/hadoop-hdfs-nfs/src/main/java/org/apache/hadoop/hdfs/nfs/nfs3/RpcProgramNfs3.java

▼  Rohith Sharma K S added a comment - 08/Sep/19 09:49 - edited


weichiu OneisAll starphin Could anyone update the release note since it is marked as **Important** in Flags for release 3.2.1? If it nothing breaks backward compatible, better to remove the flag 'Important'.

▼  Rohith Sharma K S added a comment - 09/Sep/19 03:00


Removed Flag which was set as **Important!**

▼ People

Assignee:

 Anuhan Torgonshar

Reporter:

 Anuhan Torgonshar

Votes:

 0 Vote for this issue

Watchers:

 6 Start watching this issue

▼ Dates

Created:

06/Mar/19 12:39

Updated:

02/Oct/19 17:14

Resolved:

18/Jun/19 05:47