Hadoop Map/Reduce / MAPREDUCE-3683

Capacity scheduler LeafQueues maximum capacity calculation issues

Details

Type:	Bug	Status:	CLOSED
Priority:	Blocker	Resolution:	Fixed
Affects Version/s:	0.23.0	Fix Version/s:	0.23.1
Component/s:	mrv2		
Labels:	None		
Target Version/s:	0.23.1		

Description

In the Capacity scheduler if you configure the queues to be hierarchical where you have root -> parent queue -> leaf queue, the leaf queue doesn't take into account its parents maximum capacity when calculate its own maximum capacity, instead it seems to use the parents capacity. Looking at the code its using the parents absoluteCapacity and I think it should be using the parents absoluteMaximumCapacity.

It also seems to only use the parents capacity in the leaf queues max capacity calculation when the leaf queue has a max capacity configured. If the leaf queues maximum-capacity is not configured, then it can use 100% of the cluster.

Attachments

MAPREDUCE-3683.patch	24 kB	24/Jan/12 21:27
MAPREDUCE-3683.patch	13 kB	24/Jan/12 07:04
MAPREDUCE-3683.patch	4 kB	19/Jan/12 08:22

Activity

Thomas Graves created issue - 17/Jan/12 18:57

Arun Murthy made changes - 17/Jan/12 22:24

Field	Original Value	New Value
Assignee		Arun C Murthy [acmurthy]

Arun Murthy added a comment - 19/Jan/12 06:48

> leaf queue doesn't take into account its parents maximum capacity when calculate its own maximum capacity, instead it seems to use the parents capacity - I think it should be using the parents absoluteMaximumCapacity.

Good catch. I'll fix this.

However, if the parent's max-capacity is undefined i.e. UNLIMITED (-1), we have two choices:

1. Use the parent's 'current' capacity to compute the max-capacity of the child at runtime
2. Set the child's max-capacity to be UNLIMITED too.

I'm leaning towards (2) since it's simpler and easier to explain to users i.e. it's more consistent. Otherwise, we run the risk of differing max-capacity (due to parent's capacity changing) which seems unpredictable and hard to explain to users/admins.

Thoughts?

> If the leaf queues maximum-capacity is not configured, then it can use 100% of the cluster.

That isn't true. It just means there is no max-cap for the child and it can use as much capacity as the parent queue gets, which might be capped by the parent/grandparent capacity/max-capacity and so on. This seems like the correct behaviour. Right?

Arun Murthy added a comment - 19/Jan/12 08:22

Straight fwd patch for review.
Note that I had to special case absMaxCap for 'root'.

Still testing, will need to rebase after ~~MAPREDUCE-3681~~.

Arun Murthy made changes - 19/Jan/12 08:22

Attachment	MAPREDUCE-3683.patch [12511106]

Thomas Graves added a comment - 19/Jan/12 19:17

> However, if the parent's max-capacity is undefined i.e. UNLIMITED (-1), we have two choices:
>
> 1. Use the parent's 'current' capacity to compute the max-capacity of the child at runtime
> 2. Set the child's max-capacity to be UNLIMITED too.
>
> I'm leaning towards (2) since it's simpler and easier to explain to users i.e. it's more consistent. Otherwise, we run the risk of differing max-capacity (due to parent's capacity changing) which seems unpredictable and hard to explain to users/admins.
>
> Thoughts?

I guess I expected if the parent wasn't set (its max-capacity defaults to -1), which I thought implies a queue can use complete capacity of the cluster, so if a childs max-capacity was set, then it would just use that childs max-capacity - which would be % of the entire cluster. If you go with option 2 that means that if I set my max-capacity for my child queue it basically does nothing, which isn't what I would expect either.

> That isn't true. It just means there is no max-cap for the child and it can use as much capacity as the parent queue gets, which might be capped by the parent/grandparent capacity/max-capacity and so on. This seems like the correct behaviour. Right?

Ok, sorry I was wrong, it does enforce it, what threw me off here was the web UI shows the absolute maximum capacity as 100% for the leaf queue and I expected this to be parents absolute max capacity. But I think that must be a web UI bug. I will file separate issue for that.

Nathan Roberts added a comment - 19/Jan/12 21:36

My feeling is that an undefined max-capacity should be the same as a max-capacity of 100%. Then it would follow that a child's max-capacity would be the product of its ancestors' max-capacity values. This would be very similar to how absolutecapacity is calculated and I think that consistency would help reduce confusion.

Nathan Roberts added a comment - 23/Jan/12 22:27

Let me clarify that last comment. Is there a specific reason we need to distinguish between a setting of 100% and UNDEFINED? I'm just wondering because if we're not benefiting from the distinction, it's adding confusion:

- Log messages of the form: "absoluteMaxCapacity = 3.4028235E38 [= Float.MAX_VALUE if maximumCapacity undefined, (parentAbsoluteCapacity * maximumCapacity) / 100 otherwise]"
- webservices and webui both report 100% in the case where it's UNDEFINED
- Code like the following in a few places:

```
float absoluteMaxCapacity =
  (Math.round(maximumCapacity * 100) ==
    CapacitySchedulerConfiguration.UNDEFINED) ?
   Float.MAX_VALUE :
   (parent.getAbsoluteMaximumCapacity() * maximumCapacity);
```

Dheeren Beborrtha added a comment - 23/Jan/12 23:34

If capacity=maximumCapacity for one of the container queues, is this capapcity up for grab/stealing if other queues need it and this queue is unutilized?

Arun Murthy added a comment - 24/Jan/12 07:04

Updated with tests.

Arun Murthy made changes - 24/Jan/12 07:04

Attachment	MAPREDUCE-3683.patch [12511641]

Arun Murthy added a comment - 24/Jan/12 07:10

nroberts 100% is all of parent's capacitY, UNDEFINED is effectively unlimited (infinity i.e. Float.MAX_VALUE).

Arun Murthy added a comment - 24/Jan/12 07:11

dheeren Yes, that's always the case. maxCap for a queue just prevents the queue itself.

Arun Murthy made changes - 24/Jan/12 07:11

Status	Open [1]	Patch Available [10002]

Hadoop QA added a comment - 24/Jan/12 07:43

-1 overall. Here are the results of testing the latest attachment
http://issues.apache.org/jira/secure/attachment/12511641/MAPREDUCE-3683.patch
against trunk revision .

+1 @author. The patch does not contain any @author tags.

+1 tests included. The patch appears to include 3 new or modified tests.

+1 javadoc. The javadoc tool did not generate any warning messages.

+1 javac. The applied patch does not increase the total number of javac compiler warnings.

+1 eclipse:eclipse. The patch built with eclipse:eclipse.

+1 findbugs. The patch does not introduce any new Findbugs (version 1.3.9) warnings.

+1 release audit. The applied patch does not increase the total number of release audit warnings.

-1 core tests. The patch failed these unit tests:
org.apache.hadoop.yarn.server.resourcemanager.webapp.TestRMWebServicesCapacitySched

+1 contrib tests. The patch passed contrib unit tests.

Test results: https://builds.apache.org/job/PreCommit-MAPREDUCE-Build/1658//testReport/
Console output: https://builds.apache.org/job/PreCommit-MAPREDUCE-Build/1658//console

This message is automatically generated.

Nathan Roberts added a comment - 24/Jan/12 15:14

Sorry if I'm just being dense here..

> 100% is all of parent's capacitY, UNDEFINED is effectively unlimited (infinity i.e. Float.MAX_VALUE).

I don't see how that lines up with previous comments:

> Tom: If the leaf queues maximum-capacity is not configured, then it can use 100% of the cluster.
>
> Arun: That isn't true. It just means there is no max-cap for the child and it can use as much capacity as the parent queue gets, which might be capped by the parent/grandparent capacity/max-capacity and so on. This seems like the correct behaviour. Right?

I think it works the way this last exchange describes it. Given the following queue configuration (cap,maxcap):
a (25,25)
b (75,75)
a.a1 (50,UNDEFINED)

a.a1 gets 25% of the cluster and nothing more.

I believe it works this way because the resources in the 'a' subtree are exhausted when it's trying to assign containers to that subtree(ParentQueue.assignToQueue).

I actually think it's correct that a leaf queue should not be able to break out of the limits set by its ancestors. However, if this is the case then UNDEFINED is the same as 100%.

Arun Murthy added a comment - 24/Jan/12 16:49

nroberts the interesting case was:
a (25, UNDEFINED)
a.a1 (50, UNDEFINED) v/s a.a1 (50, 100)

In which case, given the previous definition of max-cap (i.e. prior to this patch) this worked very differently... since max-cap was based on parent cap rather than parent max-cap.

However, given the current change, that distinction might not be necessary. We can fix that as a separate jira to unblock this bug, makes sense?

Nathan Roberts added a comment - 24/Jan/12 17:30

I'm fine with doing things in a separate jira, but doesn't this whole issue get much simpler if we can get rid of of MAX_FLOAT?

For example, my understanding of the current patch is that a configuration like:
a (25,UNDEFINED)
a.a1 (50, 50)

will let a job submitted to a.a1 consume the entire grid. This doesn't seem intuitive to me in that we ignored the maxcap setting of a.a1. Since the parent is UNDEFINED, computeAbsoluteMaxCapacity will return Float.MAX_VALUE. Am I missing something there?

Arun Murthy made changes - 24/Jan/12 21:26

Status	Patch Available [10002]	Open [1]

Arun Murthy added a comment - 24/Jan/12 21:27

Thanks for the f/b Nathan. I've incorporated your comments.

Arun Murthy made changes - 24/Jan/12 21:27

Attachment		MAPREDUCE-3683.patch [12511729]

Arun Murthy made changes - 24/Jan/12 21:27

Status	Open [1]	Patch Available [10002]

Hadoop QA added a comment - 24/Jan/12 22:46

+1 overall. Here are the results of testing the latest attachment
http://issues.apache.org/jira/secure/attachment/12511729/MAPREDUCE-3683.patch
against trunk revision .

+1 @author. The patch does not contain any @author tags.

+1 tests included. The patch appears to include 9 new or modified tests.

+1 javadoc. The javadoc tool did not generate any warning messages.

+1 javac. The applied patch does not increase the total number of javac compiler warnings.

+1 eclipse:eclipse. The patch built with eclipse:eclipse.

+1 findbugs. The patch does not introduce any new Findbugs (version 1.3.9) warnings.

+1 release audit. The applied patch does not increase the total number of release audit warnings.

+1 core tests. The patch passed unit tests in .

+1 contrib tests. The patch passed contrib unit tests.

Test results: https://builds.apache.org/job/PreCommit-MAPREDUCE-Build/1662//testReport/
Console output: https://builds.apache.org/job/PreCommit-MAPREDUCE-Build/1662//console

This message is automatically generated.

Nathan Roberts added a comment - 25/Jan/12 15:03

+1 on the patch

One question though, what should the default for capacity be? Right now I think it's UNDEFINED, which will throw an exception. If you decide to tweak that, the debug message in getCapacity() says "setCapacity", probably just remove it.

Arun Murthy added a comment - 25/Jan/12 18:19

Thanks for the review Nathan. I'll fix the logging nit in the commit.

Arun Murthy added a comment - 25/Jan/12 18:19

I just committed this.

Arun Murthy made changes - 25/Jan/12 18:19

Resolution		Fixed [1]
Status	Patch Available [10002]	Resolved [5]

Arun Murthy made changes - 25/Jan/12 18:20

Fix Version/s	0.23.1 [12318883]

Hudson added a comment - 25/Jan/12 18:25

Integrated in Hadoop-Hdfs-0.23-Commit #408 (See https://builds.apache.org/job/Hadoop-Hdfs-0.23-Commit/408/)
Merge -c 1235858 from trunk to branch-0.23 to fix ~~MAPREDUCE-3683~~. Fixed maxCapacity of queues to be product of parent maxCapacities.

acmurthy : http://svn.apache.org/viewcvs.cgi/?root=Apache-SVN&view=rev&rev=1235860
Files :

- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/CHANGES.txt
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/CSQueueUtils.java
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/CapacitySchedulerConfiguration.java
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/LeafQueue.java
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/ParentQueue.java
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/test/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/TestLeafQueue.java
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/test/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/TestQueueParsing.java
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/test/java/org/apache/hadoop/yarn/server/resourcemanager/webapp/TestRMWebServicesCapacitySched.java

---

Hudson added a comment - 25/Jan/12 18:26

Integrated in Hadoop-Common-trunk-Commit #1583 (See https://builds.apache.org/job/Hadoop-Common-trunk-Commit/1583/)
~~MAPREDUCE-3683~~. Fixed maxCapacity of queues to be product of parent maxCapacities.

acmurthy : http://svn.apache.org/viewcvs.cgi/?root=Apache-SVN&view=rev&rev=1235858
Files :

- /hadoop/common/trunk/hadoop-mapreduce-project/CHANGES.txt
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/CSQueueUtils.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/CapacitySchedulerConfiguration.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/LeafQueue.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/ParentQueue.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/test/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/TestLeafQueue.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/test/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/TestQueueParsing.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/test/java/org/apache/hadoop/yarn/server/resourcemanager/webapp/TestRMWebServicesCapacitySched.java

---

Hudson added a comment - 25/Jan/12 18:30

Integrated in Hadoop-Common-0.23-Commit #417 (See https://builds.apache.org/job/Hadoop-Common-0.23-Commit/417/)
Merge -c 1235858 from trunk to branch-0.23 to fix ~~MAPREDUCE-3683~~. Fixed maxCapacity of queues to be product of parent maxCapacities.

acmurthy : http://svn.apache.org/viewcvs.cgi/?root=Apache-SVN&view=rev&rev=1235860
Files :

- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/CHANGES.txt
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/CSQueueUtils.java
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/CapacitySchedulerConfiguration.java
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/LeafQueue.java

- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/ParentQueue.java
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/test/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/TestLeafQueue.java
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/test/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/TestQueueParsing.java
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/test/java/org/apache/hadoop/yarn/server/resourcemanager/webapp/TestRMWebServicesCapacitySched.java

---

Hudson added a comment - 25/Jan/12 18:31

Integrated in Hadoop-Hdfs-trunk-Commit #1656 (See https://builds.apache.org/job/Hadoop-Hdfs-trunk-Commit/1656/)
~~MAPREDUCE-3683~~. Fixed maxCapacity of queues to be product of parent maxCapacities.

acmurthy : http://svn.apache.org/viewcvs.cgi/?root=Apache-SVN&view=rev&rev=1235858
Files :

- /hadoop/common/trunk/hadoop-mapreduce-project/CHANGES.txt
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/CSQueueUtils.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/CapacitySchedulerConfiguration.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/LeafQueue.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/ParentQueue.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/test/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/TestLeafQueue.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/test/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/TestQueueParsing.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/test/java/org/apache/hadoop/yarn/server/resourcemanager/webapp/TestRMWebServicesCapacitySched.java

---

Hudson added a comment - 25/Jan/12 18:38

Integrated in Hadoop-Mapreduce-trunk-Commit #1600 (See https://builds.apache.org/job/Hadoop-Mapreduce-trunk-Commit/1600/)
~~MAPREDUCE-3683~~. Fixed maxCapacity of queues to be product of parent maxCapacities.

acmurthy : http://svn.apache.org/viewcvs.cgi/?root=Apache-SVN&view=rev&rev=1235858
Files :

- /hadoop/common/trunk/hadoop-mapreduce-project/CHANGES.txt
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/CSQueueUtils.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/CapacitySchedulerConfiguration.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/LeafQueue.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/ParentQueue.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/test/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/TestLeafQueue.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/test/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/TestQueueParsing.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/test/java/org/apache/hadoop/yarn/server/resourcemanager/webapp/TestRMWebServicesCapacitySched.java

---

Hudson added a comment - 25/Jan/12 18:38

Integrated in Hadoop-Mapreduce-0.23-Commit #433 (See https://builds.apache.org/job/Hadoop-Mapreduce-0.23-Commit/433/)
Merge -c 1235858 from trunk to branch-0.23 to fix ~~MAPREDUCE-3683~~. Fixed maxCapacity of queues to be product of parent maxCapacities.

acmurthy : http://svn.apache.org/viewcvs.cgi/?root=Apache-SVN&view=rev&rev=1235860
Files :

- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/CHANGES.txt
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/CSQueueUtils.java

- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/CapacitySchedulerConfiguration.java
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/LeafQueue.java
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/ParentQueue.java
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/test/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/TestLeafQueue.java
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/test/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/TestQueueParsing.java
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/test/java/org/apache/hadoop/yarn/server/resourcemanager/webapp/TestRMWebServicesCapacitySched.java

---

Hudson added a comment - 26/Jan/12 12:42

Integrated in Hadoop-Hdfs-trunk #937 (See https://builds.apache.org/job/Hadoop-Hdfs-trunk/937/)
~~MAPREDUCE-3683~~. Fixed maxCapacity of queues to be product of parent maxCapacities.

acmurthy : http://svn.apache.org/viewcvs.cgi/?root=Apache-SVN&view=rev&rev=1235858
Files :

- /hadoop/common/trunk/hadoop-mapreduce-project/CHANGES.txt
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/CSQueueUtils.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/CapacitySchedulerConfiguration.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/LeafQueue.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/ParentQueue.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/test/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/TestLeafQueue.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/test/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/TestQueueParsing.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/test/java/org/apache/hadoop/yarn/server/resourcemanager/webapp/TestRMWebServicesCapacitySched.java

---

Hudson added a comment - 26/Jan/12 12:48

Integrated in Hadoop-Hdfs-0.23-Build #150 (See https://builds.apache.org/job/Hadoop-Hdfs-0.23-Build/150/)
Merge -c 1235858 from trunk to branch-0.23 to fix ~~MAPREDUCE-3683~~. Fixed maxCapacity of queues to be product of parent maxCapacities.

acmurthy : http://svn.apache.org/viewcvs.cgi/?root=Apache-SVN&view=rev&rev=1235860
Files :

- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/CHANGES.txt
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/CSQueueUtils.java
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/CapacitySchedulerConfiguration.java
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/LeafQueue.java
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/ParentQueue.java
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/test/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/TestLeafQueue.java
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/test/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/TestQueueParsing.java
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/test/java/org/apache/hadoop/yarn/server/resourcemanager/webapp/TestRMWebServicesCapacitySched.java

---

Hudson added a comment - 26/Jan/12 13:20

Integrated in Hadoop-Mapreduce-0.23-Build #172 (See https://builds.apache.org/job/Hadoop-Mapreduce-0.23-Build/172/)
Merge -c 1235858 from trunk to branch-0.23 to fix ~~MAPREDUCE-3683~~. Fixed maxCapacity of queues to be product of parent maxCapacities.

acmurthy : http://svn.apache.org/viewcvs.cgi/?root=Apache-SVN&view=rev&rev=1235860
Files :

- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/CHANGES.txt
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/CSQueueUtils.java
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/CapacitySchedulerConfiguration.java
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/LeafQueue.java
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/ParentQueue.java
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/test/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/TestLeafQueue.java
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/test/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/TestQueueParsing.java
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/test/java/org/apache/hadoop/yarn/server/resourcemanager/webapp/TestRMWebServicesCapacitySched.java

---

Hudson added a comment - 26/Jan/12 13:23

Integrated in Hadoop-Mapreduce-trunk #970 (See https://builds.apache.org/job/Hadoop-Mapreduce-trunk/970/)
~~MAPREDUCE-3683~~. Fixed maxCapacity of queues to be product of parent maxCapacities.

acmurthy : http://svn.apache.org/viewcvs.cgi/?root=Apache-SVN&view=rev&rev=1235858
Files :

- /hadoop/common/trunk/hadoop-mapreduce-project/CHANGES.txt
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/CSQueueUtils.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/CapacitySchedulerConfiguration.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/LeafQueue.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/main/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/ParentQueue.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/test/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/TestLeafQueue.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/test/java/org/apache/hadoop/yarn/server/resourcemanager/scheduler/capacity/TestQueueParsing.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-resourcemanager/src/test/java/org/apache/hadoop/yarn/server/resourcemanager/webapp/TestRMWebServicesCapacitySched.java

---

Arun Murthy made changes - 05/Mar/12 02:48

Status	Resolved [5]	Closed [6]

Transition	Time In Source Status	Execution Times
Arun Murthy made transition - 24/Jan/12 21:26 PATCH AVAILABLE → OPEN	14h 14m	1
Arun Murthy made transition - 24/Jan/12 21:27 OPEN → PATCH AVAILABLE	6d 12h 15m	2
Arun Murthy made transition - 25/Jan/12 18:19 PATCH AVAILABLE → RESOLVED	20h 52m	1
Arun Murthy made transition - 05/Mar/12 02:48 RESOLVED → CLOSED	39d 8h 29m	1

People

Assignee:

Arun Murthy

Reporter:

Thomas Graves

Votes:

0 Vote for this issue

Watchers:

6 Start watching this issue

Dates

Created:

17/Jan/12 18:57

Updated:

05/Mar/12 02:48

Resolved:

25/Jan/12 18:19