

Public signup for this instance is **disabled**. Go to our Self serve sign up page to request an account.



Hadoop Map/Reduce / MAPREDUCE-4423

Potential infinite fetching of map output

Details

Type: □ Bug

Priority: ♠ Critical

0.23.3, 2.0.0-alpha, 3.0.0-alpha1

Component/s: mrv2
Labels: None

Target Version/s: 0.23.3

Status: CLOSED
Resolution: Fixed

Fix Version/s:

0.23.3, 2.0.2-alpha

Description

Affects Version/s:

Inside Fetcher.java there are a few cases where an error can happen and the corresponding map task is not marked as a fetch failure. One of these is if the Shuffle server returns a malformed result.

MAPREDUCE-3992 makes this case a lot less common, but it is still possible. IF the shuffle handler always returns a malformed result, but a OK response the Fetcher will never stop trying to fetch those results.

Attachments

■ MR-4423.txt	12 kB	25/Jul/12 21:06
■ MR-4423.txt	11 kB	19/Jul/12 17:55
■ MR-4423.txt	11 kB	17/Jul/12 22:01
■ MR-4423-branch-0.23.txt	15 kB	26/Jul/12 21:43
MR-4423-trunk.txt	15 kB	26/Jul/12 21:43
MR-4423-trunk.txt	12 kB	26/Jul/12 15:42

Activity

▼ ○ Robert Joseph Evans added a comment - 17/Jul/12 22:01

This patch adds in a test, and marks tasks failed on any error while trying to fetch the data.

- Madoop QA added a comment 17/Jul/12 22:24
 - +1 overall. Here are the results of testing the latest attachment http://issues.apache.org/jira/secure/attachment/12536897/MR-4423.txt against trunk revision .
 - +1 @author. The patch does not contain any @author tags.
 - +1 tests included. The patch appears to include 1 new or modified test files.
 - +1 javac. The applied patch does not increase the total number of javac compiler warnings.
 - +1 javadoc. The javadoc tool did not generate any warning messages.
 - +1 eclipse:eclipse. The patch built with eclipse:eclipse.
 - +1 findbugs. The patch does not introduce any new Findbugs (version 1.3.9) warnings.
 - +1 release audit. The applied patch does not increase the total number of release audit warnings.
 - +1 core tests. The patch passed unit tests in hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-core.
 - +1 contrib tests. The patch passed contrib unit tests.

https://issues.apache.org/jira/browse/MAPREDUCE-4423

Test results: https://builds.apache.org/job/PreCommit-MAPREDUCE-Build/2616//testReport/Console output: https://builds.apache.org/job/PreCommit-MAPREDUCE-Build/2616//console

This message is automatically generated.

✓ O Jason Darrell Lowe added a comment - 18/Jul/12 21:17

Looks good overall. One small nit in that the good boolean was removed, but a comment remains that refers to it.

Robert Joseph Evans added a comment - 19/Jul/12 17:55

This patch updates the comment that mentioned the deleted good variable.

- O Hadoop QA added a comment 19/Jul/12 18:15
 - +1 overall. Here are the results of testing the latest attachment http://issues.apache.org/jira/secure/attachment/12537218/MR-4423.txt against trunk revision .
 - +1 @author. The patch does not contain any @author tags.
 - +1 tests included. The patch appears to include 1 new or modified test files.
 - +1 javac. The applied patch does not increase the total number of javac compiler warnings.
 - +1 javadoc. The javadoc tool did not generate any warning messages.
 - +1 eclipse:eclipse. The patch built with eclipse:eclipse.
 - +1 findbugs. The patch does not introduce any new Findbugs (version 1.3.9) warnings.
 - +1 release audit. The applied patch does not increase the total number of release audit warnings.
 - +1 core tests. The patch passed unit tests in hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client/
 - +1 contrib tests. The patch passed contrib unit tests.

Test results: https://builds.apache.org/job/PreCommit-MAPREDUCE-Build/2627//testReport/Console output: https://builds.apache.org/job/PreCommit-MAPREDUCE-Build/2627//console

This message is automatically generated.

- ➤ O Thomas Graves added a comment 19/Jul/12 18:19
 - +1 looks good. Thanks Bobby! I'll commit this shortly.
- ▼ Hudson added a comment 19/Jul/12 18:25

Integrated in Hadoop-Common-trunk-Commit #2500 (See https://builds.apache.org/job/Hadoop-Common-trunk-Commit/2500/)

MAPREDUCE-4423. Potential infinite fetching of map output (Robert Evans via tgraves) (Revision 1363454)

Result = SUCCESS

tgraves: http://svn.apache.org/viewcvs.cgi/?root=Apache-SVN&view=rev&rev=1363454

Files :

- /hadoop/common/trunk/hadoop-mapreduce-project/CHANGES.txt
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-clientcore/src/main/java/org/apache/hadoop/mapreduce/task/reduce/Fetcher.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-core/src/test/java/org/apache/hadoop/mapreduce/task/reduce/TestFetcher.java
- O Hudson added a comment 19/Jul/12 18:27

Integrated in Hadoop-Hdfs-trunk-Commit #2565 (See https://builds.apache.org/job/Hadoop-Hdfs-trunk-Commit/2565/)
MAPREDUCE-4423. Potential infinite fetching of map output (Robert Evans via tgraves) (Revision 1363454)

Result = SUCCESS

tgraves: http://svn.apache.org/viewcvs.cgi/?root=Apache-SVN&view=rev&rev=1363454 Files:

/hadoop/common/trunk/hadoop-mapreduce-project/CHANGES.txt

- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-core/src/main/java/org/apache/hadoop/mapreduce/task/reduce/Fetcher.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-core/src/test/java/org/apache/hadoop/mapreduce/task/reduce/TestFetcher.java
- O Hudson added a comment 19/Jul/12 19:14

Integrated in Hadoop-Mapreduce-trunk-Commit #2522 (See https://builds.apache.org/job/Hadoop-Mapreduce-trunk-Commit/2522/)

MAPREDUCE 4423. Potential infinite fetching of map output (Robert Evans via tgraves) (Revision 1363454)

Result = FAILURE

tgraves: http://svn.apache.org/viewcvs.cgi/?root=Apache-SVN&view=rev&rev=1363454

Files:

- /hadoop/common/trunk/hadoop-mapreduce-project/CHANGES.txt
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-clientcore/src/main/java/org/apache/hadoop/mapreduce/task/reduce/Fetcher.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-core/src/test/java/org/apache/hadoop/mapreduce/task/reduce/TestFetcher.java
- ▼ Hudson added a comment 20/Jul/12 12:40

Integrated in Hadoop-Hdfs-0.23-Build #319 (See https://builds.apache.org/job/Hadoop-Hdfs-0.23-Build/319/) merge -r 1363454:1363455 from branch-2. FIXES: MAPREDUCE-4423 (Revision 1363456)

Result = SUCCESS

tgraves: http://svn.apache.org/viewcvs.cgi/?root=Apache-SVN&view=rev&rev=1363456 Files:

- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/CHANGES.txt
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-core/src/main/java/org/apache/hadoop/mapreduce/task/reduce/Fetcher.java
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-core/src/test/java/org/apache/hadoop/mapreduce/task/reduce/TestFetcher.java
- ✓ O Hudson added a comment 20/Jul/12 14:05

Integrated in Hadoop-Mapreduce-trunk #1142 (See https://builds.apache.org/job/Hadoop-Mapreduce-trunk/1142/) MAPREDUCE 4423. Potential infinite fetching of map output (Robert Evans via tgraves) (Revision 1363454)

Result = FAILURE

tgraves: http://svn.apache.org/viewcvs.cgi/?root=Apache-SVN&view=rev&rev=1363454

Files:

- /hadoop/common/trunk/hadoop-mapreduce-project/CHANGES.txt
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-core/src/main/java/org/apache/hadoop/mapreduce/task/reduce/Fetcher.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-core/src/test/java/org/apache/hadoop/mapreduce/task/reduce/TestFetcher.java
- Thomas Graves added a comment 20/Jul/12 20:16

I'm going to revert this because this seems to have uncovered another bug where map output comes back as OK but has internal error in the body. We need to investigate that more.

▼ ○ Hudson added a comment - 20/Jul/12 20:27

Integrated in Hadoop-Hdfs-trunk-Commit #2573 (See https://builds.apache.org/job/Hadoop-Hdfs-trunk-Commit/2573/) svn merge --change -1363454 for reverting MAPREDUCE-4423 (Revision 1363935)

Result = SUCCESS

tgraves: http://svn.apache.org/viewcvs.cgi/?root=Apache-SVN&view=rev&rev=1363935

Files:

- /hadoop/common/trunk/hadoop-mapreduce-project/CHANGES.txt
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-clientcore/src/main/java/org/apache/hadoop/mapreduce/task/reduce/Fetcher.java

/hadoop/common/trunk/hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-core/src/test/java/org/apache/hadoop/mapreduce/task/reduce/TestFetcher.java

Hudson added a comment - 20/Jul/12 20:34

Integrated in Hadoop-Common-trunk-Commit #2508 (See https://builds.apache.org/job/Hadoop-Common-trunk-Commit/2508/) svn merge --change -1363454 for reverting MAPREDUCE-4423 (Revision 1363935)

Result = SUCCESS

tgraves: http://svn.apache.org/viewcvs.cgi/?root=Apache-SVN&view=rev&rev=1363935

Files:

- /hadoop/common/trunk/hadoop-mapreduce-project/CHANGES.txt
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-core/src/main/java/org/apache/hadoop/mapreduce/task/reduce/Fetcher.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-core/src/test/java/org/apache/hadoop/mapreduce/task/reduce/TestFetcher.java

✓ O Hudson added a comment - 20/Jul/12 21:43

Integrated in Hadoop-Mapreduce-trunk-Commit #2529 (See https://builds.apache.org/job/Hadoop-Mapreduce-trunk-Commit/2529/)

svn merge --change -1363454 for reverting MAPREDUCE-4423 (Revision 1363935)

Result = FAILURE

tgraves: http://svn.apache.org/viewcvs.cgi/?root=Apache-SVN&view=rev&rev=1363935

Files:

- /hadoop/common/trunk/hadoop-mapreduce-project/CHANGES.txt
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-core/src/main/java/org/apache/hadoop/mapreduce/task/reduce/Fetcher.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-core/src/test/java/org/apache/hadoop/mapreduce/task/reduce/TestFetcher.java

O Hudson added a comment - 21/Jul/12 12:39

Integrated in Hadoop-Hdfs-0.23-Build #320 (See https://builds.apache.org/job/Hadoop-Hdfs-0.23-Build/320/) svn merge --change -1363456 for reverting MAPREDUCE 4423 (Revision 1363937)

Result = UNSTABLE

tgraves: http://svn.apache.org/viewcvs.cgi/?root=Apache-SVN&view=rev&rev=1363937

Files:

- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/CHANGES.txt
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-core/src/main/java/org/apache/hadoop/mapreduce/task/reduce/Fetcher.java
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-core/src/test/java/org/apache/hadoop/mapreduce/task/reduce/TestFetcher.java

▼ ○ Hudson added a comment - 21/Jul/12 12:48

Integrated in Hadoop-Hdfs-trunk #1111 (See https://builds.apache.org/job/Hadoop-Hdfs-trunk/1111/) svn merge --change -1363454 for reverting MAPREDUCE 4423 (Revision 1363935)

Result = FAILURE

tgraves: http://svn.apache.org/viewcvs.cgi/?root=Apache-SVN&view=rev&rev=1363935

Files:

- /hadoop/common/trunk/hadoop-mapreduce-project/CHANGES.txt
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-clientcore/src/main/java/org/apache/hadoop/mapreduce/task/reduce/Fetcher.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-core/src/test/java/org/apache/hadoop/mapreduce/task/reduce/TestFetcher.java

Hudson added a comment - 21/Jul/12 14:03

Integrated in Hadoop-Mapreduce-trunk #1143 (See https://builds.apache.org/job/Hadoop-Mapreduce-trunk/1143/) svn merge --change -1363454 for reverting MAPREDUCE-4423 (Revision 1363935)

Result = FAILURE

tgraves: http://svn.apache.org/viewcvs.cgi/?root=Apache-SVN&view=rev&rev=1363935

Files ·

- /hadoop/common/trunk/hadoop-mapreduce-project/CHANGES.txt
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-core/src/main/java/org/apache/hadoop/mapreduce/task/reduce/Fetcher.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-core/src/test/java/org/apache/hadoop/mapreduce/task/reduce/TestFetcher.java
- Robert Joseph Evans added a comment 25/Jul/12 21:06

I determined what the problem was. It was not that this patch is flawed, but it was exposing another transient issue that was committed recently MAPREDUCE 4384.

I did find a few minor issue with the previous patch that I have fixed in this one.

- O Hadoop QA added a comment 25/Jul/12 21:24
 - +1 overall. Here are the results of testing the latest attachment http://issues.apache.org/jira/secure/attachment/12537898/MR-4423.txt against trunk revision .
 - +1 @author. The patch does not contain any @author tags.
 - +1 tests included. The patch appears to include 1 new or modified test files.
 - +1 javac. The applied patch does not increase the total number of javac compiler warnings.
 - +1 javadoc. The javadoc tool did not generate any warning messages.
 - +1 eclipse:eclipse. The patch built with eclipse:eclipse.
 - +1 findbugs. The patch does not introduce any new Findbugs (version 1.3.9) warnings.
 - +1 release audit. The applied patch does not increase the total number of release audit warnings.
 - +1 core tests. The patch passed unit tests in hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client/
 - +1 contrib tests. The patch passed contrib unit tests.

Test results: https://builds.apache.org/job/PreCommit-MAPREDUCE-Build/2660//testReport/Console output: https://builds.apache.org/job/PreCommit-MAPREDUCE-Build/2660//console

This message is automatically generated.

▼ ○ Thomas Graves added a comment - 26/Jul/12 14:17

doesn't apply cleanly to trunk anymore, can you upmerge.

Robert Joseph Evans added a comment - 26/Jul/12 15:42

This version has been upmerged to the latest version of trunk. Please use the original patch for branch-0.23.

Thomas Graves added a comment - 26/Jul/12 16:17

Still looking, but minor comment, it might be a bit confusing to still log this when you return the empty array when told to wait. I realize the array is empty but still says failed task when it didn't really fail any. Perhaps add a check for empty too.

if(failedTasks != null) {

LOG.warn("copyMapOutput failed for tasks "+Arrays.toString(failedTasks));

Robert Joseph Evans added a comment - 26/Jul/12 16:55

That is a great catch. I'll update the patches. If you have anything else please let me know.

- ✓ Hadoop QA added a comment 26/Jul/12 21:17
 - +1 overall. Here are the results of testing the latest attachment http://issues.apache.org/jira/secure/attachment/12538008/MR-4423-trunk.txt against trunk revision .

- +1 @author. The patch does not contain any @author tags.
- +1 tests included. The patch appears to include 1 new or modified test files.
- +1 javac. The applied patch does not increase the total number of javac compiler warnings.
- +1 javadoc. The javadoc tool did not generate any warning messages.
- +1 eclipse:eclipse. The patch built with eclipse:eclipse.
- +1 findbugs. The patch does not introduce any new Findbugs (version 1.3.9) warnings.
- +1 release audit. The applied patch does not increase the total number of release audit warnings.
- +1 core tests. The patch passed unit tests in hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client/nadoop-nado-client/nadoop-nado-client/nadoop-nado-client/nadoop-nado-client/nadoop-nado-client/nadoop-n
- +1 contrib tests. The patch passed contrib unit tests.

Test results: https://builds.apache.org/job/PreCommit-MAPREDUCE-Build/2667//testReport/Console output: https://builds.apache.org/job/PreCommit-MAPREDUCE-Build/2667//console

This message is automatically generated.

Robert Joseph Evans added a comment - 26/Jul/12 21:43

Updated the logging to not warn on WAIT, and added in a test specifically for WAIT.

- ▼ Hadoop QA added a comment 27/Jul/12 00:44
 - +1 overall. Here are the results of testing the latest attachment http://issues.apache.org/jira/secure/attachment/12538083/MR-4423-trunk.txt against trunk revision .
 - +1 @author. The patch does not contain any @author tags.
 - +1 tests included. The patch appears to include 1 new or modified test files.
 - +1 javac. The applied patch does not increase the total number of javac compiler warnings.
 - +1 javadoc. The javadoc tool did not generate any warning messages.
 - +1 eclipse:eclipse. The patch built with eclipse:eclipse.
 - +1 findbugs. The patch does not introduce any new Findbugs (version 1.3.9) warnings.
 - +1 release audit. The applied patch does not increase the total number of release audit warnings.
 - +1 core tests. The patch passed unit tests in hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client/nadoop-nado-client/nadoop-nado-client/nadoop-nado-client/nadoop-nado-client/nadoop-nado-client/nadoop-n
 - +1 contrib tests. The patch passed contrib unit tests.

Test results: https://builds.apache.org/job/PreCommit-MAPREDUCE-Build/2668//testReport/Console output: https://builds.apache.org/job/PreCommit-MAPREDUCE-Build/2668//console

This message is automatically generated.

- ▼ Thomas Graves added a comment 27/Jul/12 01:48
 - +1 thanks Bobby!
- Hudson added a comment 27/Jul/12 02:15

Integrated in Hadoop-Common-trunk-Commit #2526 (See https://builds.apache.org/job/Hadoop-Common-trunk-Commit/2526/)

MAPREDUCE-4423. Potential infinite fetching of map output (Robert Evans via tgraves) (Revision 1366258)

Result = SUCCESS

tgraves: http://svn.apache.org/viewcvs.cgi/?root=Apache-SVN&view=rev&rev=1366258 Files:

- /hadoop/common/trunk/hadoop-mapreduce-project/CHANGES.txt
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-core/src/main/java/org/apache/hadoop/mapreduce/task/reduce/Fetcher.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-clientcore/src/test/java/org/apache/hadoop/mapreduce/task/reduce/TestFetcher.java

▼ ○ Hudson added a comment - 27/Jul/12 02:19

Integrated in Hadoop-Hdfs-trunk-Commit #2590 (See https://builds.apache.org/job/Hadoop-Hdfs-trunk-Commit/2590/)

MAPREDUCE-4423. Potential infinite fetching of map output (Robert Evans via tgraves) (Revision 1366258)

Result = SUCCESS

tgraves: http://svn.apache.org/viewcvs.cgi/?root=Apache-SVN&view=rev&rev=1366258 Files:

- /hadoop/common/trunk/hadoop-mapreduce-project/CHANGES.txt
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-core/src/main/java/org/apache/hadoop/mapreduce/task/reduce/Fetcher.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-core/src/test/java/org/apache/hadoop/mapreduce/task/reduce/TestFetcher.java
- ✓ O Hudson added a comment 27/Jul/12 02:44

Integrated in Hadoop-Mapreduce-trunk-Commit #2546 (See https://builds.apache.org/job/Hadoop-Mapreduce-trunk-Commit/2546/)

MAPREDUCE 4423. Potential infinite fetching of map output (Robert Evans via tgraves) (Revision 1366258)

Result = FAILURE

tgraves: http://svn.apache.org/viewcvs.cgi/?root=Apache-SVN&view=rev&rev=1366258 Files:

- /hadoop/common/trunk/hadoop-mapreduce-project/CHANGES.txt
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-core/src/main/java/org/apache/hadoop/mapreduce/task/reduce/Fetcher.java
- /hadoop/common/trunk/hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-clientcore/src/test/java/org/apache/hadoop/mapreduce/task/reduce/TestFetcher.java
- ➤ O Hudson added a comment 27/Jul/12 12:41

Integrated in Hadoop-Hdfs-0.23-Build #326 (See https://builds.apache.org/job/Hadoop-Hdfs-0.23-Build/326/) MAPREDUCE-4423. Potential infinite fetching of map output (Robert Evans via tgraves) (Revision 1366263)

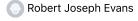
Result = SUCCESS

tgraves: http://svn.apache.org/viewcvs.cgi/?root=Apache-SVN&view=rev&rev=1366263 Files:

- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/CHANGES.txt
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-core/src/main/java/org/apache/hadoop/mapreduce/task/reduce/Fetcher.java
- /hadoop/common/branches/branch-0.23/hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-core/src/test/java/org/apache/hadoop/mapreduce/task/reduce/TestFetcher.java

People

Assignee:



Reporter:

Robert Joseph Evans

Votes:

Vote for this issue

Watchers:

6 Start watching this issue

Dates

Created:

10/Jul/12 15:26

Updated:

12/May/16 18:24

Resolved:

27/Jul/12 01:51