



Provide new timeline plugin storage to support fine-grained entity caching

▼ Details

Type:	Sub-task	Status:	RESOLVED
Priority:	Major	Resolution:	Fixed
Affects Version/s:	None	Fix Version/s:	2.8.0, 3.0.0-alpha1
Component/s:	timelineserver		
Labels:	None		
Hadoop Flags:	Reviewed		

▼ Description

To support the newly proposed APIs in [YARN-4234](#), we need to create a new plugin timeline store. The store may have similar behavior as the EntityFileTimelineStore proposed in [YARN-3942](#), but cache date in cache id granularity, instead of application id granularity. Let's have this storage as a standalone one, instead of updating EntityFileTimelineStore, to keep the existing store (EntityFileTimelineStore) stable.

▼ Attachments

YANR-4265-branch-2.8-addendum.patch	1.0 kB	18/Jan/16 11:00
YARN-4265.YARN-4234.001.patch	104 kB	13/Nov/15 00:01
YARN-4265.YARN-4234.002.patch	105 kB	11/Dec/15 00:45
YARN-4265-branch-2-addendum.patch	1.0 kB	18/Jan/16 11:00
YARN-4265-trunk.001.patch	106 kB	23/Dec/15 21:01
YARN-4265-trunk.002.patch	104 kB	05/Jan/16 00:34
YARN-4265-trunk.003.patch	104 kB	05/Jan/16 21:23
YARN-4265-trunk.004.patch	107 kB	09/Jan/16 01:33
YARN-4265-trunk.005.patch	112 kB	12/Jan/16 01:56
YARN-4265-trunk.006.patch	112 kB	13/Jan/16 22:28
YARN-4265-trunk.007.patch	114 kB	14/Jan/16 02:43
YARN-4265-trunk.008.patch	114 kB	14/Jan/16 17:51

▼ Issue Links

blocks
YARN-4545 Allow YARN distributed shell to use ATS v1.5 APIs RESOLVED
depends upon
YARN-4234 New put APIs in TimelineClient for ats v1.5 RESOLVED
is depended upon by
YARN-4219 New levelDB cache storage for timeline v1.5 RESOLVED
supercedes
YARN-3942 Timeline store to read events from HDFS RESOLVED

▼ Activity

▼ Li Lu added a comment - 29/Oct/15 21:52
We need to use the finalized APIs in YARN-4234 to build the reading code routine for ATS 1.5 storage.
▼ Li Lu added a comment - 29/Oct/15 22:43

This is a proof-of-concept patch for the proposed improvement. This patch is largely based on the work in [YARN-3942](#). In this patch we added a new storage plugin to accommodate the newly added cache id concept. As in [YARN-4234](#), entity logs are stored in <active/done dir name>/app_id/attempt_id/<prefix>-cache_id.log. The newly added plugin can with this new directory structure, and can refresh (reload) entity logs in a cache id granularity (rather than application granularity).

Li Lu added a comment - 29/Oct/15 22:51

Another key addition in this patch, compare to [YARN-3942](#), is a "cache id plugin" mechanism. One general problem the plugin needs to solve is "which exact storage should I look at on a read request like ...". The user can add a plugin (for example, a Tez plugin) in the configuration. At run time, on each read request, the plugin can help deciding if we need to look at the summary storage only or the detailed (entity) storage. In this way, the whole storage decision process of ATS 1.5 fine-grain cache storage will be:

- On writes, the application can decide if an entity should go to the summary storage or the detailed storage. If it should go to the latter, the application need to provide a cache id for this entity.
- On reads, the application provides a plugin to decide if the look up can be performed on the summary storage, or on the detailed storage. If it's the latter, the plugin provides a cache id so that the storage should look at.

Li Lu added a comment - 29/Oct/15 23:10

This patch relies on [YARN-3942.002](#) and [YARN-4234.3](#) patches.

Jason Darrell Lowe added a comment - 11/Nov/15 21:24

Thanks for the patch, [gtCarrera9](#)! This looks like most of the patch is a copy of the entity timeline store from [YARN-3942](#) with a few edits, so I'm sorta reviewing my own code here. As such I did a diff of the patch from this JIRA and the one from [YARN-3942](#) so I could focus on what's changed. I'll defer to others to review the parts that are identical to [YARN-3942](#). Eventually I can see this being a superset of [YARN-3942](#), since it can cache to memory and either cache everything or a subset based on what the plugins decide.

TIMELINE_SERVICE_PLUGIN_ENABLED and DEFAULT_TIMELINE_SERVICE_PLUGIN_ENABLED are not needed.

Is there a reason we're not using the Configuration.getInstance method or the ReflectionUtils methods to handle plugin loading?

If no plugins are configured (which is the default behavior), do we want a fallback plugin that emulates what [YARN-3942](#) is doing?

Are there plans to support the leveldb store as an alternative to the memory store for the detail timeline store? There was concern that a single dag could overwhelm the server, and storing it to leveldb instead of the memory store would be one way to try to mitigate that. I'm wondering if the class to use for the detail log timeline store should be configurable with MemoryTimelineStore as the default. Could also do this as a followup JIRA if necessary.

Should add entries to yarn-default.xml for the new properties?

Do we want to log at the info level that a path is being skipped during the scan? The store can end up scanning fairly often in practice, so this could end up logging a lot for just one path per scan. I'm wondering if making it a debug log is more appropriate.

Comment in getDoneAppPath mentions a cache ID but it's using an app ID.

Nit: Indentation is off at the start of the YarnConfiguration patch hunk.

Li Lu added a comment - 11/Nov/15 21:36

Thanks [jlowe](#)! Yes the current POC version is heavily based on [YARN-3942](#). I'm doing some refactoring right now but the general logic should be pretty similar to your logic (thank for the work!). I'll address your comments in the next version of the patch.

If no plugins are configured (which is the default behavior), do we want a fallback plugin that emulates what [YARN-3942](#) is doing?

That's a nice suggestion. We can build a plugin that simply return the appid to simulate this.

Are there plans to support the leveldb store as an alternative to the memory store for the detail timeline store? There was concern that a single dag could overwhelm the server, and storing it to leveldb instead of the memory store would be one way to try to mitigate that. I'm wondering if the class to use for the detail log timeline store should be configurable with MemoryTimelineStore as the default. Could also do this as a followup JIRA if necessary.

Will do.

Should add entries to yarn-default.xml for the new properties?

Yes. Will add them.

Do we want to log at the info level that a path is being skipped during the scan? The store can end up scanning fairly often in practice, so this could end up logging a lot for just one path per scan. I'm wondering if making it a debug log is more appropriate.

Sure.

Li Lu added a comment - 13/Nov/15 00:01

I refactored my code and added some unit tests. The current patch only depends on [YARN-4234](#). I addressed comments raised by [jlowe](#) in the previous round of review, with two pending actions that I think we can address in separate JIRAs:

1. Build a "fall-back" plugin to have the same behavior of [YARN-3942](#).
2. Make caching storage pluggable.

For the second item, right now my patch does *not* depend on [YARN-4219](#). So I'm setting the v1.5 plugin storage to use the memory storage system as the "caching" storage, similar to [YARN-3942](#).

I'm adding the v1.5 plugin storage (EntityGroupFSTimelineStore) in a module called hadoop-yarn-server-timeline-pluginstorage. This is slightly different to [YARN-3942](#). We need a separate module because the new v1.5 storage depends on yarn-client, but we don't want ats server depends on yarn-client (The v1.5 storage is not a purely server side storage.). I'm naming it as "pluginstorage" because I'm considering to put the leveldb caching storage into this extension as well.

Li Lu added a comment - 11/Dec/15 00:45

Update a new patch to fix some log issues and implement the getTimeline API.

Junping Du added a comment - 23/Dec/15 13:39

I just commit [YARN-4234](#). [gtCarrera9](#), would you rebase your patch on latest trunk branch? Thanks!

Hi [jlowe](#), I saw your comments above: "This looks like most of the patch is a copy of the entity timeline store from [YARN-3942](#) with a few edits, so I'm sorta reviewing my own code here. As such I did a diff of the patch from this JIRA and the one from [YARN-3942](#) so I could focus on what's changed. I'll defer to others to review the parts that are identical to [YARN-3942](#). Eventually I can see this being a superset of [YARN-3942](#), since it can cache to memory and either cache everything or a subset based on what the plugins decide." Are you OK with continue the review effort going with this patch? Or you have some other preferences?

Hadoop QA added a comment - 23/Dec/15 13:48

-1 overall

Vote	Subsystem	Runtime	Comment
-1	patch	0m 5s	YARN-4265 does not apply to trunk. Rebase required? Wrong Branch? See https://wiki.apache.org/hadoop/HowToContribute for help.

Subsystem	Report/Notes
JIRA Patch URL	https://issues.apache.org/jira/secure/attachment/12776931/YARN-4265.YARN-4234.002.patch
JIRA Issue	YARN-4265
Powered by	Apache Yetus 0.2.0-SNAPSHOT http://yetus.apache.org
Console output	https://builds.apache.org/job/PreCommit-YARN-Build/10080/console

This message was automatically generated.

Li Lu added a comment - 23/Dec/15 21:01 - edited

Thanks [junping_du](#)! I just rebased my patch to the latest trunk.

Hadoop QA added a comment - 24/Dec/15 00:51

-1 overall

Vote	Subsystem	Runtime	Comment
0	reexec	0m 0s	Docker mode activated.
+1	@author	0m 0s	The patch does not contain any @author tags.
+1	test4tests	0m 0s	The patch appears to include 8 new or modified test files.
+1	mvninstall	8m 11s	trunk passed
+1	compile	2m 2s	trunk passed with JDK v1.8.0_66
+1	compile	2m 15s	trunk passed with JDK v1.7.0_91
+1	checkstyle	0m 29s	trunk passed
+1	mvnsite	2m 43s	trunk passed
+1	mvneclipse	0m 45s	trunk passed

-1	findbugs	3m 32s	branch/hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server no findbugs output file (hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/target/findbugsXml.xml)
+1	javadoc	2m 13s	trunk passed with JDK v1.8.0_66
+1	javadoc	4m 45s	trunk passed with JDK v1.7.0_91
+1	mvninstall	2m 24s	the patch passed
+1	compile	2m 4s	the patch passed with JDK v1.8.0_66
+1	javac	2m 4s	the patch passed
+1	compile	2m 21s	the patch passed with JDK v1.7.0_91
+1	javac	2m 21s	the patch passed
-1	checkstyle	0m 29s	Patch generated 54 new checkstyle issues in hadoop-yarn-project/hadoop-yarn (total was 292, now 345).
+1	mvnsite	2m 47s	the patch passed
+1	mvnecclipse	0m 45s	the patch passed
+1	whitespace	0m 0s	Patch has no whitespace issues.
+1	xml	0m 2s	The patch has no ill-formed XML file.
-1	findbugs	3m 42s	patch/hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server no findbugs output file (hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/target/findbugsXml.xml)
+1	javadoc	2m 11s	the patch passed with JDK v1.8.0_66
+1	javadoc	4m 46s	the patch passed with JDK v1.7.0_91
+1	unit	0m 26s	hadoop-yarn-api in the patch passed with JDK v1.8.0_66.
+1	unit	2m 5s	hadoop-yarn-common in the patch passed with JDK v1.8.0_66.
-1	unit	79m 12s	hadoop-yarn-server in the patch failed with JDK v1.8.0_66.
+1	unit	0m 26s	hadoop-yarn-api in the patch passed with JDK v1.7.0_91.
+1	unit	2m 19s	hadoop-yarn-common in the patch passed with JDK v1.7.0_91.
-1	unit	81m 40s	hadoop-yarn-server in the patch failed with JDK v1.7.0_91.
-1	asflicense	0m 34s	Patch generated 1 ASF License warnings.
		222m 18s	

Reason	Tests
JDK v1.8.0_66 Failed junit tests	hadoop.yarn.server.resourcemanager.TestClientRMTokens
	hadoop.yarn.server.resourcemanager.TestAMAuthorization
JDK v1.7.0_91 Failed junit tests	hadoop.yarn.server.resourcemanager.TestClientRMTokens
	hadoop.yarn.server.resourcemanager.TestAMAuthorization
	hadoop.yarn.server.resourcemanager.TestRM

Subsystem	Report/Notes
Docker	Image:yetus/hadoop:0ca8df7
JIRA Patch URL	https://issues.apache.org/jira/secure/attachment/12779316/YARN-4265-trunk.001.patch
JIRA Issue	YARN-4265
Optional Tests	asflicense compile javac javadoc mvninstall mvnsite unit findbugs checkstyle xml
uname	Linux 597a5836ed64 3.13.0-36-lowlatency #63-Ubuntu SMP PREEMPT Wed Sep 3 21:56:12 UTC 2014 x86_64 x86_64 x86_64 GNU/Linux
Build tool	maven
Personality	/testptch/hadoop/patchprocess/precommit/personality/provided.sh
git revision	trunk / bb5df27
Default Java	1.7.0_91
Multi-JDK versions	/usr/lib/jvm/java-8-oracle:1.8.0_66 /usr/lib/jvm/java-7-openjdk-amd64:1.7.0_91
findbugs	v3.0.0
checkstyle	https://builds.apache.org/job/PreCommit-YARN-Build/10085/artifact/patchprocess/diff-checkstyle-hadoop-yarn-project_hadoop-

	yarn.txt
unit	https://builds.apache.org/job/PreCommit-YARN-Build/10085/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.8.0_66.txt
unit	https://builds.apache.org/job/PreCommit-YARN-Build/10085/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.7.0_91.txt
unit test logs	https://builds.apache.org/job/PreCommit-YARN-Build/10085/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.8.0_66.txt https://builds.apache.org/job/PreCommit-YARN-Build/10085/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.7.0_91.txt
JDK v1.7.0_91 Test Results	https://builds.apache.org/job/PreCommit-YARN-Build/10085/testReport/
asflicense	https://builds.apache.org/job/PreCommit-YARN-Build/10085/artifact/patchprocess/patch-asflicense-problems.txt
modules	C: hadoop-yarn-project/hadoop-yarn/hadoop-yarn-api hadoop-yarn-project/hadoop-yarn/hadoop-yarn-common hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-applicationhistoryservice hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-timeline-pluginstorage U: hadoop-yarn-project/hadoop-yarn
Max memory used	76MB
Powered by	Apache Yetus 0.2.0-SNAPSHOT http://yetus.apache.org
Console output	https://builds.apache.org/job/PreCommit-YARN-Build/10085/console

This message was automatically generated.

Li Lu added a comment - 30/Dec/15 02:32

BTW, I've got some offline discussion with [jlowe](#). In general, we agreed to move forward with only [YARN-4265](#), which is built on top of [YARN-4234](#). If I've got any misunderstanding, please feel free to let me know, Jason. Thanks!

Li Lu added a comment - 30/Dec/15 02:33

Sorry by saying [YARN-4234](#) I'm actually meaning [YARN-3942](#).

Junping Du added a comment - 04/Jan/16 17:54

Thanks [gtCarrera9](#) for clarification. Assume Jason is fine with going on with this patch, I quickly go through the v3 patch. with following comments so far (I haven't finished my review yet as patch is pretty big):

In YarnConfiguration.java,

```
TIMELINE_SERVICE_ENTITYGROUP_FS_STORE_SCAN_INTERVAL_SECONDS_DEFAULT = 60;
```

I noticed that we are setting 1 minutes as default scan interval but original patch in [HDFS-3942](#) is 5 minutes. Why shall we do any update here? The same question on "app-cache-size", the default value in [HDFS-3942](#) is 5 but here is 10. Any reason to update the value?

In yarn-default.xml,

```
+ <description>DFS path to store active application's timeline data</description>
...
+ <description>DFS path to store done application's timeline data</description>
```

DFS is very old name, use HDFS instead to be more clear.

Why we don't have any default value specified in property of "yarn.timeline-service.entity-group-fs-store.group-id-plugin-classes"?

In hadoop-yarn-server-timeline-pluginstorage/pom.xml,

For EmptyTimelineEntityGroupPlugin.java, why we need this class? I didn't see any help provided even in tests. We should remove it if useless.

In EntityCacheItem.java,

We should have a description for this class in Javadoc.

Can we optimize the synchronization logic here? Like in synchronized method refreshCache, we are initialize/start/stop TimelineDataManager (and MemoryTimelineStore) which is quite expensive and unnecessary to block other synchronized operations. Shall we move these operations out of synchronized block?

```
+ LOG.warn("Error closing datamanager", e);
```

I think we are closing store here instead of datamanager. Isn't it?

```
+ public boolean needRefresh() {
+     //TODO: make a config for cache freshness
+     return (Time.monotonicNow() - lastRefresh > 10000);
+ }
```

Does refresh interval here need to do any coordination with scan interval specified in "yarn.timeline-service.entity-group-fs-store.scan-interval-seconds"?

In EntityGroupFSTimelineStore.java,

```
+     if (appState != AppState.UNKNOWN) {
+         appLogs = new AppLogs(applicationId, appDirPath, appState);
+         LOG.debug("Create and try to add new appLogs to appIdLogMap for {}",
+             applicationId);
+         AppLogs oldAppLogs = appIdLogMap.putIfAbsent(applicationId, appLogs);
+         if (oldAppLogs != null) {
+             appLogs = oldAppLogs;
+         }
+     }
```

This logic is very similiar with method of getAndSetActiveLog(). Can we consolidate them together?

If parseSummaryLogs() is synchronized, it seems getSummaryLogs() should be synchronized too or the getter will get stale(half-done) result.

Still checking if other multi-threads issues. More comments will come soon.

Li Lu added a comment - 05/Jan/16 00:34 - edited

Thanks [junping_du](#) for the review! In the 002 patch I addressed most of the checkstyle problems, as well as most existing comments. Please feel free to add more. Some comments:

I noticed that we are setting 1 minutes as default scan interval but original patch in [HDFS-3942](#) is 5 minutes. Why shall we do any update here?

For now I increased the default frequency to scan HDFS and pull timeline data. Having a 5-minute time interval means users are less likely to see any running status for apps that finish within 5 minutes. Right now I'm setting this value to 1 minute to reduce reader react time.

The same question on "app-cache-size", the default value in [HDFS-3942](#) is 5 but here is 10. Any reason to update the value?

In [YARN-3942](#), caching is performed on application level. In this patch, caching is performed in entity groups. Each application may have a few to tens of entity groups. Normally, there are slightly more active entity groups than active applications in the system. For now, I'm increasing this default value to hold slightly more entity groups in cache.

Why we don't have any default value specified in property of "yarn.timeline-service.entity-group-fs-store.group-id-plugin-classes"?

Plugins are provided by third-party applications such as Tez. Right now we cannot assume which exact entity group plugin the user is using, therefore we have to conservatively leave this config as empty.

For EmptyTimelineEntityGroupPlugin.java, why we need this class? I didn't see any help provided even in tests. We should remove it if useless.

Ah, nice catch. Removed it.

Can we optimize the synchronization logic here? Like in synchronized method refreshCache, we are intialize/start/stop TimelineDataManager (and MemoryTimelineStore) which is quite expensive and unnecessary to block other synchronized operations. Shall we move these operations out of synchronized block?

It's certainly doable. Right now I have yet to optimize this part because it's a little bit tricky to fine tune synchronization performance before we have a relatively stable starting point. Also, we're using fine-grained locking for each cached item in the reader cache, and cache refresh only happens infrequently (~10 secs by default), so maybe we'd like to stabilize the whole synchronization story before fine tune this part?

Hadoop QA added a comment - 05/Jan/16 04:25

✖

-1 overall

Vote	Subsystem	Runtime	Comment
0	reexec	0m 0s	Docker mode activated.
+1	@author	0m 0s	The patch does not contain any @author tags.
+1	test4tests	0m 0s	The patch appears to include 8 new or modified test files.
+1	mvninstall	7m 36s	trunk passed
+1	compile	1m 45s	trunk passed with JDK v1.8.0_66
+1	compile	2m 4s	trunk passed with JDK v1.7.0_91
+1	checkstyle	0m 30s	trunk passed
+1	mvnsite	2m 28s	trunk passed

+1	mvneclipse	0m 44s	trunk passed
-1	findbugs	3m 19s	branch/hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server no findbugs output file (hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/target/findbugsXml.xml)
+1	javadoc	2m 0s	trunk passed with JDK v1.8.0_66
+1	javadoc	4m 26s	trunk passed with JDK v1.7.0_91
+1	mvninstall	2m 9s	the patch passed
+1	compile	1m 48s	the patch passed with JDK v1.8.0_66
+1	javac	1m 48s	the patch passed
+1	compile	2m 10s	the patch passed with JDK v1.7.0_91
+1	javac	2m 10s	the patch passed
-1	checkstyle	0m 31s	Patch generated 9 new checkstyle issues in hadoop-yarn-project/hadoop-yarn (total was 292, now 300).
+1	mvnsite	2m 30s	the patch passed
+1	mvneclipse	0m 38s	the patch passed
+1	whitespace	0m 1s	Patch has no whitespace issues.
+1	xml	0m 1s	The patch has no ill-formed XML file.
-1	findbugs	3m 25s	patch/hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server no findbugs output file (hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/target/findbugsXml.xml)
-1	javadoc	3m 32s	hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.8.0_66 with JDK v1.8.0_66 generated 32 new issues (was 544, now 544).
+1	javadoc	1m 51s	the patch passed with JDK v1.8.0_66
+1	javadoc	4m 23s	the patch passed with JDK v1.7.0_91
+1	unit	0m 19s	hadoop-yarn-api in the patch passed with JDK v1.8.0_66.
+1	unit	1m 56s	hadoop-yarn-common in the patch passed with JDK v1.8.0_66.
-1	unit	79m 35s	hadoop-yarn-server in the patch failed with JDK v1.8.0_66.
+1	unit	0m 23s	hadoop-yarn-api in the patch passed with JDK v1.7.0_91.
+1	unit	2m 10s	hadoop-yarn-common in the patch passed with JDK v1.7.0_91.
-1	unit	80m 43s	hadoop-yarn-server in the patch failed with JDK v1.7.0_91.
+1	asflicense	0m 20s	Patch does not generate ASF License warnings.
		226m 56s	

Reason	Tests
JDK v1.8.0_66 Failed junit tests	hadoop.yarn.server.resourcemanager.TestClientRMTokens
	hadoop.yarn.server.resourcemanager.TestAMAuthorization
JDK v1.7.0_91 Failed junit tests	hadoop.yarn.server.resourcemanager.TestClientRMTokens
	hadoop.yarn.server.resourcemanager.TestAMAuthorization

Subsystem	Report/Notes
Docker	Image:yetus/hadoop:0ca8df7
JIRA Patch URL	https://issues.apache.org/jira/secure/attachment/12780433/YARN-4265-trunk.002.patch
JIRA Issue	YARN-4265
Optional Tests	asflicense compile javac javadoc mvninstall mvnsite unit findbugs checkstyle xml
uname	Linux 14719e8d400e 3.13.0-36-lowlatency #63-Ubuntu SMP PREEMPT Wed Sep 3 21:56:12 UTC 2014 x86_64 x86_64 x86_64 GNU/Linux
Build tool	maven
Personality	/testptch/hadoop/patchprocess/precommit/personality/provided.sh
git revision	trunk / 778146e
Default Java	1.7.0_91
Multi-JDK versions	/usr/lib/jvm/java-8-oracle:1.8.0_66 /usr/lib/jvm/java-7-openjdk-amd64:1.7.0_91
findbugs	v3.0.0

checkstyle	https://builds.apache.org/job/PreCommit-YARN-Build/10149/artifact/patchprocess/diff-checkstyle-hadoop-yarn-project_hadoop-yarn.txt
javadoc	hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.8.0_66: https://builds.apache.org/job/PreCommit-YARN-Build/10149/artifact/patchprocess/diff-javadoc-javadoc-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.8.0_66.txt
unit	https://builds.apache.org/job/PreCommit-YARN-Build/10149/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.8.0_66.txt
unit	https://builds.apache.org/job/PreCommit-YARN-Build/10149/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.7.0_91.txt
unit test logs	https://builds.apache.org/job/PreCommit-YARN-Build/10149/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.8.0_66.txt https://builds.apache.org/job/PreCommit-YARN-Build/10149/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.7.0_91.txt
JDK v1.7.0_91 Test Results	https://builds.apache.org/job/PreCommit-YARN-Build/10149/testReport/
modules	C: hadoop-yarn-project/hadoop-yarn/hadoop-yarn-api hadoop-yarn-project/hadoop-yarn/hadoop-yarn-common hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-applicationhistoryservice hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-timeline-pluginstorage U: hadoop-yarn-project/hadoop-yarn
Max memory used	76MB
Powered by	Apache Yetus 0.2.0-SNAPSHOT http://yetus.apache.org
Console output	https://builds.apache.org/job/PreCommit-YARN-Build/10149/console

This message was automatically generated.

Junping Du added a comment - 05/Jan/16 14:51

Thanks [gtCarrera9](#) for reply and updating the patch.
I am still reviewing the new patch but have some quick response below:

Plugins are provided by third-party applications such as Tez. Right now we cannot assume which exact entity group plugin the user is using, therefore we have to conservatively leave this config as empty.

Thanks for clarification on this. Can we provide a default one (could be simple) for built-in YARN applications, i.e. distributed shell? This can be helpful to demonstrate how to use it for applications other than Tez. Or other apps have to refer Tez code on how to use it. IMO, this is not a very good practice for YARN.

we're using fine-grained locking for each cached item in the reader cache, and cache refresh only happens infrequently (~10 secs by default), so maybe we'd like to stabilize the whole synchronization story before fine tune this part?

Sounds like a plan. Looks like we still need to make effort on checking the whole synchronization mechanism. Isn't it?

Li Lu added a comment - 05/Jan/16 18:04

Can we provide a default one (could be simple) for built-in YARN applications, i.e. distributed shell?

We can certainly do that, given that the new APIs are in and we've got this patch. Since this is a separate work than introducing the whole new storage (as in this JIRA), maybe we can address this as a new JIRA? I can start off doing that right away.

Looks like we still need to make effort on checking the whole synchronization mechanism. Isn't it?

I checked all synchronization logic a few times before I submit the original patch. We're using fine-grained locking in this patch, so yes it's more error pruning. I'll check again later on and fix any problem I discovered. Feel free to raise up any concerns in your review. Thanks!

Li Lu added a comment - 05/Jan/16 21:23

Refresh the patch once more to:

- 1. Add documentations about cache item and its main features
- 2. Double check synchronizations and fix unsynchronized accessors.
- 3. Change accessor modifiers to some methods to make them stricter.

Run findbugs locally and it looks fine.

Hadoop QA added a comment - 06/Jan/16 01:41

-1 overall

Vote	Subsystem	Runtime	Comment
0	reexec	0m 0s	Docker mode activated.
+1	@author	0m 0s	The patch does not contain any @author tags.
+1	test4tests	0m 0s	The patch appears to include 8 new or modified test files.
+1	mvninstall	8m 19s	trunk passed
+1	compile	2m 2s	trunk passed with JDK v1.8.0_66
+1	compile	2m 26s	trunk passed with JDK v1.7.0_91
+1	checkstyle	0m 35s	trunk passed
+1	mvnsite	2m 40s	trunk passed
+1	mvneclipse	0m 48s	trunk passed
-1	findbugs	3m 25s	branch/hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server no findbugs output file (hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/target/findbugsXml.xml)
+1	javadoc	2m 14s	trunk passed with JDK v1.8.0_66
+1	javadoc	4m 50s	trunk passed with JDK v1.7.0_91
+1	mvninstall	2m 29s	the patch passed
+1	compile	2m 6s	the patch passed with JDK v1.8.0_66
+1	javac	2m 6s	the patch passed
+1	compile	2m 25s	the patch passed with JDK v1.7.0_91
+1	javac	2m 25s	the patch passed
-1	checkstyle	0m 33s	Patch generated 9 new checkstyle issues in hadoop-yarn-project/hadoop-yarn (total was 292, now 300).
+1	mvnsite	2m 47s	the patch passed
+1	mvneclipse	0m 42s	the patch passed
+1	whitespace	0m 0s	Patch has no whitespace issues.
+1	xml	0m 2s	The patch has no ill-formed XML file.
-1	findbugs	3m 41s	patch/hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server no findbugs output file (hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/target/findbugsXml.xml)
-1	javadoc	3m 58s	hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.8.0_66 with JDK v1.8.0_66 generated 28 new issues (was 544, now 552).
+1	javadoc	2m 10s	the patch passed with JDK v1.8.0_66
-1	javadoc	9m 4s	hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.7.0_91 with JDK v1.7.0_91 generated 1 new issues (was 3, now 4).
+1	javadoc	4m 43s	the patch passed with JDK v1.7.0_91
+1	unit	0m 24s	hadoop-yarn-api in the patch passed with JDK v1.8.0_66.
+1	unit	2m 5s	hadoop-yarn-common in the patch passed with JDK v1.8.0_66.
-1	unit	77m 24s	hadoop-yarn-server in the patch failed with JDK v1.8.0_66.
+1	unit	0m 26s	hadoop-yarn-api in the patch passed with JDK v1.7.0_91.
+1	unit	2m 10s	hadoop-yarn-common in the patch passed with JDK v1.7.0_91.
-1	unit	79m 38s	hadoop-yarn-server in the patch failed with JDK v1.7.0_91.
+1	asflicense	0m 22s	Patch does not generate ASF License warnings.
		218m 47s	

Reason	Tests
JDK v1.8.0_66 Failed junit tests	hadoop.yarn.server.resourcemanager.TestClientRMTokens
	hadoop.yarn.server.resourcemanager.TestAMAuthorization
JDK v1.7.0_91 Failed junit tests	hadoop.yarn.server.resourcemanager.TestClientRMTokens
	hadoop.yarn.server.resourcemanager.TestAMAuthorization

Subsystem	Report/Notes
-----------	--------------

Docker	Image:yetus/hadoop:0ca8df7
JIRA Patch URL	https://issues.apache.org/jira/secure/attachment/12780628/YARN-4265-trunk.003.patch
JIRA Issue	YARN-4265
Optional Tests	asflicense compile javac javadoc mvninstall mvnsite unit findbugs checkstyle xml
uname	Linux 1abf81a2424a 3.13.0-36-lowlatency #63-Ubuntu SMP PREEMPT Wed Sep 3 21:56:12 UTC 2014 x86_64 x86_64 x86_64 GNU/Linux
Build tool	maven
Personality	/testptch/hadoop/patchprocess/precommit/personality/provided.sh
git revision	trunk / dec8fed
Default Java	1.7.0_91
Multi-JDK versions	/usr/lib/jvm/java-8-oracle:1.8.0_66 /usr/lib/jvm/java-7-openjdk-amd64:1.7.0_91
findbugs	v3.0.0
checkstyle	https://builds.apache.org/job/PreCommit-YARN-Build/10161/artifact/patchprocess/diff-checkstyle-hadoop-yarn-project_hadoop-yarn.txt
javadoc	hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.8.0_66: https://builds.apache.org/job/PreCommit-YARN-Build/10161/artifact/patchprocess/diff-javadoc-javadoc-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.8.0_66.txt
javadoc	hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.7.0_91: https://builds.apache.org/job/PreCommit-YARN-Build/10161/artifact/patchprocess/diff-javadoc-javadoc-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.7.0_91.txt
unit	https://builds.apache.org/job/PreCommit-YARN-Build/10161/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.8.0_66.txt
unit	https://builds.apache.org/job/PreCommit-YARN-Build/10161/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.7.0_91.txt
unit test logs	https://builds.apache.org/job/PreCommit-YARN-Build/10161/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.8.0_66.txt https://builds.apache.org/job/PreCommit-YARN-Build/10161/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.7.0_91.txt
JDK v1.7.0_91 Test Results	https://builds.apache.org/job/PreCommit-YARN-Build/10161/testReport/
modules	C: hadoop-yarn-project/hadoop-yarn/hadoop-yarn-api hadoop-yarn-project/hadoop-yarn/hadoop-yarn-common hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-applicationhistoryservice hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-timeline-pluginstorage U: hadoop-yarn-project/hadoop-yarn
Max memory used	75MB
Powered by	Apache Yetus 0.2.0-SNAPSHOT http://yetus.apache.org
Console output	https://builds.apache.org/job/PreCommit-YARN-Build/10161/console

This message was automatically generated.

Junping Du added a comment - 06/Jan/16 17:18

Thanks gtCarrera9 for updating the patch.

Since this is a separate work than introducing the whole new storage (as in this JIRA), maybe we can address this as a new JIRA?

+1. Address this in a separated JIRA sounds good to me, and I can take a look at that demo patch in reviewing this JIRA.

Li Lu added a comment - 06/Jan/16 17:24

Sure. Linking this the [YARN-4545](#) which works on the distributed shell integration.

Junping Du added a comment - 08/Jan/16 16:03

Thanks gtCarrera9 for updating the patch.

trunk.003 patch is getting closer, more comments:

In yarn-default.xml,

+ How long the ATS v1.5 entity group file system storage systemwill keep an
+ application's data in the done directory.

"file system storage systemwill" should be "file system storage will"

In EntityCacheItem.java, inside of refreshCache():

```
+ if (needRefresh()) {
+   if (!appLogs.isDone()) {
+     appLogs.parseSummaryLogs();
+   } else if (appLogs.getDetailLogs().isEmpty()) {
+     appLogs.scanForLogs();
+   }
+   if (!appLogs.getDetailLogs().isEmpty()) {
```

I am a bit confused with logic here: if appLogs is not done yet, but its detail logs is empty, do we need to scanForLogs? If not, we should document the reason at the least.

```
+ if (appLogs.getDetailLogs().isEmpty()) {
+   LOG.debug("cache id {}'s detail log is empty!", groupId);
+ }
```

This is inside of case {{ if (!appLogs.getDetailLogs().isEmpty()) }}, does it possible for detailLogs to become empty again? If not, we should remove it?

```
+ // Only refresh the log that matches the cache id
+ if (log.getFilename().contains(groupId.toString())) {
```

If we have two groupIds: 114859476_01_1 and 114859476_01_11, the later one's log file name can match with previous groupId as well? If so, we may consider to match file name with cache id more exactly? The same case with code below {{if (log.getFilename().contains(groupId.toString())) }}

appLogs.getDetailLogs().removeAll(removeList); should be move out of for each, or removeList will be useless here.

In EntityGroupFSTimelineStore.java,

```
+ + "%04d" + Path.SEPARATOR // app num / 10000000
```

I think it should be app num / 1000,000 (1 million) but not 10 millions.

```
+ private final Map<TimelineEntityGroupId, EntityCacheItem> cachedLogs
+ = Collections.synchronizedMap(
+ new LinkedHashMap<TimelineEntityGroupId, EntityCacheItem>(
+ appCacheMaxSize + 1, 0.75f, true) {
```

It should move cachedLogs initialization into serviceInit after appCacheMaxSize is setup already.

```
+ @SuppressWarnings("serial")
+ @Override
+ protected void serviceInit(Configuration conf) throws Exception {
```

SuppressWarnings("serial") seems to be unnecessary?

In loadPlugins(),

```
+ } catch (Exception e) {
+   LOG.info("Error loading plugin " + name, e);
+ }
```

It should use warn instead of info level for log message.

```
+ LOG.debug("Load plugin class {}", cacheIdPlugin.getClass().getName());
```

Better to use info instead of debug to follow the same practice with other messages in serviceInit.

For cleanLogs(Path dirpath), it seems like the execution result of cleanup log depends on the order of files/directories returned. Say an app dir include: file A, dir B, file A is a fresh one and all files in dir B are older than logRetainMillis. If file A get return first, the cleanLogs() do nothing, but if dir B get return first, cleanLogs() will cleanup dir B. Give fs.listFilesIterator(dirpath) could return file A, dir B in randomly order, is this randomly behavior expected?

```
// scans for new logs and returns the modification timestamp of the
// most recently modified log
@InterfaceAudience.Private
@VisibleForTesting
long scanForLogs() throws IOException {
```

Can we directly return appDirPath's modification time instead of go through all sub directories? Besides domain, summary, entity logs, anything else could exist under this directory?

```
for (LogInfo log : summaryLogs) {
  if (log.getFilename().equals(filename)
      && log.getAttemptDirName().equals(attemptDirName)) {
    return;
  }
}
```

Is it a common case for a AppLogs have many summaryLogs (and detail logs)? I guess not, but if my guess is wrong, we may should to use hashmap instead of list which costs to much time in search.

In LogInfo.parsePath(),

```
try {
    in.seek(offset);
    try {
        parser = jsonFactory.createJsonParser(in);
        parser.configure(JsonParser.Feature.AUTO_CLOSE_SOURCE, false);
    } catch (IOException e) {
        // if app hasn't completed then there may be errors due to the
        // incomplete file which are treated as EOF until app completes
        if (appCompleted) {
            throw e;
        } else {
            LOG.debug("Exception in parse path: {}", e.getMessage());
        }
    }
}
return doParse(tdm, parser, objMapper, ugi, appCompleted);
```

Even we get rid of throw exception for app incomplected case, doParse() below will throw a NPE exception. Is this expected?

In PluginStoreTestUtils.java,

```
static TimelineEntities generateTestEntities() {
    TimelineEntities entities = new TimelineEntities();
```

Indentation issue here.

Li Lu added a comment - 09/Jan/16 01:33 - edited

Thanks [junping_du](#) for the review! I updated my patch according to your comments. Some quick comments:

I am a bit confused with logic here: if appLogs is not done yet, but its detail logs is empty, do we need to scanForLogs? If not, we should document the reason at the least.

Yes, we only update summary logs when the app is running. Updated comments for this.

If we have two groupIds: 114859476_01_1 and 114859476_01_11, the later one's log file name can match with previous groupId as well? If so, we may consider to match file name with cache id more exactly? The same case with code below {{if (log.getFilename().contains(groupId.toString())) }}

Nice catch! What I'm trying to address here is the names with entity group id and a sequence number. I've updated related logic here.

For cleanLogs(Path dirpath), it seems like the execution result of cleanup log depends on the order of files/directories returned. Say an app dir include: file A, dir B, file A is a fresh one and all files in dir B are older than logRetainMillis. If file A get return first, the cleanLogs() do nothing, but if dir B get return first, cleanLogs() will clenup dir B. Give fs.listFiles(dirpath) could return file A, dir B in randomly order, is this randomly behavior expected?

This is not possible because in the first part of cleanLogs(), we're only doing a DFS to decide if we need to remove this dir. If any file in the directory is new, we will not remove it. The detailed remove logic happens after the DFS process.

Is it a common case for a AppLogs have many summaryLogs (and detail logs)?

Right now we're not facing this kind of use case. We can certainly optimize this logic in future though.

Can we directly return appDirPath's modification time instead of go through all sub directories?

I believe we cannot. We're trying to return the latest time any file within a directory has been changed to decide if the app is in UNKNOWN state for long enough in parseSummaryLogs.

Hadoop QA added a comment - 09/Jan/16 05:27

-1 overall

Vote	Subsystem	Runtime	Comment
0	reexec	0m 0s	Docker mode activated.
+1	@author	0m 0s	The patch does not contain any @author tags.
+1	test4tests	0m 0s	The patch appears to include 8 new or modified test files.
+1	mvninstall	7m 27s	trunk passed
+1	compile	1m 45s	trunk passed with JDK v1.8.0_66
+1	compile	2m 3s	trunk passed with JDK v1.7.0_91
+1	checkstyle	0m 30s	trunk passed
+1	mvnsite	2m 27s	trunk passed

+1	mvneclipse	0m 43s	trunk passed
-1	findbugs	3m 16s	branch/hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server no findbugs output file (hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/target/findbugsXml.xml)
-1	javadoc	0m 44s	hadoop-yarn-server in trunk failed with JDK v1.8.0_66.
+1	javadoc	4m 21s	trunk passed with JDK v1.7.0_91
+1	mvninstall	2m 4s	the patch passed
+1	compile	1m 46s	the patch passed with JDK v1.8.0_66
+1	javac	1m 46s	the patch passed
+1	compile	2m 6s	the patch passed with JDK v1.7.0_91
+1	javac	2m 6s	the patch passed
-1	checkstyle	0m 31s	Patch generated 9 new checkstyle issues in hadoop-yarn-project/hadoop-yarn (total was 292, now 300).
+1	mvnsite	2m 29s	the patch passed
+1	mvneclipse	0m 37s	the patch passed
+1	whitespace	0m 0s	Patch has no whitespace issues.
+1	xml	0m 2s	The patch has no ill-formed XML file.
-1	findbugs	3m 23s	patch/hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server no findbugs output file (hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/target/findbugsXml.xml)
-1	javadoc	0m 42s	hadoop-yarn-server in the patch failed with JDK v1.8.0_66.
-1	javadoc	7m 54s	hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.7.0_91 with JDK v1.7.0_91 generated 1 new issues (was 3, now 4).
+1	javadoc	4m 18s	the patch passed with JDK v1.7.0_91
+1	unit	0m 20s	hadoop-yarn-api in the patch passed with JDK v1.8.0_66.
+1	unit	1m 50s	hadoop-yarn-common in the patch passed with JDK v1.8.0_66.
-1	unit	72m 13s	hadoop-yarn-server in the patch failed with JDK v1.8.0_66.
+1	unit	0m 23s	hadoop-yarn-api in the patch passed with JDK v1.7.0_91.
+1	unit	2m 8s	hadoop-yarn-common in the patch passed with JDK v1.7.0_91.
-1	unit	75m 2s	hadoop-yarn-server in the patch failed with JDK v1.7.0_91.
+1	asflicense	0m 22s	Patch does not generate ASF License warnings.
		202m 18s	

Reason	Tests
JDK v1.8.0_66 Failed junit tests	hadoop.yarn.server.resourcemanager.TestClientRMTokens
	hadoop.yarn.server.resourcemanager.TestAMAuthorization
JDK v1.7.0_91 Failed junit tests	hadoop.yarn.server.resourcemanager.TestClientRMTokens
	hadoop.yarn.server.resourcemanager.TestAMAuthorization

Subsystem	Report/Notes
Docker	Image:yetus/hadoop:0ca8df7
JIRA Patch URL	https://issues.apache.org/jira/secure/attachment/12781349/YARN-4265-trunk.004.patch
JIRA Issue	YARN-4265
Optional Tests	asflicense compile javac javadoc mvninstall mvnsite unit findbugs checkstyle xml
uname	Linux e01d1240fbac 3.13.0-36-lowlatency #63-Ubuntu SMP PREEMPT Wed Sep 3 21:56:12 UTC 2014 x86_64 x86_64 x86_64 GNU/Linux
Build tool	maven
Personality	/testptch/hadoop/patchprocess/precommit/personality/provided.sh
git revision	trunk / 109e528
Default Java	1.7.0_91
Multi-JDK versions	/usr/lib/jvm/java-8-oracle:1.8.0_66 /usr/lib/jvm/java-7-openjdk-amd64:1.7.0_91
findbugs	v3.0.0

javadoc	https://builds.apache.org/job/PreCommit-YARN-Build/10217/artifact/patchprocess/branch-javadoc-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.8.0_66.txt
checkstyle	https://builds.apache.org/job/PreCommit-YARN-Build/10217/artifact/patchprocess/diff-checkstyle-hadoop-yarn-project_hadoop-yarn.txt
javadoc	https://builds.apache.org/job/PreCommit-YARN-Build/10217/artifact/patchprocess/patch-javadoc-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.8.0_66.txt
javadoc	hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.7.0_91: https://builds.apache.org/job/PreCommit-YARN-Build/10217/artifact/patchprocess/diff-javadoc-javadoc-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.7.0_91.txt
unit	https://builds.apache.org/job/PreCommit-YARN-Build/10217/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.8.0_66.txt
unit	https://builds.apache.org/job/PreCommit-YARN-Build/10217/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.7.0_91.txt
unit test logs	https://builds.apache.org/job/PreCommit-YARN-Build/10217/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.8.0_66.txt https://builds.apache.org/job/PreCommit-YARN-Build/10217/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.7.0_91.txt
JDK v1.7.0_91 Test Results	https://builds.apache.org/job/PreCommit-YARN-Build/10217/testReport/
modules	C: hadoop-yarn-project/hadoop-yarn/hadoop-yarn-api hadoop-yarn-project/hadoop-yarn/hadoop-yarn-common hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-applicationhistoryservice hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-timeline-pluginstorage U: hadoop-yarn-project/hadoop-yarn
Max memory used	76MB
Powered by	Apache Yetus 0.2.0-SNAPSHOT http://yetus.apache.org
Console output	https://builds.apache.org/job/PreCommit-YARN-Build/10217/console

This message was automatically generated.

Junping Du added a comment - 11/Jan/16 16:43

Thanks [gtCarrera9](#) for reply my comments above.

This is not possible because in the first part of `cleanLogs()`, we're only doing a DFS to decide if we need to remove this dir. If any file in the directory is new, we will not remove it. The detailed remove logic happens after the DFS process.

If our intention is to a DFS (deep first search), then I would say the implementation of `cleanLogs()` failed to achieve this because the flag of `shouldClean` in recursively calling `cleanLogs()` doesn't passed up to upper caller to reflect `shouldClean` flag there. We should fix it with a unit test.

I believe we cannot. We're trying to return the latest time any file within a directory has been changed to decide if the app is in UNKNOWN state for long enough in `parseSummaryLogs`.

For HDFS, any file update should reflect the modification time to parent directory. Any exception case here?

Mingliang Liu added a comment - 11/Jan/16 22:05 - edited

Hi [junping_du](#) and [gtCarrera9](#),

For HDFS, any file update should reflect the modification time to parent directory. Any exception case here?

For creating/moving files in an existing directory, the modification time of the parent directory will be updated. Refer to [code here](#).

However, should this jira consider appending to an existing file, the modification time of the parent directory will not be updated. Refer to [code here](#).

Thanks.

Li Lu added a comment - 12/Jan/16 01:56 - edited


Thanks [junping_du](#) and [liuml07](#) for the help! Since we need to handle files with appends, we cannot directly use directory modification time to decide if the contents of a directory has been changed. This means we also need to change some logics in the `cleanLogs` method. I redesigned `cleanLogs` method to perform a log scan with two methods:

Method 1: For the given directory, search (in a depth first fashion) to find out application log directories. For each of them, call method 2.

Method 2: For the given application log directory, search all files inside. If there exists a file that has been "recently" (as defined by the configs) updates, skip removing this directory. Otherwise, remove this application log directory.

In this way we can search inside a directory for all application log directories that need to be reclaimed.

According to Junping's suggestion, I've also added a new unit test (`testCleanLogs`) to cover common cases for the `cleanLogs` method.

▼  Hadoop QA added a comment - 12/Jan/16 05:46

 -1 overall

Vote	Subsystem	Runtime	Comment
0	reexec	0m 0s	Docker mode activated.
+1	@author	0m 0s	The patch does not contain any @author tags.
+1	test4tests	0m 0s	The patch appears to include 8 new or modified test files.
+1	mvninstall	7m 34s	trunk passed
+1	compile	1m 52s	trunk passed with JDK v1.8.0_66
+1	compile	2m 6s	trunk passed with JDK v1.7.0_91
+1	checkstyle	0m 30s	trunk passed
+1	mvnsite	2m 30s	trunk passed
+1	mvneclipse	0m 42s	trunk passed
-1	findbugs	3m 17s	branch/hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server no findbugs output file (hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/target/findbugsXml.xml)
-1	javadoc	0m 44s	hadoop-yarn-server in trunk failed with JDK v1.8.0_66.
+1	javadoc	4m 29s	trunk passed with JDK v1.7.0_91
+1	mvninstall	2m 7s	the patch passed
+1	compile	1m 50s	the patch passed with JDK v1.8.0_66
+1	javac	1m 50s	the patch passed
+1	compile	2m 8s	the patch passed with JDK v1.7.0_91
+1	javac	2m 8s	the patch passed
-1	checkstyle	0m 31s	Patch generated 9 new checkstyle issues in hadoop-yarn-project/hadoop-yarn (total was 292, now 300).
+1	mvnsite	2m 28s	the patch passed
+1	mvneclipse	0m 36s	the patch passed
+1	whitespace	0m 0s	Patch has no whitespace issues.
+1	xml	0m 2s	The patch has no ill-formed XML file.
-1	findbugs	3m 25s	patch/hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server no findbugs output file (hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/target/findbugsXml.xml)
-1	javadoc	1m 3s	hadoop-yarn-server in the patch failed with JDK v1.8.0_66.
-1	javadoc	8m 49s	hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.7.0_91 with JDK v1.7.0_91 generated 1 new issues (was 3, now 4).
+1	javadoc	4m 32s	the patch passed with JDK v1.7.0_91
+1	unit	0m 28s	hadoop-yarn-api in the patch passed with JDK v1.8.0_66.
+1	unit	2m 14s	hadoop-yarn-common in the patch passed with JDK v1.8.0_66.
-1	unit	72m 10s	hadoop-yarn-server in the patch failed with JDK v1.8.0_66.
+1	unit	0m 23s	hadoop-yarn-api in the patch passed with JDK v1.7.0_91.
+1	unit	2m 7s	hadoop-yarn-common in the patch passed with JDK v1.7.0_91.
-1	unit	74m 7s	hadoop-yarn-server in the patch failed with JDK v1.7.0_91.
+1	asflicense	0m 22s	Patch does not generate ASF License warnings.
		203m 27s	

Reason	Tests
JDK v1.8.0_66 Failed junit tests	hadoop.yarn.server.resourcemanager.TestClientRMTokens
	hadoop.yarn.server.resourcemanager.scheduler.fair.TestFairScheduler
	hadoop.yarn.server.resourcemanager.TestAMAuthorization
JDK v1.7.0_91 Failed junit tests	hadoop.yarn.server.resourcemanager.TestClientRMTokens
	hadoop.yarn.server.resourcemanager.TestAMAuthorization

Subsystem	Report/Notes
Docker	Image:yetus/hadoop:0ca8df7
JIRA Patch URL	https://issues.apache.org/jira/secure/attachment/12781700/YARN-4265-trunk.005.patch
JIRA Issue	YARN-4265
Optional Tests	asflicense compile javac javadoc mvninstall mvnsite unit findbugs checkstyle xml
uname	Linux e2474be5d70e 3.13.0-36-lowlatency #63-Ubuntu SMP PREEMPT Wed Sep 3 21:56:12 UTC 2014 x86_64 x86_64 x86_64 GNU/Linux
Build tool	maven
Personality	/testptch/hadoop/patchprocess/precommit/personality/provided.sh
git revision	trunk / c2e2e13
Default Java	1.7.0_91
Multi-JDK versions	/usr/lib/jvm/java-8-oracle:1.8.0_66 /usr/lib/jvm/java-7-openjdk-amd64:1.7.0_91
findbugs	v3.0.0
javadoc	https://builds.apache.org/job/PreCommit-YARN-Build/10237/artifact/patchprocess/branch-javadoc-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.8.0_66.txt
checkstyle	https://builds.apache.org/job/PreCommit-YARN-Build/10237/artifact/patchprocess/diff-checkstyle-hadoop-yarn-project_hadoop-yarn.txt
javadoc	https://builds.apache.org/job/PreCommit-YARN-Build/10237/artifact/patchprocess/patch-javadoc-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.8.0_66.txt
javadoc	hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.7.0_91: https://builds.apache.org/job/PreCommit-YARN-Build/10237/artifact/patchprocess/diff-javadoc-javadoc-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.7.0_91.txt
unit	https://builds.apache.org/job/PreCommit-YARN-Build/10237/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.8.0_66.txt
unit	https://builds.apache.org/job/PreCommit-YARN-Build/10237/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.7.0_91.txt
unit test logs	https://builds.apache.org/job/PreCommit-YARN-Build/10237/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.8.0_66.txt https://builds.apache.org/job/PreCommit-YARN-Build/10237/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.7.0_91.txt
JDK v1.7.0_91 Test Results	https://builds.apache.org/job/PreCommit-YARN-Build/10237/testReport/
modules	C: hadoop-yarn-project/hadoop-yarn/hadoop-yarn-api hadoop-yarn-project/hadoop-yarn/hadoop-yarn-common hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server/hadoop-yarn-server-applicationhistoryservice hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-timeline-pluginstorage U: hadoop-yarn-project/hadoop-yarn
Max memory used	76MB
Powered by	Apache Yetus 0.2.0-SNAPSHOT http://yetus.apache.org
Console output	https://builds.apache.org/job/PreCommit-YARN-Build/10237/console

This message was automatically generated.

Junping Du added a comment - 13/Jan/16 18:04

Thanks [gtCarrera9](#) for patient and updating the patch.

A few minor comments:

For new algorithm in cleanLogs(), I think it only works for specific log directory/file structure. Isn't it? If so, can you document it as javadoc on specific log structure it applies on?

For new added test case, please correct some naming issues with DIRs in each case. Also, instead of calling `fs.getFileStatus(appDirClean)` and catch exceptions for file/directory non-existing, it is more convenient to use `fs.exists(Path)`.

Please correct javadoc warning reported by Jenkins also.

Li Lu added a comment - 13/Jan/16 22:28 - edited

Thanks [junping_du](#) for the review! I addressed all of you concerns in the latest patch. Any other commnets? cc/[xgong](#), [jlowe](#)

Hadoop QA added a comment - 14/Jan/16 01:28

 -1 overall

Vote	Subsystem	Runtime	Comment
0	reexec	0m 0s	Docker mode activated.
+1	@author	0m 0s	The patch does not contain any @author tags.
+1	test4tests	0m 0s	The patch appears to include 8 new or modified test files.
0	mvndep	0m 34s	Maven dependency ordering for branch
+1	mvninstall	7m 27s	trunk passed
+1	compile	1m 53s	trunk passed with JDK v1.8.0_66
+1	compile	2m 7s	trunk passed with JDK v1.7.0_91
+1	checkstyle	0m 30s	trunk passed
+1	mvnsite	2m 52s	trunk passed
+1	mvneclipse	0m 55s	trunk passed
-1	findbugs	3m 16s	branch/hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server no findbugs output file (hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/target/findbugsXml.xml)
+1	javadoc	2m 13s	trunk passed with JDK v1.8.0_66
+1	javadoc	4m 48s	trunk passed with JDK v1.7.0_91
0	mvndep	0m 43s	Maven dependency ordering for patch
+1	mvninstall	2m 47s	the patch passed
+1	compile	1m 54s	the patch passed with JDK v1.8.0_66
+1	javac	1m 54s	the patch passed
+1	compile	2m 11s	the patch passed with JDK v1.7.0_91
+1	javac	2m 11s	the patch passed
-1	checkstyle	0m 31s	Patch generated 9 new checkstyle issues in hadoop-yarn-project/hadoop-yarn (total was 292, now 300).
+1	mvnsite	3m 19s	the patch passed
+1	mvneclipse	1m 10s	the patch passed
+1	whitespace	0m 0s	Patch has no whitespace issues.
+1	xml	0m 2s	The patch has no ill-formed XML file.
-1	findbugs	3m 58s	patch/hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server no findbugs output file (hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/target/findbugsXml.xml)
-1	javadoc	4m 1s	hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.8.0_66 with JDK v1.8.0_66 generated 7 new issues (was 544, now 551).
-1	javadoc	4m 1s	hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server_hadoop-yarn-server-applicationhistoryservice-jdk1.8.0_66 with JDK v1.8.0_66 generated 100 new issues (was 100, now 200).
-1	javadoc	4m 1s	hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server_hadoop-yarn-server-applicationhistoryservice-jdk1.8.0_66 with JDK v1.8.0_66 generated 100 new issues (was 100, now 200).
+1	javadoc	2m 53s	the patch passed with JDK v1.8.0_66
+1	javadoc	5m 20s	the patch passed with JDK v1.7.0_91
+1	unit	0m 24s	hadoop-yarn-api in the patch passed with JDK v1.8.0_66.
+1	unit	2m 11s	hadoop-yarn-common in the patch passed with JDK v1.8.0_66.
-1	unit	16m 33s	hadoop-yarn-server in the patch failed with JDK v1.8.0_66.
+1	unit	3m 10s	hadoop-yarn-server-applicationhistoryservice in the patch passed with JDK v1.8.0_66.
+1	unit	3m 16s	hadoop-yarn-server-applicationhistoryservice in the patch passed with JDK v1.8.0_66.
+1	unit	0m 29s	hadoop-yarn-api in the patch passed with JDK v1.7.0_91.
+1	unit	2m 27s	hadoop-yarn-common in the patch passed with JDK v1.7.0_91.
-1	unit	76m 3s	hadoop-yarn-server in the patch failed with JDK v1.7.0_91.
+1	unit	3m 26s	hadoop-yarn-server-applicationhistoryservice in the patch passed with JDK v1.7.0_91.
+1	unit	3m 21s	hadoop-yarn-server-applicationhistoryservice in the patch passed with JDK v1.7.0_91.
+1	asflicense	0m 22s	Patch does not generate ASF License warnings.
		172m 55s	

Reason	Tests
--------	-------

JDK v1.7.0_91 Failed junit tests	hadoop.yarn.server.resourcemanager.TestClientRMTokens
	hadoop.yarn.server.resourcemanager.TestAMAuthorization

Subsystem	Report/Notes
Docker	Image:yetus/hadoop:0ca8df7
JIRA Patch URL	https://issues.apache.org/jira/secure/attachment/12782144/YARN-4265-trunk.006.patch
JIRA Issue	YARN-4265
Optional Tests	asflicense compile javac javadoc mvninstall mvnsite unit findbugs checkstyle xml
uname	Linux 2ee81e11226f 3.13.0-36-lowlatency #63-Ubuntu SMP PREEMPT Wed Sep 3 21:56:12 UTC 2014 x86_64 x86_64 x86_64 GNU/Linux
Build tool	maven
Personality	/testptch/hadoop/patchprocess/precommit/personality/provided.sh
git revision	trunk / 321072b
Default Java	1.7.0_91
Multi-JDK versions	/usr/lib/jvm/java-8-oracle:1.8.0_66 /usr/lib/jvm/java-7-openjdk-amd64:1.7.0_91
findbugs	v3.0.0
checkstyle	https://builds.apache.org/job/PreCommit-YARN-Build/10269/artifact/patchprocess/diff-checkstyle-hadoop-yarn-project_hadoop-yarn.txt
javadoc	hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.8.0_66: https://builds.apache.org/job/PreCommit-YARN-Build/10269/artifact/patchprocess/diff-javadoc-javadoc-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.8.0_66.txt
javadoc	hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server_hadoop-yarn-server-applicationhistoryservice-jdk1.8.0_66: https://builds.apache.org/job/PreCommit-YARN-Build/10269/artifact/patchprocess/diff-javadoc-javadoc-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server_hadoop-yarn-server-applicationhistoryservice-jdk1.8.0_66.txt
javadoc	hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server_hadoop-yarn-server-applicationhistoryservice-jdk1.8.0_66: https://builds.apache.org/job/PreCommit-YARN-Build/10269/artifact/patchprocess/diff-javadoc-javadoc-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server_hadoop-yarn-server-applicationhistoryservice-jdk1.8.0_66.txt
unit	https://builds.apache.org/job/PreCommit-YARN-Build/10269/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.8.0_66.txt
unit	https://builds.apache.org/job/PreCommit-YARN-Build/10269/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.7.0_91.txt
unit test logs	https://builds.apache.org/job/PreCommit-YARN-Build/10269/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.7.0_91.txt
JDK v1.7.0_91 Test Results	https://builds.apache.org/job/PreCommit-YARN-Build/10269/testReport/
modules	C: hadoop-yarn-project/hadoop-yarn/hadoop-yarn-api hadoop-yarn-project/hadoop-yarn/hadoop-yarn-common hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-applicationhistoryservice hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-timeline-pluginstorage U: hadoop-yarn-project/hadoop-yarn
Max memory used	76MB
Powered by	Apache Yetus 0.2.0-SNAPSHOT http://yetus.apache.org
Console output	https://builds.apache.org/job/PreCommit-YARN-Build/10269/console

This message was automatically generated.

Li Lu added a comment - 14/Jan/16 02:43

New patch adds the missing package-info.java that the check-style script is mourning about.

Li Lu added a comment - 14/Jan/16 02:43

Also, the failing unit tests are not related to the changes in this JIRA.

Hadoop QA added a comment - 14/Jan/16 07:46

-1 overall

Vote	Subsystem	Runtime	Comment
0	reexec	0m 0s	Docker mode activated.
+1	@author	0m 0s	The patch does not contain any @author tags.
+1	test4tests	0m 0s	The patch appears to include 8 new or modified test files.
0	mvndep	0m 50s	Maven dependency ordering for branch
+1	mvninstall	7m 41s	trunk passed
+1	compile	1m 46s	trunk passed with JDK v1.8.0_66
+1	compile	2m 7s	trunk passed with JDK v1.7.0_91
+1	checkstyle	0m 31s	trunk passed
+1	mvnsite	2m 55s	trunk passed
+1	mvneclipse	0m 58s	trunk passed
-1	findbugs	3m 16s	branch/hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server no findbugs output file (hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/target/findbugsXml.xml)
+1	javadoc	2m 15s	trunk passed with JDK v1.8.0_66
+1	javadoc	4m 59s	trunk passed with JDK v1.7.0_91
-1	mvndep	0m 27s	patch's hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server dependency:list failed
-1	mvndep	0m 44s	patch's hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-timeline-pluginstorage dependency:list failed
0	mvndep	0m 44s	Maven dependency ordering for patch
+1	mvninstall	2m 48s	the patch passed
+1	compile	1m 48s	the patch passed with JDK v1.8.0_66
+1	javac	1m 48s	the patch passed
+1	compile	2m 8s	the patch passed with JDK v1.7.0_91
+1	javac	2m 8s	the patch passed
-1	checkstyle	0m 33s	Patch generated 9 new checkstyle issues in hadoop-yarn-project/hadoop-yarn (total was 293, now 301).
+1	mvnsite	3m 21s	the patch passed
+1	mvneclipse	1m 5s	the patch passed
+1	whitespace	0m 0s	Patch has no whitespace issues.
+1	xml	0m 1s	The patch has no ill-formed XML file.
-1	findbugs	3m 32s	patch/hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server no findbugs output file (hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/target/findbugsXml.xml)
-1	javadoc	3m 21s	hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.8.0_66 with JDK v1.8.0_66 generated 7 new issues (was 544, now 551).
-1	javadoc	3m 22s	hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server_hadoop-yarn-server-applicationhistoryservice-jdk1.8.0_66 with JDK v1.8.0_66 generated 100 new issues (was 100, now 200).
-1	javadoc	3m 22s	hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server_hadoop-yarn-server-applicationhistoryservice-jdk1.8.0_66 with JDK v1.8.0_66 generated 100 new issues (was 100, now 200).
+1	javadoc	2m 21s	the patch passed with JDK v1.8.0_66
+1	javadoc	5m 5s	the patch passed with JDK v1.7.0_91
+1	unit	0m 20s	hadoop-yarn-api in the patch passed with JDK v1.8.0_66.
+1	unit	1m 53s	hadoop-yarn-common in the patch passed with JDK v1.8.0_66.
-1	unit	78m 36s	hadoop-yarn-server in the patch failed with JDK v1.8.0_66.
+1	unit	3m 41s	hadoop-yarn-server-applicationhistoryservice in the patch passed with JDK v1.8.0_66.
+1	unit	3m 43s	hadoop-yarn-server-applicationhistoryservice in the patch passed with JDK v1.8.0_66.
+1	unit	0m 24s	hadoop-yarn-api in the patch passed with JDK v1.7.0_91.
+1	unit	2m 9s	hadoop-yarn-common in the patch passed with JDK v1.7.0_91.
-1	unit	79m 59s	hadoop-yarn-server in the patch failed with JDK v1.7.0_91.
+1	unit	4m 8s	hadoop-yarn-server-applicationhistoryservice in the patch passed with JDK v1.7.0_91.
+1	unit	4m 5s	hadoop-yarn-server-applicationhistoryservice in the patch passed with JDK v1.7.0_91.
+1	asflicense	0m 19s	Patch does not generate ASF License warnings.
		239m 23s	

Reason	Tests
JDK v1.8.0_66 Failed junit tests	hadoop.yarn.server.resourcemanager.TestClientRMTokens
	hadoop.yarn.server.resourcemanager.TestAMAuthorization
JDK v1.7.0_91 Failed junit tests	hadoop.yarn.server.resourcemanager.TestClientRMTokens
	hadoop.yarn.server.resourcemanager.TestAMAuthorization

Subsystem	Report/Notes
Docker	Image:yetus/hadoop:0ca8df7
JIRA Patch URL	https://issues.apache.org/jira/secure/attachment/12782186/YARN-4265-trunk.007.patch
JIRA Issue	YARN-4265
Optional Tests	asflicense compile javac javadoc mvninstall mvnsite unit findbugs checkstyle xml
uname	Linux 44a210f396e3 3.13.0-36-lowlatency #63-Ubuntu SMP PREEMPT Wed Sep 3 21:56:12 UTC 2014 x86_64 x86_64 x86_64 GNU/Linux
Build tool	maven
Personality	/testptch/hadoop/patchprocess/precommit/personality/provided.sh
git revision	trunk / 8315582
Default Java	1.7.0_91
Multi-JDK versions	/usr/lib/jvm/java-8-oracle:1.8.0_66 /usr/lib/jvm/java-7-openjdk-amd64:1.7.0_91
findbugs	v3.0.0
mvndep	hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server: https://builds.apache.org/job/PreCommit-YARN-Build/10273/artifact/patchprocess/maven-patch-dependencylist-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server.txt
mvndep	hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server_hadoop-yarn-server-timeline-pluginstorage: https://builds.apache.org/job/PreCommit-YARN-Build/10273/artifact/patchprocess/maven-patch-dependencylist-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server_hadoop-yarn-server-timeline-pluginstorage.txt
checkstyle	https://builds.apache.org/job/PreCommit-YARN-Build/10273/artifact/patchprocess/diff-checkstyle-hadoop-yarn-project_hadoop-yarn.txt
javadoc	hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.8.0_66: https://builds.apache.org/job/PreCommit-YARN-Build/10273/artifact/patchprocess/diff-javadoc-javadoc-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.8.0_66.txt
javadoc	hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server_hadoop-yarn-server-applicationhistoryservice-jdk1.8.0_66: https://builds.apache.org/job/PreCommit-YARN-Build/10273/artifact/patchprocess/diff-javadoc-javadoc-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server_hadoop-yarn-server-applicationhistoryservice-jdk1.8.0_66.txt
javadoc	hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server_hadoop-yarn-server-applicationhistoryservice-jdk1.8.0_66: https://builds.apache.org/job/PreCommit-YARN-Build/10273/artifact/patchprocess/diff-javadoc-javadoc-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server_hadoop-yarn-server-applicationhistoryservice-jdk1.8.0_66.txt
unit	https://builds.apache.org/job/PreCommit-YARN-Build/10273/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.8.0_66.txt
unit	https://builds.apache.org/job/PreCommit-YARN-Build/10273/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.7.0_91.txt
unit test logs	https://builds.apache.org/job/PreCommit-YARN-Build/10273/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.8.0_66.txt https://builds.apache.org/job/PreCommit-YARN-Build/10273/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.7.0_91.txt
JDK v1.7.0_91 Test Results	https://builds.apache.org/job/PreCommit-YARN-Build/10273/testReport/
modules	C: hadoop-yarn-project/hadoop-yarn/hadoop-yarn-api hadoop-yarn-project/hadoop-yarn/hadoop-yarn-common hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-applicationhistoryservice hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-project/hadoop-yarn-server/hadoop-yarn-server-applicationhistoryservice hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-timeline-pluginstorage U: hadoop-yarn-project/hadoop-yarn
Max memory used	76MB
Powered by	Apache Yetus 0.2.0-SNAPSHOT http://yetus.apache.org
Console output	https://builds.apache.org/job/PreCommit-YARN-Build/10273/console

This message was automatically generated.

Junping Du added a comment - 14/Jan/16 15:11

The checkstyle complain:

```
./hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-timeline-pluginstorage/src/main/java/org/apache/hadoop/yarn/server
```

should be taken care of. Also, it sounds like maven dependency failed?

Li Lu added a comment - 14/Jan/16 17:51

New patch to address the checkstyle hidden field warning. I also tested locally for the maven dependency issue. I cleaned my local m2 repository, rebuilt hadoop, and ran maven dependency. I cannot reproduce the problem. Also, the problem appears to be intermittent in the past Jenkins runs.

Hadoop QA added a comment - 14/Jan/16 21:51

 -1 overall

Vote	Subsystem	Runtime	Comment
0	reexec	0m 0s	Docker mode activated.
+1	@author	0m 0s	The patch does not contain any @author tags.
+1	test4tests	0m 0s	The patch appears to include 8 new or modified test files.
0	mvndep	0m 38s	Maven dependency ordering for branch
+1	mvninstall	8m 39s	trunk passed
+1	compile	2m 31s	trunk passed with JDK v1.8.0_66
+1	compile	2m 25s	trunk passed with JDK v1.7.0_91
+1	checkstyle	0m 33s	trunk passed
+1	mvnsite	3m 11s	trunk passed
+1	mvneclipse	0m 59s	trunk passed
0	findbugs	0m 0s	Skipped branch modules with no Java source: hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server
+1	findbugs	3m 26s	trunk passed
+1	javadoc	2m 29s	trunk passed with JDK v1.8.0_66
+1	javadoc	5m 2s	trunk passed with JDK v1.7.0_91
-1	mvndep	0m 27s	patch's hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server dependency:list failed
-1	mvndep	0m 44s	patch's hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-timeline-pluginstorage dependency:list failed
0	mvndep	0m 44s	Maven dependency ordering for patch
+1	mvninstall	2m 58s	the patch passed
+1	compile	2m 3s	the patch passed with JDK v1.8.0_66
+1	javac	2m 3s	the patch passed
+1	compile	2m 23s	the patch passed with JDK v1.7.0_91
+1	javac	2m 23s	the patch passed
-1	checkstyle	0m 34s	Patch generated 8 new checkstyle issues in hadoop-yarn-project/hadoop-yarn (total was 292, now 299).
+1	mvnsite	3m 32s	the patch passed
+1	mvneclipse	1m 7s	the patch passed
+1	whitespace	0m 0s	Patch has no whitespace issues.
+1	xml	0m 2s	The patch has no ill-formed XML file.
0	findbugs	0m 0s	Skipped patch modules with no Java source: hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server
+1	findbugs	4m 34s	the patch passed
-1	javadoc	3m 39s	hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.8.0_66 with JDK v1.8.0_66 generated 7 new issues (was 544, now 551).
-1	javadoc	3m 40s	hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server_hadoop-yarn-server-applicationhistoryservice-jdk1.8.0_66 with JDK v1.8.0_66 generated 100 new issues (was 100, now 200).
-1	javadoc	3m 40s	hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server_hadoop-yarn-server-applicationhistoryservice-jdk1.8.0_66 with JDK v1.8.0_66 generated 100 new issues (was 100, now 200).
+1	javadoc	2m 35s	the patch passed with JDK v1.8.0_66

+1	javadoc	5m 17s	the patch passed with JDK v1.7.0_91
+1	unit	0m 22s	hadoop-yarn-api in the patch passed with JDK v1.8.0_66.
+1	unit	2m 0s	hadoop-yarn-common in the patch passed with JDK v1.8.0_66.
-1	unit	74m 3s	hadoop-yarn-server in the patch failed with JDK v1.8.0_66.
+1	unit	3m 8s	hadoop-yarn-server-applicationhistoryservice in the patch passed with JDK v1.8.0_66.
+1	unit	3m 9s	hadoop-yarn-server-applicationhistoryservice in the patch passed with JDK v1.8.0_66.
+1	unit	0m 26s	hadoop-yarn-api in the patch passed with JDK v1.7.0_91.
+1	unit	2m 18s	hadoop-yarn-common in the patch passed with JDK v1.7.0_91.
-1	unit	75m 59s	hadoop-yarn-server in the patch failed with JDK v1.7.0_91.
+1	unit	3m 20s	hadoop-yarn-server-applicationhistoryservice in the patch passed with JDK v1.7.0_91.
+1	unit	3m 19s	hadoop-yarn-server-applicationhistoryservice in the patch passed with JDK v1.7.0_91.
+1	asflicense	0m 20s	Patch does not generate ASF License warnings.
		225m 51s	

Reason	Tests
JDK v1.8.0_66 Failed junit tests	hadoop.yarn.server.resourcemanager.TestClientRMTokens
	hadoop.yarn.server.resourcemanager.TestAMAuthorization
JDK v1.7.0_91 Failed junit tests	hadoop.yarn.server.resourcemanager.TestClientRMTokens
	hadoop.yarn.server.resourcemanager.TestAMAuthorization

Subsystem	Report/Notes
Docker	Image:yetus/hadoop:0ca8df7
JIRA Patch URL	https://issues.apache.org/jira/secure/attachment/12782321/YARN-4265-trunk.008.patch
JIRA Issue	YARN-4265
Optional Tests	asflicense compile javac javadoc mvninstall mvnsite unit findbugs checkstyle xml
uname	Linux 360a943f7467 3.13.0-36-lowlatency #63-Ubuntu SMP PREEMPT Wed Sep 3 21:56:12 UTC 2014 x86_64 x86_64 x86_64 GNU/Linux
Build tool	maven
Personality	/testptch/hadoop/patchprocess/precommit/personality/provided.sh
git revision	trunk / 9d04f26
Default Java	1.7.0_91
Multi-JDK versions	/usr/lib/jvm/java-8-oracle:1.8.0_66 /usr/lib/jvm/java-7-openjdk-amd64:1.7.0_91
findbugs	v3.0.0
mvnprop	hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server: https://builds.apache.org/job/PreCommit-YARN-Build/10286/artifact/patchprocess/maven-patch-dependencylist-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server.txt
mvnprop	hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server_hadoop-yarn-server-timeline-pluginstorage: https://builds.apache.org/job/PreCommit-YARN-Build/10286/artifact/patchprocess/maven-patch-dependencylist-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server_hadoop-yarn-server-timeline-pluginstorage.txt
checkstyle	https://builds.apache.org/job/PreCommit-YARN-Build/10286/artifact/patchprocess/diff-checkstyle-hadoop-yarn-project_hadoop-yarn.txt
javadoc	hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.8.0_66: https://builds.apache.org/job/PreCommit-YARN-Build/10286/artifact/patchprocess/diff-javadoc-javadoc-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.8.0_66.txt
javadoc	hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server_hadoop-yarn-server-applicationhistoryservice-jdk1.8.0_66: https://builds.apache.org/job/PreCommit-YARN-Build/10286/artifact/patchprocess/diff-javadoc-javadoc-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server_hadoop-yarn-server-applicationhistoryservice-jdk1.8.0_66.txt
javadoc	hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server_hadoop-yarn-server-applicationhistoryservice-jdk1.8.0_66: https://builds.apache.org/job/PreCommit-YARN-Build/10286/artifact/patchprocess/diff-javadoc-javadoc-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server_hadoop-yarn-server-applicationhistoryservice-jdk1.8.0_66.txt
unit	https://builds.apache.org/job/PreCommit-YARN-Build/10286/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.8.0_66.txt

unit	https://builds.apache.org/job/PreCommit-YARN-Build/10286/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.7.0_91.txt
unit test logs	https://builds.apache.org/job/PreCommit-YARN-Build/10286/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.8.0_66.txt https://builds.apache.org/job/PreCommit-YARN-Build/10286/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server-jdk1.7.0_91.txt
JDK v1.7.0_91 Test Results	https://builds.apache.org/job/PreCommit-YARN-Build/10286/testReport/
modules	C: hadoop-yarn-project/hadoop-yarn/hadoop-yarn-api hadoop-yarn-project/hadoop-yarn/hadoop-yarn-common hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-applicationhistoryservice hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-timeline-pluginstorage U: hadoop-yarn-project/hadoop-yarn
Max memory used	76MB
Powered by	Apache Yetus 0.2.0-SNAPSHOT http://yetus.apache.org
Console output	https://builds.apache.org/job/PreCommit-YARN-Build/10286/console

This message was automatically generated.

Li Lu added a comment - 14/Jan/16 21:59

The maven dependency check uses a separate maven repository `-Dmaven.repo.local=/home/jenkins/yetus-m2/hadoop-trunk-0`. I tested the whole workflow for building from the scratch and it worked. The AHS-test jar has not been published to apache SNAPSHOT (timeline plugin is the first module to depend on this module) yet. Therefore the mvn dependency test appears to be unrelated here.

Junping Du added a comment - 15/Jan/16 15:41

Tried `mvn dependency:list` locally for `hadoop-yarn-server-timeline-pluginstorage` and `hadoop-yarn-server`, both get passed successfully. Agree that test failures are also unrelated. Looks like all my major comments has been addressed. [jlowe](#) and [xgong](#), do you have further comments here?

Jason Darrell Lowe added a comment - 15/Jan/16 21:49


+1 lgtn.

Hudson added a comment - 18/Jan/16 01:36


FAILURE: Integrated in Hadoop-trunk-Commit #9128 (See <https://builds.apache.org/job/Hadoop-trunk-Commit/9128/>)
[YARN-4265](#). Provide new timeline plugin storage to support fine-grained (junping_du: rev 02f597c5db36ded385413958bdee793ad7eda40e)

- hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-timeline-pluginstorage/src/main/java/org/apache/hadoop/yarn/server/timeline/package-info.java
- hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-applicationhistoryservice/pom.xml
- hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-timeline-pluginstorage/src/test/java/org/apache/hadoop/yarn/server/timeline/TestEntityGroupFSTimelineStore.java
- hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-timeline-pluginstorage/src/main/java/org/apache/hadoop/yarn/server/timeline/TimelineEntityGroupPlugin.java
- hadoop-yarn-project/hadoop-yarn/hadoop-yarn-api/src/main/java/org/apache/hadoop/yarn/conf/YarnConfiguration.java
- hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-timeline-pluginstorage/src/main/java/org/apache/hadoop/yarn/server/timeline/EntityCacheItem.java
- hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-applicationhistoryservice/src/test/java/org/apache/hadoop/yarn/server/applicationhistoryservice/webapp/TestAHSWebServices.java
- hadoop-yarn-project/hadoop-yarn/hadoop-yarn-common/src/main/resources/yarn-default.xml
- hadoop-yarn-project/CHANGES.txt
- hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-timeline-pluginstorage/src/test/java/org/apache/hadoop/yarn/server/timeline/TestLogInfo.java
- hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-timeline-pluginstorage/src/main/java/org/apache/hadoop/yarn/server/timeline/EntityGroupFSTimelineStore.java
- hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-timeline-pluginstorage/src/test/java/org/apache/hadoop/yarn/server/timeline/PluginStoreTestUtils.java
- hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-applicationhistoryservice/src/test/java/org/apache/hadoop/yarn/server/timeline/TestTimelineDataManager.java
- hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-timeline-pluginstorage/src/test/java/org/apache/hadoop/yarn/server/timeline/EntityGroupPlugInForTest.java
- hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-applicationhistoryservice/src/main/java/org/apache/hadoop/yarn/server/timeline/TimelineDataManager.java
- hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-timeline-pluginstorage/pom.xml
- hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-applicationhistoryservice/src/main/java/org/apache/hadoop/yarn/server/applicationhistoryservice/ApplicationHistoryServer.java


- `hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-timeline-pluginstorage/src/main/java/org/apache/hadoop/yarn/server/timeline/LogInfo.java`
- `hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-applicationhistoryservice/src/test/java/org/apache/hadoop/yarn/server/applicationhistoryservice/TestApplicationHistoryClientService.java`
- `hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-applicationhistoryservice/src/test/java/org/apache/hadoop/yarn/server/applicationhistoryservice/TestApplicationHistoryManagerOnTimelineStore.java`
- `hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/pom.xml`
- `hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-applicationhistoryservice/src/main/java/org/apache/hadoop/yarn/server/timeline/TimelineDataManagerMetrics.java`

▼  Junping Du added a comment - 18/Jan/16 01:41

Thanks [gtCarrera9](#) and [jlowe](#) for the patch contribution! I have commit the 08 patch to trunk, branch-2 and branch-2.8.

▼  Karthik Kambatla added a comment - 18/Jan/16 08:51

This seems to break branch-2.8 compilation (mvn clean install -DskipTests). Can someone look into it?

▼  Rohith Sharma K S added a comment - 18/Jan/16 10:43

I found the issue, failure is because newly added project pom.xml version still 3.0.0-SNAPSHOT in branch-2.8. It should be 2.8.0-SNAPSHOT. So, dependency download will fail. After making change to 2.8.0-SNAPSHOT, it passes me locally.

▼  Rohith Sharma K S added a comment - 18/Jan/16 10:47


The same issue is there in branch-2 also.

▼  Rohith Sharma K S added a comment - 18/Jan/16 11:00

Updating the addendum patch fixing compilation issue in branch-2/branch-2.8

▼  Rohith Sharma K S added a comment - 18/Jan/16 11:21

Committed addendum patch to branch-2 and branch-2.8 respectively to unblock compilation issue asap..

▼  Junping Du added a comment - 18/Jan/16 12:50

It is weird that my build on branch-2/branch-2.8 locally get passed before check in... may due to local cache from build against trunk before? Anyway, thanks kasha@cloudera.com and [rohithsharma](#) to report the issue and fix it.

▼ People

Assignee:

 Li Lu

Reporter:

 Li Lu

Votes:

 0 Vote for this issue

Watchers:

 9 Start watching this issue

▼ Dates

Created:

14/Oct/15 20:48

Updated:

25/Oct/19 20:27

Resolved:

18/Jan/16 01:41