



ShuffleHandler is not aware of disks that are added

Details

Type:	Bug	Status:	RESOLVED
Priority:	Major	Resolution:	Fixed
Affects Version/s:	None	Fix Version/s:	2.9.0, 2.8.3, 3.0.0
Component/s:	None		
Labels:	None		
Hadoop Flags:	Reviewed		

Description

The ShuffleHandler permanently remembers the list of "good" disks on NM startup. If disks later are added to the node then map tasks will start using them but the ShuffleHandler will not be aware of them. The end result is that the data cannot be shuffled from the node leading to fetch failures and re-runs of the map tasks.

Attachments

YARN-7244.001.patch	13 kB	22/Sep/17 15:04
YARN-7244.002.patch	14 kB	22/Sep/17 20:01
YARN-7244.003.patch	18 kB	11/Oct/17 22:31
YARN-7244.004.patch	18 kB	12/Oct/17 01:05
YARN-7244.005.patch	27 kB	12/Oct/17 17:51
YARN-7244.006.patch	30 kB	15/Oct/17 20:32
YARN-7244.007.patch	33 kB	16/Oct/17 02:17
YARN-7244.008.patch	33 kB	16/Oct/17 11:53
YARN-7244.009.patch	33 kB	16/Oct/17 11:59
YARN-7244.010.patch	32 kB	16/Oct/17 21:42
YARN-7244.011.patch	32 kB	17/Oct/17 18:39
YARN-7244.012.patch	32 kB	17/Oct/17 21:16
YARN-7244.013.patch	32 kB	18/Oct/17 00:54
YARN-7244-branch-2.8.001.patch	38 kB	30/Oct/17 15:31
YARN-7244-branch-2.8.002.patch	38 kB	30/Oct/17 15:58

Activity

↑


▼

Kuhu Shukla

added a comment - 22/Sep/17 15:14

v1 patch that adds a new LocalDirAllocator#getLocalPathToRead() that decides to filter the bad directories based on a boolean. Changing the original call would be more pervasive. The patch does modify the AllocatorPerContext#getLocalPathToRead() signature since that is a private static class to LocalDirAllocator. The ShuffleHandler uses a yarn config to decide whether or not to filter bad dirs. This value , when false will never take out bad directories and hence any changes to local dirs would not impact the shuffle handler reads. Even if the mkdirs and exists check fails we want the dirs to be listed in the localdirs member when the config is false. For testing reasons, I have added a getter to the IDirAllocator which is package private. Appreciate any comments/corrections to this patch.

Another way to handle this would have been to change the AuxiliaryServices to pass the NMContext or the LocalDirAllocator from the NM . The former approach needs nodemanager dependencies to be added and the latter is tricky as I am not sure how the

▼  Hadoop QA added a comment - 22/Sep/17 17:14

 -1 overall


Vote	Subsystem	Runtime	Comment
0	reexec	0m 20s	Docker mode activated.
			Prechecks
+1	@author	0m 0s	The patch does not contain any @author tags.
+1	test4tests	0m 0s	The patch appears to include 1 new or modified test files.
			trunk Compile Tests
0	mvndep	1m 40s	Maven dependency ordering for branch
+1	mvninstall	15m 29s	trunk passed
+1	compile	15m 39s	trunk passed
+1	checkstyle	2m 16s	trunk passed
+1	mvnsite	2m 13s	trunk passed
+1	findbugs	3m 22s	trunk passed
+1	javadoc	2m 1s	trunk passed
			Patch Compile Tests
0	mvndep	0m 20s	Maven dependency ordering for patch
+1	mvninstall	1m 38s	the patch passed
+1	compile	16m 50s	the patch passed
+1	javac	16m 50s	the patch passed
-0	checkstyle	2m 33s	root: The patch generated 2 new + 312 unchanged - 0 fixed = 314 total (was 312)
+1	mvnsite	2m 45s	the patch passed
+1	whitespace	0m 0s	The patch has no whitespace issues.
+1	findbugs	4m 47s	the patch passed
+1	javadoc	2m 7s	the patch passed
			Other Tests
+1	unit	10m 9s	hadoop-common in the patch passed.
-1	unit	0m 51s	hadoop-yarn-api in the patch failed.
+1	unit	0m 49s	hadoop-mapreduce-client-shuffle in the patch passed.
+1	asflicense	0m 48s	The patch does not generate ASF License warnings.
		113m 5s	

Reason	Tests
Failed junit tests	hadoop.yarn.conf.TestYarnConfigurationFields


Subsystem	Report/Notes
Docker	Image:yetus/hadoop:71bbb86
JIRA Issue	YARN-7244
JIRA Patch URL	https://issues.apache.org/jira/secure/attachment/12888518/YARN-7244.001.patch

Optional Tests	asflicense compile javac javadoc mvninstall mvnsite unit findbugs checkstyle
uname	Linux 114687c9751c 3.13.0-116-generic #163-Ubuntu SMP Fri Mar 31 14:13:22 UTC 2017 x86_64 x86_64 x86_64 GNU/Linux
Build tool	maven
Personality	/testptch/hadoop/patchprocess/precommit/personality/provided.sh
git revision	trunk / 908d8e9
Default Java	1.8.0_144
findbugs	v3.1.0-RC1
checkstyle	https://builds.apache.org/job/PreCommit-YARN-Build/17589/artifact/patchprocess/diff-checkstyle-root.txt
unit	https://builds.apache.org/job/PreCommit-YARN-Build/17589/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-api.txt
Test Results	https://builds.apache.org/job/PreCommit-YARN-Build/17589/testReport/
modules	C: hadoop-common-project/hadoop-common hadoop-yarn-project/hadoop-yarn/hadoop-yarn-api hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-shuffle U: .
Console output	https://builds.apache.org/job/PreCommit-YARN-Build/17589/console
Powered by	Apache Yetus 0.6.0-SNAPSHOT http://yetus.apache.org

This message was automatically generated.

▼  **Kuhu Shukla** added a comment - 22/Sep/17 20:01

Fixing minor test issue for newly added yarn config key.

▼  **Hadoop QA** added a comment - 22/Sep/17 21:59


 **+1 overall**

Vote	Subsystem	Runtime	Comment
0	reexec	0m 12s	Docker mode activated.
			Prechecks
+1	@author	0m 0s	The patch does not contain any @author tags.
+1	test4tests	0m 0s	The patch appears to include 1 new or modified test files.
			trunk Compile Tests
0	mvndep	0m 29s	Maven dependency ordering for branch
+1	mvninstall	13m 36s	trunk passed
+1	compile	15m 22s	trunk passed
+1	checkstyle	2m 10s	trunk passed
+1	mvnsite	2m 50s	trunk passed
+1	findbugs	4m 38s	trunk passed
+1	javadoc	2m 27s	trunk passed
			Patch Compile Tests
0	mvndep	0m 17s	Maven dependency ordering for patch
+1	mvninstall	2m 2s	the patch passed
+1	compile	11m 24s	the patch passed
+1	javac	11m 24s	the patch passed
-0	checkstyle	2m 10s	root: The patch generated 2 new + 312 unchanged - 0 fixed = 314 total (was 312)
+1	mvnsite	3m 2s	the patch passed
+1	whitespace	0m 0s	The patch has no whitespace issues.

+1	xml	0m 2s	The patch has no ill-formed XML file.
+1	findbugs	5m 30s	the patch passed
+1	javadoc	2m 45s	the patch passed
			Other Tests
+1	unit	8m 32s	hadoop-common in the patch passed.
+1	unit	0m 41s	hadoop-yarn-api in the patch passed.
+1	unit	2m 58s	hadoop-yarn-common in the patch passed.
+1	unit	0m 33s	hadoop-mapreduce-client-shuffle in the patch passed.
+1	asflicense	0m 34s	The patch does not generate ASF License warnings.
		108m 11s	

Subsystem	Report/Notes
Docker	Image:yetus/hadoop:71bbb86
JIRA Issue	YARN-7244
JIRA Patch URL	https://issues.apache.org/jira/secure/attachment/12888555/YARN-7244.002.patch
Optional Tests	asflicense compile javac javadoc mvninstall mvnsite unit findbugs checkstyle xml
uname	Linux 41fb21dc1d76 3.13.0-119-generic #166-Ubuntu SMP Wed May 3 12:18:55 UTC 2017 x86_64 x86_64 x86_64 GNU/Linux
Build tool	maven
Personality	/testptch/hadoop/patchprocess/precommit/personality/provided.sh
git revision	trunk / 08fca50
Default Java	1.8.0_144
findbugs	v3.1.0-RC1
checkstyle	https://builds.apache.org/job/PreCommit-YARN-Build/17598/artifact/patchprocess/diff-checkstyle-root.txt
Test Results	https://builds.apache.org/job/PreCommit-YARN-Build/17598/testReport/
modules	C: hadoop-common-project/hadoop-common hadoop-yarn-project/hadoop-yarn/hadoop-yarn-api hadoop-yarn-project/hadoop-yarn/hadoop-yarn-common hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-shuffle U: .
Console output	https://builds.apache.org/job/PreCommit-YARN-Build/17598/console
Powered by	Apache Yetus 0.6.0-SNAPSHOT http://yetus.apache.org

This message was automatically generated.

 Bibin Chundatt added a comment - 26/Sep/17 09:07

Thank you kshukla for patch. Minor comments from my side

1. Better to check directory exists first if we are not concerned of permission . thoughts?

```
348         if (ctx.localFS.mkdirs(tmpDir) || ctx.localFS.exists(tmpDir))
```

2. Testcase is successful even if YARN_SHUFFLE_BAD_DIRS_FILTER_ENABLED set to false.

```
conf.setBoolean(
    YarnConfiguration.YARN_SHUFFLE_BAD_DIRS_FILTER_ENABLED, false);
```

 Kuhu Shukla added a comment - 26/Sep/17 16:23 - edited

Thank you bibinchundatt for the review comments!

Better to check directory exists first if we are not concerned of permission . thoughts?

Makes sense to me.

Testcase is successful even if YARN_SHUFFLE_BAD_DIRS_FILTER_ENABLED set to false.


That is expected. This config essentially decides if we should stick to existing behavior (value=true) of removing directories if they are bad or not (value=false).

▼  Sunil G added a comment - 26/Sep/17 17:21

Thanks [kshukla](#) for this improvement.

Some doubts here.

1. ShuffleHandler is in mapred package. I think its not that safe to read YARN_SHUFFLE_BAD_DIRS_FILTER_ENABLED from MR end. We could push this config name to LocalDirAllocator and then read from NM end. Please correct me if I missed something.
2. A pull model to get all changed dirs from NM needs context as you mentioned. `LocalDirAlloc.getLocalPathToRead` could definitely skip failed/bad dirs on fly. If all dirs are bad, then MR has to be notified via exception. Do you think, we can improve this to skip as default behavior itself?

▼  Kuhu Shukla added a comment - 27/Sep/17 12:06


Thank you [sunilg](#) for the review comments!

We could push this config name to LocalDirAllocator and then read from NM end

I am not sure how we can initialize the config specifically for LocalDirAllocator (may be add a constructor?) and what would reading it from NM end mean. Agreed that code separation is important here. May be not having this as a config might help?

Do you think, we can improve this to skip as default behavior itself?

I did not fully get what you have in mind here. Could you help me understand and elaborate a bit. Thanks a lot!

▼  Jason Darrell Lowe added a comment - 27/Sep/17 14:56

Thanks for the patch!

The core issue here is that the NM is handing out directories to tasks that the shuffle manager is unaware of. This filtering-or-not approach doesn't completely solve the issue, since the ShuffleHandler will still attempt to visit disks that the NM has already determined are bad. That could cause performance problems if the ShuffleHandler tries to read a particularly problematic disk over and over as it searches for outputs to shuffle for every shuffle request.

It would be more ideal if the NM could convey to aux services what directories are in use. Then the ShuffleHandler and NM would be in sync with respect to what disks should or should not be used.

Another way to handle this would have been to change the AuxiliaryServices to pass the NMContext or the LocalDirAllocator from the NM .

That would be nice, as there are probably other things in the NMContext that aux services may want to know about. However we could always go with a much more direct route. We could add an API to AuxiliaryService that can set a callback object that can be leveraged to retrieve the current list of paths that are good for reading or writing, or we can an API to AuxiliaryService that the NM can call to update that service on the list of paths good for reading and writing. (i.e.: either a 'pull' or 'push' model for exposing the current good directories to aux services).

The 'pull' model requires an interface or abstract class in yarn-api that defines the API aux services can call to retrieve the directories, and we would put the actual implementation of that interface in yarn-server-nodemanager. Ideally the interface would look a lot like the existing `getLocalDirsForRead()`, `getLocalDirsForWrite()`, etc. of the `LocalDirsHandlerService` so it's an easy pass-through to implement on the nodemanager side.

The 'push' model requires adding a listener interface to `LocalDirsHandlerService` so we know when a disk is added or removed and can callback into each aux service to update them on the current list of dirs for reading and writing.

Haven't had a lot of time to figure out which would be more ideal in practice in terms of ease-of-use and performance, but I think I'd rather see the aux services be more in sync with the rest of the NM wrt. local dirs being actively used.

▼  Sunil G added a comment - 27/Sep/17 15:17

Thanks [jlowe](#) for adding more clarity on this.

'pull' model may be better and could work for all such cases. As Jason suggested if apps could know the latest dirs from `getLocalDirsForRead/Write`, shuffle handler will have a list of valid dirs always. Only potential issue which I see is that, once a set of dirs are pulled from `LocalDirAllocator#ctx.localDirs`, these dirs will be validated only when one more `getLocalPathForWrite/Read` is invoked. So there could be a window where we may get a stale dirs. If new api `LocalDirAllocator#getLocalDirsForRead` could call `confChanged`, then i think it should be a source of truth for localDirs for given time snapshot.

Do you think, we can improve this to skip as default behavior itself

Currently in this patch, you are trying to avoid disk validation check when `shouldFilter` is false. To add more context, may be we could skip this check here provided we have a valid dirs in `ShuffleHandler` end based on earlier api.

▼  Jason Darrell Lowe added a comment - 27/Sep/17 18:09


Only potential issue which I see is that, once a set of dirs are pulled from `LocalDirAllocator#ctx.localDirs`, these dirs will be validated only when one more `getLocalPathForWrite/Read` is invoked. So there could be a window where we may get a stale dirs.

I wouldn't worry too much about that window. Think of the much larger window a container gets, since it is only told once, on startup, what the list of valid dirs are. I think we're fine as long as aux services are notified fairly soon after a disk fails. It doesn't have to be instantaneous nor atomic. We could make a pull API where the aux service can essentially directly call the NM's `LocalDirHandlerService` for getting a path to read or a path to write, then the aux service doesn't even have to manage the directories itself if all it cares about is finding a place to write or read.

▼  Sunil G added a comment - 28/Sep/17 12:37

the aux service doesn't even have to manage the directories itself if all it cares about is finding a place to write or read

Yes. Thats perfectly fine. Thank you very much for clarifying. One more thought here, currently `ShuffleHandler` creates **`LocalDirAllocator`** and use that to get the better dirs to operate on. `LocalDirHandlerService` is not used by `ShuffleHandler` to now. I think we might not need `LocalDirHandlerService`, rather a new api as you mentioned in `LocalDirAllocator` named **`getLocalDirsForRead`** could enough to get valid dirs as it pulls all configured `NM_LOCAL_DIRS` and validates same.

▼  Kuhu Shukla added a comment - 28/Sep/17 14:22

Thank you [jlowe](#), [sunilg](#) for the review/comments.

We could make a pull API where the aux service can essentially directly call the NM's `LocalDirHandlerService` for getting a path to read or a path to write, then the aux service doesn't even have to manage the directories itself if all it cares about is finding a place to write or read.

A pull model where the `Shuffle handler` /aux service does not maintain valid dirs state would be my preference but the other pull approach would work too. I will start reworking the patch in the meantime and will finalize based on what we decide. Appreciate your thoughts.

▼  Jason Darrell Lowe added a comment - 28/Sep/17 15:31

rather a new api as you mentioned in `LocalDirAllocator` named `getLocalDirsForRead` could enough to get valid dirs as it pulls all configured `NM_LOCAL_DIRS` and validates same.

`LocalDirAllocator` should not have the new API, IMHO. That class is in `hadoop-common` and shouldn't be involved in solving this `nodemanager`-specific problem.

I'm thinking we go with the pull approach with something like the following. Note that I'm not stuck on the specific names of new interfaces/classes, they're just examples for reference.

1. In `AuxiliaryService` add new methods to get and set the API object to interact with the NM's local dirs management, e.g.:

```
public AuxiliaryLocalPathHandler getAuxiliaryLocalPathHandler();
```

```
public void setAuxiliaryLocalPathHandler(AuxiliaryLocalPathHandler);
```

2. The new AuxiliaryLocalPathHandler object would be in hadoop-yarn-api and look something like this:

```
public interface AuxiliaryLocalPathHandler {
    Path getLocalPathForRead(String);
    Path getLocalPathForWrite(String);
    Path getLocalPathForWrite(String, long);
}
```

- 3. AuxiliaryService would implement a LocalDirsHandler that maps the AuxiliaryLocalDirsHandler calls to the NMs LocalDirsHandlerService.
- 4. The ShuffleHandler can leverage the new AuxiliaryLocalPathHandler to find shuffle input files rather than manage its own LocalDirAllocator.

Kuhu Shukla added a comment - 11/Oct/17 22:32

Updated patch closer to the design Jason mentioned earlier. Adds a new Path Handler that is passed from the Containermanager -> AuxServices -> AuxiliaryService -> ShuffleHandler.
Appreciate any comments on the approach/patch. Thanks a lot!

Hadoop QA added a comment - 11/Oct/17 23:16

-1 overall

Vote	Subsystem	Runtime	Comment
0	reexec	0m 0s	Docker mode activated.
-1	patch	0m 9s	YARN-7244 does not apply to trunk. Rebase required? Wrong Branch? See https://wiki.apache.org/hadoop/HowToContribute for help.

Subsystem	Report/Notes
JIRA Issue	YARN-7244
JIRA Patch URL	https://issues.apache.org/jira/secure/attachment/12891577/YARN-7244.003.patch
Console output	https://builds.apache.org/job/PreCommit-YARN-Build/17866/console
Powered by	Apache Yetus 0.6.0-SNAPSHOT http://yetus.apache.org

This message was automatically generated.

Kuhu Shukla added a comment - 12/Oct/17 01:11

Rebasing patch on trunk.

Hadoop QA added a comment - 12/Oct/17 03:22

-1 overall

Vote	Subsystem	Runtime	Comment
0	reexec	0m 14s	Docker mode activated.
			Prechecks
+1	@author	0m 0s	The patch does not contain any @author tags.
+1	test4tests	0m 0s	The patch appears to include 3 new or modified test files.
			trunk Compile Tests
0	mvndep	0m 20s	Maven dependency ordering for branch
+1	mvninstall	13m 21s	trunk passed


+1	compile	14m 17s	trunk passed
+1	checkstyle	2m 3s	trunk passed
+1	mvnsite	1m 36s	trunk passed
+1	shadedclient	15m 4s	branch has no errors when building and testing our client artifacts.
+1	findbugs	2m 54s	trunk passed
+1	javadoc	1m 20s	trunk passed
Patch Compile Tests			
0	mvndep	0m 18s	Maven dependency ordering for patch
+1	mvninstall	1m 26s	the patch passed
+1	compile	12m 31s	the patch passed
+1	javac	12m 31s	the patch passed
-0	checkstyle	2m 12s	root: The patch generated 16 new + 325 unchanged - 1 fixed = 341 total (was 326)
+1	mvnsite	1m 52s	the patch passed
+1	whitespace	0m 0s	The patch has no whitespace issues.
+1	shadedclient	11m 29s	patch has no errors when building and testing our client artifacts.
+1	findbugs	3m 8s	the patch passed
+1	javadoc	1m 27s	the patch passed
Other Tests			
+1	unit	0m 43s	hadoop-yarn-api in the patch passed.
-1	unit	16m 46s	hadoop-yarn-server-nodemanager in the patch failed.
-1	unit	0m 35s	hadoop-mapreduce-client-shuffle in the patch failed.
+1	asflicense	0m 37s	The patch does not generate ASF License warnings.
		121m 8s	

Reason	Tests
Failed junit tests	hadoop.yarn.server.nodemanager.scheduler.TestDistributedScheduler
	hadoop.mapred.TestShuffleHandler


Subsystem	Report/Notes
Docker	Image:yetus/hadoop:3d04c00
JIRA Issue	YARN-7244
JIRA Patch URL	https://issues.apache.org/jira/secure/attachment/12891601/YARN-7244.004.patch
Optional Tests	asflicense compile javac javadoc mvninstall mvnsite unit shadedclient findbugs checkstyle
uname	Linux 0188628fbd85 3.13.0-119-generic #166-Ubuntu SMP Wed May 3 12:18:55 UTC 2017 x86_64 x86_64 x86_64 GNU/Linux
Build tool	maven
Personality	/testptch/hadoop/patchprocess/precommit/personality/provided.sh
git revision	trunk / 075358e
Default Java	1.8.0_144
findbugs	v3.1.0-RC1
checkstyle	https://builds.apache.org/job/PreCommit-YARN-Build/17876/artifact/patchprocess/diff-checkstyle-root.txt
unit	https://builds.apache.org/job/PreCommit-YARN-Build/17876/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server_hadoop-yarn-server-nodemanager.txt

unit	https://builds.apache.org/job/PreCommit-YARN-Build/17876/artifact/patchprocess/patch-unit-hadoop-mapreduce-project_hadoop-mapreduce-client_hadoop-mapreduce-client-shuffle.txt
Test Results	https://builds.apache.org/job/PreCommit-YARN-Build/17876/testReport/
modules	C: hadoop-yarn-project/hadoop-yarn/hadoop-yarn-api hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-nodemanager hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-shuffle U: .
Console output	https://builds.apache.org/job/PreCommit-YARN-Build/17876/console
Powered by	Apache Yetus 0.6.0-SNAPSHOT http://yetus.apache.org


This message was automatically generated.

▼  Kuhu Shukla added a comment - 12/Oct/17 12:32

Test failures are related. Will update shortly.

▼  Kuhu Shukla added a comment - 12/Oct/17 17:52

Fixing TestShuffleHandler failures. The TestDistributedScheduler failure is documented in [YARN-7299](#).

▼  Hadoop QA added a comment - 12/Oct/17 20:43

 -1 overall

Vote	Subsystem	Runtime	Comment
0	reexec	0m 21s	Docker mode activated.
			Prechecks
+1	@author	0m 0s	The patch does not contain any @author tags.
+1	test4tests	0m 0s	The patch appears to include 4 new or modified test files.
			trunk Compile Tests
0	mvndep	0m 26s	Maven dependency ordering for branch
+1	mvninstall	19m 53s	trunk passed
+1	compile	22m 27s	trunk passed
+1	checkstyle	2m 43s	trunk passed
+1	mvnsite	2m 43s	trunk passed
+1	shadedclient	19m 35s	branch has no errors when building and testing our client artifacts.
+1	findbugs	4m 11s	trunk passed
+1	javadoc	2m 1s	trunk passed
			Patch Compile Tests
0	mvndep	0m 26s	Maven dependency ordering for patch
+1	mvninstall	2m 3s	the patch passed
+1	compile	17m 18s	the patch passed
+1	javac	17m 18s	the patch passed
-0	checkstyle	2m 58s	root: The patch generated 43 new + 348 unchanged - 2 fixed = 391 total (was 350)
+1	mvnsite	2m 54s	the patch passed
+1	whitespace	0m 0s	The patch has no whitespace issues.
+1	shadedclient	13m 48s	patch has no errors when building and testing our client artifacts.
+1	findbugs	5m 20s	the patch passed
+1	javadoc	1m 52s	the patch passed
			Other Tests
+1	unit	1m 9s	hadoop-yarn-api in the patch passed.

-1	unit	18m 41s	hadoop-yarn-server-nodemanager in the patch failed.
-1	unit	1m 2s	hadoop-mapreduce-client-shuffle in the patch failed.
+1	asflicense	1m 2s	The patch does not generate ASF License warnings.
		160m 15s	

Reason	Tests
Failed junit tests	hadoop.yarn.server.nodemanager.scheduler.TestDistributedScheduler
	hadoop.mapred.TestShuffleHandler

Subsystem	Report/Notes
Docker	Image:yetus/hadoop:3d04c00
JIRA Issue	YARN-7244
JIRA Patch URL	https://issues.apache.org/jira/secure/attachment/12891761/YARN-7244.005.patch
Optional Tests	asflicense compile javac javadoc mvninstall mvnsite unit shadedclient findbugs checkstyle
uname	Linux 8bd46ea15b5d 3.13.0-119-generic #166-Ubuntu SMP Wed May 3 12:18:55 UTC 2017 x86_64 x86_64 x86_64 GNU/Linux
Build tool	maven
Personality	/testptch/hadoop/patchprocess/precommit/personality/provided.sh
git revision	trunk / 13fcfb3
Default Java	1.8.0_144
findbugs	v3.1.0-RC1
checkstyle	https://builds.apache.org/job/PreCommit-YARN-Build/17890/artifact/patchprocess/diff-checkstyle-root.txt
unit	https://builds.apache.org/job/PreCommit-YARN-Build/17890/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server_hadoop-yarn-server-nodemanager.txt
unit	https://builds.apache.org/job/PreCommit-YARN-Build/17890/artifact/patchprocess/patch-unit-hadoop-mapreduce-project_hadoop-mapreduce-client_hadoop-mapreduce-client-shuffle.txt
Test Results	https://builds.apache.org/job/PreCommit-YARN-Build/17890/testReport/
modules	C: hadoop-yarn-project/hadoop-yarn/hadoop-yarn-api hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-nodemanager hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-shuffle U: .
Console output	https://builds.apache.org/job/PreCommit-YARN-Build/17890/console
Powered by	Apache Yetus 0.6.0-SNAPSHOT http://yetus.apache.org

This message was automatically generated.

Jason Darrell Lowe added a comment - 13/Oct/17 15:44

Thanks for updating the patch! The test failure appears to be related, and it would be good to cleanup the checkstyle issues.

In TestShuffleHandler the code to build absLogDir is replicated a lot of times and is always the same value. This should be a final static computed once in TestShuffleHandler and referenced in the places that need it rather than copy-n-paste of the code. This also removes the need for the getAbsLogDir method. When computing absLogDir the code should not assume the target directory is the right place but rather look at the environment for direction where to go. GenericTestUtils#getTestDir would help here.

Rather than needing a TestShuffleHandler1 class that overrides getAuxiliaryLocalPathHandler, I think we can just use ShuffleHandler directly and call setAuxiliaryLocalPathHandler after creating it. That's what the nodemanager is going to do anyway during normal execution, and the unit test would be closer to the real-world scenario.

Do we really need to provide both a `TestAuxiliaryLocalPathHandler` and `TestAuxiliaryLocalPathHandler1`? The original test code used a `LocalDirAllocator` directly, and it seems like we could use one test handler class that passes through to the same local dir allocator that was originally setup in the unit test.

`AuxiliaryService#getAuxiliaryLocalPathHandler` and `AuxiliaryService#setAuxiliaryLocalPathHandler` need javadocs since this is a public class that we expect app framework providers to use.

`AuxServices` should take the path handler in its constructor since it's not legal to call any other methods before it has been set. This allows the handler field to be marked final and removes the need for get/set methods.

In `BaseContainerManagerTest` it's more appropriate to start the node health checker service rather than the dirs handler, since the health checker is responsible for initing and starting the dirs handler. Otherwise it looks weird that we started a service without calling init first.

`TestAuxServices` should use Mockito to build the `AuxiliaryLocalPathHandler` mock object.

`AuxService#getServices` returns a collection of `AuxiliaryService` objects, so it would be cleaner to iterate on `AuxiliaryService` objects directly rather than iterate on `Service` objects and cast them in the loop.

I'm not a fan of `testContainerAuxPathHandler` since it uses long sleeps to essentially test that the aux path handler is hooked up to the nodemanager's dir handler. It would be much better if we could eliminate the need for sleeps and instead set a mocked/controllable dirs handler in the containermanager then change the behavior of the mocked object and verify that change is reflected in the aux path handler is well. Or even simpler, `AuxiliaryLocalPathHandlerImpl` could take a dir handler as a constructor argument and change to a static class so we can instantiate it and test it directly rather than needing to bring up the entire container manager for this test.

- Jason Darrell Lowe added a comment - 13/Oct/17 20:03
- I forgot to mention that `AuxiliaryLocalPathHandler` should be marked Public, Evolving and also have javadocs since app framework providers will be expected to call these methods.
- Kuhu Shukla added a comment - 15/Oct/17 20:32
- Thank you for the comments/review [jlowe](#)! Updated patch. Will wait for PreCommit before any review requests.

- Hadoop QA added a comment - 15/Oct/17 22:43

-1 overall


Vote	Subsystem	Runtime	Comment
0	reexec	0m 17s	Docker mode activated.
			Prechecks
+1	@author	0m 0s	The patch does not contain any @author tags.
+1	test4tests	0m 0s	The patch appears to include 3 new or modified test files.
			trunk Compile Tests
0	mvndep	1m 36s	Maven dependency ordering for branch
+1	mvninstall	13m 40s	trunk passed
+1	compile	15m 5s	trunk passed
+1	checkstyle	2m 7s	trunk passed
+1	mvnsite	1m 38s	trunk passed
+1	shadedclient	15m 22s	branch has no errors when building and testing our client artifacts.
+1	findbugs	2m 38s	trunk passed
+1	javadoc	1m 19s	trunk passed
			Patch Compile Tests
0	mvndep	0m 19s	Maven dependency ordering for patch
+1	mvninstall	1m 13s	the patch passed
+1	compile	11m 36s	the patch passed
+1	javac	11m 36s	the patch passed

-0	checkstyle	2m 11s	root: The patch generated 8 new + 317 unchanged - 2 fixed = 325 total (was 319)
+1	mvnsite	1m 47s	the patch passed
+1	whitespace	0m 0s	The patch has no whitespace issues.
+1	shadedclient	11m 35s	patch has no errors when building and testing our client artifacts.
+1	findbugs	3m 9s	the patch passed
-1	javadoc	0m 32s	hadoop-yarn-project_hadoop-yarn_hadoop-yarn-api generated 11 new + 123 unchanged - 0 fixed = 134 total (was 123)
Other Tests			
+1	unit	0m 42s	hadoop-yarn-api in the patch passed.
-1	unit	16m 55s	hadoop-yarn-server-nodemanager in the patch failed.
-1	unit	0m 37s	hadoop-mapreduce-client-shuffle in the patch failed.
+1	asflicense	0m 36s	The patch does not generate ASF License warnings.
		123m 5s	


Reason	Tests
Failed junit tests	hadoop.yarn.server.nodemanager.scheduler.TestDistributedScheduler
	hadoop.mapred.TestShuffleHandler

Subsystem	Report/Notes
Docker	Image:yetus/hadoop:0de40f0
JIRA Issue	YARN-7244
JIRA Patch URL	https://issues.apache.org/jira/secure/attachment/12892309/YARN-7244.006.patch
Optional Tests	asflicense compile javac javadoc mvninstall mvnsite unit shadedclient findbugs checkstyle
uname	Linux 0b8bc629617a 3.13.0-116-generic #163-Ubuntu SMP Fri Mar 31 14:13:22 UTC 2017 x86_64 x86_64 x86_64 GNU/Linux
Build tool	maven
Personality	/testptch/hadoop/patchprocess/precommit/personality/provided.sh
git revision	trunk / 87ea1df
Default Java	1.8.0_144
findbugs	v3.1.0-RC1
checkstyle	https://builds.apache.org/job/PreCommit-YARN-Build/17945/artifact/patchprocess/diff-checkstyle-root.txt
javadoc	https://builds.apache.org/job/PreCommit-YARN-Build/17945/artifact/patchprocess/diff-javadoc-javadoc-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-api.txt
unit	https://builds.apache.org/job/PreCommit-YARN-Build/17945/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server_hadoop-yarn-server-nodemanager.txt
unit	https://builds.apache.org/job/PreCommit-YARN-Build/17945/artifact/patchprocess/patch-unit-hadoop-mapreduce-project_hadoop-mapreduce-client_hadoop-mapreduce-client-shuffle.txt
Test Results	https://builds.apache.org/job/PreCommit-YARN-Build/17945/testReport/
modules	C: hadoop-yarn-project/hadoop-yarn/hadoop-yarn-api hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-nodemanager hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-shuffle U: .
Console output	https://builds.apache.org/job/PreCommit-YARN-Build/17945/console
Powered by	Apache Yetus 0.6.0-SNAPSHOT http://yetus.apache.org

This message was automatically generated.

▼  Kuhu Shukla added a comment - 16/Oct/17 02:18

Updated patch that fixes checkstyles (almost all.. there is one to add getter in a test that seems excessive to me) and test failure for testMapFileAccess. My setup did not allow for that test to run and required overhauling. Verified that it passes now.

▼  Hadoop QA added a comment - 16/Oct/17 04:37


 -1 overall

Vote	Subsystem	Runtime	Comment
0	reexec	15m 10s	Docker mode activated.
Prechecks			
+1	@author	0m 0s	The patch does not contain any @author tags.
+1	test4tests	0m 0s	The patch appears to include 3 new or modified test files.
trunk Compile Tests			
0	mvndep	1m 26s	Maven dependency ordering for branch
+1	mvninstall	13m 51s	trunk passed
+1	compile	15m 44s	trunk passed
+1	checkstyle	2m 8s	trunk passed
+1	mvnsite	1m 40s	trunk passed
+1	shadedclient	15m 19s	branch has no errors when building and testing our client artifacts.
+1	findbugs	2m 47s	trunk passed
+1	javadoc	1m 17s	trunk passed
Patch Compile Tests			
0	mvndep	0m 17s	Maven dependency ordering for patch
+1	mvninstall	1m 11s	the patch passed
+1	compile	10m 54s	the patch passed
+1	javac	10m 54s	the patch passed
+1	checkstyle	2m 12s	root: The patch generated 0 new + 317 unchanged - 2 fixed = 317 total (was 319)
+1	mvnsite	1m 47s	the patch passed
+1	whitespace	0m 0s	The patch has no whitespace issues.
+1	shadedclient	11m 34s	patch has no errors when building and testing our client artifacts.
+1	findbugs	3m 10s	the patch passed
-1	javadoc	0m 31s	hadoop-yarn-project_hadoop-yarn_hadoop-yarn-api generated 11 new + 123 unchanged - 0 fixed = 134 total (was 123)
Other Tests			
+1	unit	0m 42s	hadoop-yarn-api in the patch passed.
-1	unit	16m 7s	hadoop-yarn-server-nodemanager in the patch failed.
+1	unit	0m 35s	hadoop-mapreduce-client-shuffle in the patch passed.
+1	asflicense	0m 36s	The patch does not generate ASF License warnings.
		136m 49s	


Reason	Tests
Failed junit tests	hadoop.yarn.server.nodemanager.scheduler.TestDistributedScheduler

Subsystem	Report/Notes
Docker	Image:yetus/hadoop:0de40f0
JIRA Issue	YARN-7244
JIRA Patch URL	https://issues.apache.org/jira/secure/attachment/12892321/YARN-7244.007.patch
Optional Tests	asflicense compile javac javadoc mvninstall mvnsite unit shadedclient findbugs checkstyle
uname	Linux 8a924ec4e6b5 3.13.0-119-generic #166-Ubuntu SMP Wed May 3 12:18:55 UTC 2017 x86_64 x86_64 x86_64 GNU/Linux
Build tool	maven
Personality	/testptch/hadoop/patchprocess/precommit/personality/provided.sh
git revision	trunk / 20575ec
Default Java	1.8.0_144
findbugs	v3.1.0-RC1
javadoc	https://builds.apache.org/job/PreCommit-YARN-Build/17946/artifact/patchprocess/diff-javadoc-javadoc-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-api.txt
unit	https://builds.apache.org/job/PreCommit-YARN-Build/17946/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server_hadoop-yarn-server-nodemanager.txt
Test Results	https://builds.apache.org/job/PreCommit-YARN-Build/17946/testReport/
modules	C: hadoop-yarn-project/hadoop-yarn/hadoop-yarn-api hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-nodemanager hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-shuffle U: .
Console output	https://builds.apache.org/job/PreCommit-YARN-Build/17946/console
Powered by	Apache Yetus 0.6.0-SNAPSHOT http://yetus.apache.org

This message was automatically generated.

-  **Kuhu Shukla** added a comment - 16/Oct/17 11:53

Fixing minor javadoc comments. Think the patch is ready for some review!

-  **Hadoop QA** added a comment - 16/Oct/17 14:22

 **-1 overall**


Vote	Subsystem	Runtime	Comment
0	reexec	0m 14s	Docker mode activated.
			Prechecks
+1	@author	0m 0s	The patch does not contain any @author tags.
+1	test4tests	0m 0s	The patch appears to include 3 new or modified test files.
			trunk Compile Tests
0	mvndep	0m 23s	Maven dependency ordering for branch
+1	mvninstall	17m 38s	trunk passed
+1	compile	19m 18s	trunk passed
+1	checkstyle	2m 29s	trunk passed
+1	mvnsite	2m 7s	trunk passed
+1	shadedclient	17m 17s	branch has no errors when building and testing our client artifacts.
+1	findbugs	3m 25s	trunk passed

+1	javadoc	1m 42s	trunk passed
			Patch Compile Tests
0	mvndep	0m 22s	Maven dependency ordering for patch
+1	mvninstall	1m 41s	the patch passed
+1	compile	14m 40s	the patch passed
+1	javac	14m 40s	the patch passed
+1	checkstyle	2m 35s	root: The patch generated 0 new + 318 unchanged - 2 fixed = 318 total (was 320)
+1	mvnsite	2m 19s	the patch passed
+1	whitespace	0m 0s	The patch has no whitespace issues.
+1	shadedclient	12m 36s	patch has no errors when building and testing our client artifacts.
+1	findbugs	4m 12s	the patch passed
+1	javadoc	1m 49s	the patch passed
			Other Tests
+1	unit	0m 55s	hadoop-yarn-api in the patch passed.
-1	unit	17m 29s	hadoop-yarn-server-nodemanager in the patch failed.
+1	unit	0m 43s	hadoop-mapreduce-client-shuffle in the patch passed.
+1	asflicense	0m 44s	The patch does not generate ASF License warnings.
		141m 6s	


Reason	Tests
Failed junit tests	hadoop.yarn.server.nodemanager.scheduler.TestDistributedScheduler

Subsystem	Report/Notes
Docker	Image:yetus/hadoop:0de40f0
JIRA Issue	YARN-7244
JIRA Patch URL	https://issues.apache.org/jira/secure/attachment/12892371/YARN-7244.009.patch
Optional Tests	asflicense compile javac javadoc mvninstall mvnsite unit shadedclient findbugs checkstyle
uname	Linux cc2e02d5e1c7 3.13.0-119-generic #166-Ubuntu SMP Wed May 3 12:18:55 UTC 2017 x86_64 x86_64 x86_64 GNU/Linux
Build tool	maven
Personality	/testptch/hadoop/patchprocess/precommit/personality/provided.sh
git revision	trunk / 20575ec
Default Java	1.8.0_144
findbugs	v3.1.0-RC1
unit	https://builds.apache.org/job/PreCommit-YARN-Build/17953/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server_hadoop-yarn-server-nodemanager.txt
Test Results	https://builds.apache.org/job/PreCommit-YARN-Build/17953/testReport/
modules	C: hadoop-yarn-project/hadoop-yarn/hadoop-yarn-api hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-nodemanager hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-shuffle U: .
Console output	https://builds.apache.org/job/PreCommit-YARN-Build/17953/console
Powered by	Apache Yetus 0.6.0-SNAPSHOT http://yetus.apache.org

This message was automatically generated.

👤  Kuhu Shukla added a comment - 16/Oct/17 14:24

jlowe, request for review/comments on the latest patch. Thanks again.

👤  Jason Darrell Lowe added a comment - 16/Oct/17 20:33

Thanks for updating the patch!

Nit: The whitespace separating the private static fields from the MockShuffleHandler definition was lost.

MockShuffleHandler2 should remain a static class.

Nit: it would be nice if the javadocs added in AuxiliaryService were formatted to fit the style of the other method docs in this class.

This change seems unnecessary:

```
--- a/hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-nodemanager/src/main/java/org/apache/hadoop/yarn/server/ContainerManagerImpl.java
+++ b/hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-nodemanager/src/main/java/org/apache/hadoop/yarn/server/ContainerManagerImpl.java
@@ -224,8 +227,9 @@
     private NMTimelinePublisher nmMetricsPublisher;

     public ContainerManagerImpl(Context context, ContainerExecutor exec,
-        DeletionService deletionContext, NodeStatusUpdater nodeStatusUpdater,
-        NodeManagerMetrics metrics, LocalDirsHandlerService dirsHandler) {
+        DeletionService deletionContext, NodeStatusUpdater nodeStatusUpdater,
+        NodeManagerMetrics metrics, LocalDirsHandlerService dirsHandler) {
         DeletionService deletionContext,
         NodeStatusUpdater nodeStatusUpdater, NodeManagerMetrics metrics,
         LocalDirsHandlerService dirsHandler) {
         super(ContainerManagerImpl.class.getName());
         this.context = context;
         this.dirsHandler = dirsHandler;
```

TestAuxServices.auxiliaryLocalPathHandler should be final and named like a constant, e.g.: MOCK_AUX_PATH_HANDLER


AuxiliaryLocalPathHandlerImpl should be package-private.

In TestContainerManager#testAuxPathHandler:


```
+ try {
+     Path p = auxiliaryLocalPathHandler.getLocalPathForRead("test");
+     assertTrue(p != null);
+ } catch (Exception e) {
+     fail("Unexpected Exception!" + e);
+ }
```

This should just let the exception be thrown to fail the test directly rather than catch it explicitly only to fail in a generic way.

It would be less code to use a spy object on a real LocalDirsHandlerService that is initialized with the conf rather than a mock that calls real methods on an object without properly initializing it.

👤  Kuhu Shukla added a comment - 16/Oct/17 21:43

Attaching revised patch that address review comments. Thanks a lot !

👤  Hadoop QA added a comment - 17/Oct/17 00:56

 -1 overall

Vote	Subsystem	Runtime	Comment
0	reexec	0m 18s	Docker mode activated.
Prechecks			
+1	@author	0m 0s	The patch does not contain any @author tags.
+1	test4tests	0m 0s	The patch appears to include 3 new or modified test files.
trunk Compile Tests			
0	mvndep	0m 36s	Maven dependency ordering for branch
+1	mvninstall	23m 52s	trunk passed

+1	compile	24m 52s	trunk passed
+1	checkstyle	2m 39s	trunk passed
+1	mvnsite	2m 22s	trunk passed
+1	shadedclient	19m 9s	branch has no errors when building and testing our client artifacts.
+1	findbugs	4m 24s	trunk passed
+1	javadoc	1m 58s	trunk passed
Patch Compile Tests			
0	mvndep	0m 26s	Maven dependency ordering for patch
+1	mvninstall	1m 57s	the patch passed
+1	compile	17m 30s	the patch passed
+1	javac	17m 30s	the patch passed
-0	checkstyle	3m 55s	root: The patch generated 1 new + 318 unchanged - 2 fixed = 319 total (was 320)
-1	mvnsite	0m 52s	hadoop-yarn-server-nodemanager in the patch failed.
-1	mvnsite	0m 47s	hadoop-mapreduce-client-shuffle in the patch failed.
-1	whitespace	0m 0s	The patch has 2 line(s) that end in whitespace. Use git apply --whitespace=fix <<patch_file>>. Refer https://git-scm.com/docs/git-apply
+1	shadedclient	14m 47s	patch has no errors when building and testing our client artifacts.
+1	findbugs	4m 46s	the patch passed
-1	javadoc	0m 47s	hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server_hadoop-yarn-server-nodemanager generated 7 new + 103 unchanged - 0 fixed = 110 total (was 103)
Other Tests			
+1	unit	1m 4s	hadoop-yarn-api in the patch passed.
-1	unit	0m 46s	hadoop-yarn-server-nodemanager in the patch failed.
-1	unit	0m 43s	hadoop-mapreduce-client-shuffle in the patch failed.
+1	asflicense	0m 45s	The patch does not generate ASF License warnings.
		149m 33s	

Subsystem	Report/Notes
Docker	Image:yetus/hadoop:0de40f0
JIRA Issue	YARN-7244
JIRA Patch URL	https://issues.apache.org/jira/secure/attachment/12892477/YARN-7244.010.patch
Optional Tests	asflicense compile javac javadoc mvninstall mvnsite unit shadedclient findbugs checkstyle
uname	Linux 090c49811e83 3.13.0-119-generic #166-Ubuntu SMP Wed May 3 12:18:55 UTC 2017 x86_64 x86_64 x86_64 GNU/Linux
Build tool	maven
Personality	/testptch/hadoop/patchprocess/precommit/personality/provided.sh
git revision	trunk / 4540ffd
Default Java	1.8.0_144
findbugs	v3.1.0-RC1
checkstyle	https://builds.apache.org/job/PreCommit-YARN-Build/17972/artifact/patchprocess/diff-checkstyle-root.txt
mvnsite	https://builds.apache.org/job/PreCommit-YARN-Build/17972/artifact/patchprocess/patch-mvnsite-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server_hadoop-yarn-server-nodemanager.txt
mvnsite	https://builds.apache.org/job/PreCommit-YARN-Build/17972/artifact/patchprocess/patch-mvnsite-hadoop-mapreduce-project_hadoop-mapreduce-client_hadoop-mapreduce-client-shuffle.txt

whitespace	https://builds.apache.org/job/PreCommit-YARN-Build/17972/artifact/patchprocess/whitespace-eol.txt
javadoc	https://builds.apache.org/job/PreCommit-YARN-Build/17972/artifact/patchprocess/diff-javadoc-javadoc-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server_hadoop-yarn-server-nodemanager.txt
unit	https://builds.apache.org/job/PreCommit-YARN-Build/17972/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server_hadoop-yarn-server-nodemanager.txt
unit	https://builds.apache.org/job/PreCommit-YARN-Build/17972/artifact/patchprocess/patch-unit-hadoop-mapreduce-project_hadoop-mapreduce-client_hadoop-mapreduce-client-shuffle.txt
Test Results	https://builds.apache.org/job/PreCommit-YARN-Build/17972/testReport/
modules	C: hadoop-yarn-project/hadoop-yarn/hadoop-yarn-api hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-nodemanager hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-shuffle U: .
Console output	https://builds.apache.org/job/PreCommit-YARN-Build/17972/console
Powered by	Apache Yetus 0.6.0-SNAPSHOT http://yetus.apache.org

This message was automatically generated.


▼  Sunil G added a comment - 17/Oct/17 17:24

Thanks [kshukla](#) and [jlowe](#)

Usage of `AuxiliaryLocalPathHandler` looks very clean here.

Minor comments,

1. Do you think is it better to have a setter and update `AuxiliaryLocalPathHandler` to `AuxServices` rather than changing `AuxServices` ctor.
2. `AuxiliaryLocalPathHandler` could be in `org.apache.hadoop.yarn.server.api`? any reasons to move to `api`?
3. All apis in `AuxiliaryLocalPathHandlerImpl` could have `Override` annotation.
4. Does `ContainerManagerImpl` need to have a `getAuxiliaryLocalPathHandler` ?

▼  Kuhu Shukla added a comment - 17/Oct/17 18:24

Thank you [sunilg](#) for the comments!

Do you think is it better to have a setter and update `AuxiliaryLocalPathHandler` to `AuxServices` rather than changing `AuxServices` ctor.

The constructor change makes sure that we always initialize the `pathHandler` which seems safer to me.

`AuxiliaryLocalPathHandler` could be in `org.apache.hadoop.yarn.server.api`? any reasons to move to `api`?

You are right. This needs to be in server apis.

All apis in `AuxiliaryLocalPathHandlerImpl` could have `Override` annotation.


Will do.

Does `ContainerManagerImpl` need to have a `getAuxiliaryLocalPathHandler` ?

I added the getter to assist any testing in the futur and made it package private. I can mark it as `VisibleForTesting` or take it out , either way would be fine.

▼  Kuhu Shukla added a comment - 17/Oct/17 18:40

Updated patch addressing comments from [sunilg](#).

▼  Hadoop QA added a comment - 17/Oct/17 21:09

 -1 overall


Vote	Subsystem	Runtime	Comment
0	reexec	16m 31s	Docker mode activated.
			Prechecks
+1	@author	0m 0s	The patch does not contain any @author tags.
+1	test4tests	0m 0s	The patch appears to include 3 new or modified test files.
			trunk Compile Tests
0	mvndep	1m 47s	Maven dependency ordering for branch
+1	mvninstall	17m 24s	trunk passed
+1	compile	18m 40s	trunk passed
+1	checkstyle	2m 24s	trunk passed
+1	mvnsite	2m 0s	trunk passed
+1	shadedclient	16m 57s	branch has no errors when building and testing our client artifacts.
+1	findbugs	3m 21s	trunk passed
+1	javadoc	1m 30s	trunk passed
			Patch Compile Tests
0	mvndep	0m 20s	Maven dependency ordering for patch
+1	mvninstall	1m 23s	the patch passed
+1	compile	11m 30s	the patch passed
+1	javac	11m 30s	the patch passed
-0	checkstyle	2m 12s	root: The patch generated 1 new + 317 unchanged - 2 fixed = 318 total (was 319)
+1	mvnsite	1m 45s	the patch passed
-1	whitespace	0m 0s	The patch has 2 line(s) that end in whitespace. Use git apply --whitespace=fix <<patch_file>>. Refer https://git-scm.com/docs/git-apply
+1	shadedclient	11m 36s	patch has no errors when building and testing our client artifacts.
+1	findbugs	3m 11s	the patch passed
+1	javadoc	1m 27s	the patch passed
			Other Tests
+1	unit	0m 41s	hadoop-yarn-api in the patch passed.
-1	unit	16m 7s	hadoop-yarn-server-nodemananger in the patch failed.
+1	unit	0m 34s	hadoop-mapreduce-client-shuffle in the patch passed.
+1	asflicense	0m 38s	The patch does not generate ASF License warnings.
		148m 24s	

Reason	Tests
Failed junit tests	hadoop.yarn.server.nodemananger.scheduler.TestDistributedScheduler


Subsystem	Report/Notes
Docker	Image:yetus/hadoop:0de40f0
JIRA Issue	YARN-7244
JIRA Patch URL	https://issues.apache.org/jira/secure/attachment/12892665/YARN-7244.011.patch
Optional Tests	asflicense compile javac javadoc mvninstall mvnsite unit shadedclient findbugs checkstyle
uname	Linux 8cb6586f23bf 3.13.0-119-generic #166-Ubuntu SMP Wed May 3 12:18:55 UTC 2017 x86_64 x86_64 x86_64

	GNU/Linux
Build tool	maven
Personality	/testptch/hadoop/patchprocess/precommit/personality/provided.sh
git revision	trunk / acabc65
Default Java	1.8.0_144
findbugs	v3.1.0-RC1
checkstyle	https://builds.apache.org/job/PreCommit-YARN-Build/17986/artifact/patchprocess/diff-checkstyle-root.txt
whitespace	https://builds.apache.org/job/PreCommit-YARN-Build/17986/artifact/patchprocess/whitespace-eol.txt
unit	https://builds.apache.org/job/PreCommit-YARN-Build/17986/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server_hadoop-yarn-server-nodemanager.txt
Test Results	https://builds.apache.org/job/PreCommit-YARN-Build/17986/testReport/
modules	C: hadoop-yarn-project/hadoop-yarn/hadoop-yarn-api hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-nodemanager hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-shuffle U: .
Console output	https://builds.apache.org/job/PreCommit-YARN-Build/17986/console
Powered by	Apache Yetus 0.6.0-SNAPSHOT http://yetus.apache.org

This message was automatically generated.

▼  **Kuhu Shukla** added a comment - 17/Oct/17 21:17

Fixing minor checkstyle issues. 😞


▼  **Hadoop QA** added a comment - 17/Oct/17 21:43

 **-1 overall**

Vote	Subsystem	Runtime	Comment
0	reexec	0m 0s	Docker mode activated.
-1	docker	5m 33s	Docker failed to build yetus/hadoop:0de40f0.

Subsystem	Report/Notes
JIRA Issue	YARN-7244
JIRA Patch URL	https://issues.apache.org/jira/secure/attachment/12892684/YARN-7244.012.patch
Console output	https://builds.apache.org/job/PreCommit-YARN-Build/17988/console
Powered by	Apache Yetus 0.6.0-SNAPSHOT http://yetus.apache.org

This message was automatically generated.

▼  **Hadoop QA** added a comment - 18/Oct/17 00:15

 **-1 overall**

Vote	Subsystem	Runtime	Comment
0	reexec	0m 10s	Docker mode activated.
			Prechecks
+1	@author	0m 0s	The patch does not contain any @author tags.
+1	test4tests	0m 0s	The patch appears to include 3 new or modified test files.


			trunk Compile Tests
0	mvndep	1m 32s	Maven dependency ordering for branch
+1	mvninstall	15m 17s	trunk passed
+1	compile	16m 18s	trunk passed
+1	checkstyle	2m 13s	trunk passed
+1	mvnsite	1m 41s	trunk passed
+1	shadedclient	15m 37s	branch has no errors when building and testing our client artifacts.
+1	findbugs	2m 52s	trunk passed
+1	javadoc	1m 19s	trunk passed
			Patch Compile Tests
0	mvndep	0m 19s	Maven dependency ordering for patch
+1	mvninstall	1m 27s	the patch passed
+1	compile	11m 52s	the patch passed
+1	javac	11m 52s	the patch passed
+1	checkstyle	2m 18s	root: The patch generated 0 new + 317 unchanged - 2 fixed = 317 total (was 319)
+1	mvnsite	1m 57s	the patch passed
+1	whitespace	0m 0s	The patch has no whitespace issues.
+1	shadedclient	11m 35s	patch has no errors when building and testing our client artifacts.
+1	findbugs	3m 11s	the patch passed
+1	javadoc	1m 28s	the patch passed
			Other Tests
+1	unit	0m 42s	hadoop-yarn-api in the patch passed.
-1	unit	16m 8s	hadoop-yarn-server-nodemanager in the patch failed.
-1	unit	0m 57s	hadoop-mapreduce-client-shuffle in the patch failed.
+1	asflicense	0m 35s	The patch does not generate ASF License warnings.
		126m 8s	

Reason	Tests
Failed junit tests	hadoop.yarn.server.nodemanager.scheduler.TestDistributedScheduler


Subsystem	Report/Notes
Docker	Image:yetus/hadoop:0de40f0
JIRA Issue	YARN-7244
JIRA Patch URL	https://issues.apache.org/jira/secure/attachment/12892684/YARN-7244.012.patch
Optional Tests	asflicense compile javac javadoc mvninstall mvnsite unit shadedclient findbugs checkstyle
uname	Linux 879d912a414c 3.13.0-119-generic #166-Ubuntu SMP Wed May 3 12:18:55 UTC 2017 x86_64 x86_64 x86_64 GNU/Linux
Build tool	maven
Personality	/testptch/hadoop/patchprocess/precommit/personality/provided.sh
git revision	trunk / 7532339
Default Java	1.8.0_144
findbugs	v3.1.0-RC1

unit	https://builds.apache.org/job/PreCommit-YARN-Build/17992/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server_hadoop-yarn-server-nodemanager.txt
unit	https://builds.apache.org/job/PreCommit-YARN-Build/17992/artifact/patchprocess/patch-unit-hadoop-mapreduce-project_hadoop-mapreduce-client_hadoop-mapreduce-client-shuffle.txt
Test Results	https://builds.apache.org/job/PreCommit-YARN-Build/17992/testReport/
modules	C: hadoop-yarn-project/hadoop-yarn/hadoop-yarn-api hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-nodemanager hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-shuffle U: .
Console output	https://builds.apache.org/job/PreCommit-YARN-Build/17992/console
Powered by	Apache Yetus 0.6.0-SNAPSHOT http://yetus.apache.org

This message was automatically generated.

🗲️  Kuhu Shukla added a comment - 18/Oct/17 00:56

The build failed with 2 separate and unrelated issues I believe. The second time the cache seems to be picking up the old package for AuxiliaryLocalPathHandler. Re-triggering by uploading the same patch again. Please let me know if I missed something.

🗲️  Hadoop QA added a comment - 18/Oct/17 03:00

 -1 overall

Vote	Subsystem	Runtime	Comment
0	reexec	0m 9s	Docker mode activated.
			Prechecks
+1	@author	0m 0s	The patch does not contain any @author tags.
+1	test4tests	0m 0s	The patch appears to include 3 new or modified test files.
			trunk Compile Tests
0	mvndep	1m 9s	Maven dependency ordering for branch
+1	mvninstall	13m 19s	trunk passed
+1	compile	14m 19s	trunk passed
+1	checkstyle	2m 8s	trunk passed
+1	mvnsite	1m 35s	trunk passed
+1	shadedclient	15m 22s	branch has no errors when building and testing our client artifacts.
+1	findbugs	2m 36s	trunk passed
+1	javadoc	1m 16s	trunk passed
			Patch Compile Tests
0	mvndep	0m 18s	Maven dependency ordering for patch
+1	mvninstall	1m 13s	the patch passed
+1	compile	11m 1s	the patch passed
+1	javac	11m 1s	the patch passed
+1	checkstyle	2m 8s	root: The patch generated 0 new + 317 unchanged - 2 fixed = 317 total (was 319)
+1	mvnsite	1m 47s	the patch passed
+1	whitespace	0m 0s	The patch has no whitespace issues.
+1	shadedclient	11m 31s	patch has no errors when building and testing our client artifacts.
+1	findbugs	3m 10s	the patch passed
+1	javadoc	1m 27s	the patch passed
			Other Tests
+1	unit	0m 42s	hadoop-yarn-api in the patch passed.

-1	unit	16m 10s	hadoop-yarn-server-nodemanager in the patch failed.
+1	unit	0m 35s	hadoop-mapreduce-client-shuffle in the patch passed.
+1	asflicense	0m 35s	The patch does not generate ASF License warnings.
		119m 35s	


Reason	Tests
Failed junit tests	hadoop.yarn.server.nodemanager.scheduler.TestDistributedScheduler

Subsystem	Report/Notes
Docker	Image:yetus/hadoop:0de40f0
JIRA Issue	YARN-7244
JIRA Patch URL	https://issues.apache.org/jira/secure/attachment/12892714/YARN-7244.013.patch
Optional Tests	asflicense compile javac javadoc mvninstall mvnsite unit shadedclient findbugs checkstyle
uname	Linux b20209a04aaa 3.13.0-119-generic #166-Ubuntu SMP Wed May 3 12:18:55 UTC 2017 x86_64 x86_64 x86_64 GNU/Linux
Build tool	maven
Personality	/testptch/hadoop/patchprocess/precommit/personality/provided.sh
git revision	trunk / 46eb103
Default Java	1.8.0_144
findbugs	v3.1.0-RC1
unit	https://builds.apache.org/job/PreCommit-YARN-Build/18002/artifact/patchprocess/patch-unit-hadoop-yarn-project_hadoop-yarn_hadoop-yarn-server_hadoop-yarn-server-nodemanager.txt
Test Results	https://builds.apache.org/job/PreCommit-YARN-Build/18002/testReport/
modules	C: hadoop-yarn-project/hadoop-yarn/hadoop-yarn-api hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-nodemanager hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-shuffle U: .
Console output	https://builds.apache.org/job/PreCommit-YARN-Build/18002/console
Powered by	Apache Yetus 0.6.0-SNAPSHOT http://yetus.apache.org

This message was automatically generated.

▼  Kuhu Shukla added a comment - 18/Oct/17 11:05

jlowe, sunilg request for comments/review. Thanks a lot!

▼  Kuhu Shukla added a comment - 25/Oct/17 16:05

jlowe/sunilg appreciate any comments on the latest patch! Thank you.

▼  Sunil G added a comment - 25/Oct/17 16:17

I think latest patch seems fine to me.

▼  Jason Darrell Lowe added a comment - 27/Oct/17 21:34

+1 for the latest patch. The unit test failure is unrelated and tracked by [YARN-7299](#).

Committing this.

▼  Jason Darrell Lowe added a comment - 27/Oct/17 22:12


Thanks to Kuhu for the contribution and to Bibin and Sunil for additional review! I committed this to trunk, branch-3.0, and branch-2.

▼  Hudson added a comment - 27/Oct/17 23:02


SUCCESS: Integrated in Jenkins build Hadoop-trunk-Commit #13150 (See <https://builds.apache.org/job/Hadoop-trunk-Commit/13150/>)

[YARN-7244](#). ShuffleHandler is not aware of disks that are added. (jlw: rev 665bb147aa3fc198e53335931562dd69f9f0e345)


- (edit) hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-shuffle/src/main/java/org/apache/hadoop/mapred/ShuffleHandler.java
- (edit) hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-nodemanager/src/main/java/org/apache/hadoop/yarn/server/nodemanager/containermanager/AuxServices.java
- (edit) hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-shuffle/src/test/java/org/apache/hadoop/mapred/TestShuffleHandler.java
- (edit) hadoop-yarn-project/hadoop-yarn/hadoop-yarn-api/src/main/java/org/apache/hadoop/yarn/server/api/AuxiliaryService.java
- (edit) hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-nodemanager/src/test/java/org/apache/hadoop/yarn/server/nodemanager/containermanager/TestAuxServices.java
- (edit) hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-nodemanager/src/test/java/org/apache/hadoop/yarn/server/nodemanager/containermanager/TestContainerManager.java
- (add) hadoop-yarn-project/hadoop-yarn/hadoop-yarn-api/src/main/java/org/apache/hadoop/yarn/server/api/AuxiliaryLocalPathHandler.java
- (edit) hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-nodemanager/src/main/java/org/apache/hadoop/yarn/server/nodemanager/containermanager/ContainerManagerImpl.java

▼  Kuhu Shukla added a comment - 30/Oct/17 15:31


Re-opening to attach 2.8 version of the patch.

▼  Kuhu Shukla added a comment - 30/Oct/17 15:33

Attaching 2.8 version of the patch which needed some extra changes. The important one is in LocalDirsHandlerService which was missing getLocalPathForRead() method from trunk which went in as part of [YARN-3998](#). I have added just that method rather than change the visibility of getPathToRead().

▼  Kuhu Shukla added a comment - 30/Oct/17 15:58

Fixing minor new line checkstyle.

▼  Hadoop QA added a comment - 30/Oct/17 17:49

 -1 overall

Vote	Subsystem	Runtime	Comment
0	reexec	12m 6s	Docker mode activated.
Prechecks			
+1	@author	0m 0s	The patch does not contain any @author tags.
+1	test4tests	0m 0s	The patch appears to include 3 new or modified test files.
branch-2.8 Compile Tests			
0	mvndep	0m 18s	Maven dependency ordering for branch
+1	mvninstall	9m 7s	branch-2.8 passed
+1	compile	7m 2s	branch-2.8 passed
+1	checkstyle	1m 7s	branch-2.8 passed
+1	mvnsite	1m 19s	branch-2.8 passed
+1	findbugs	2m 34s	branch-2.8 passed
+1	javadoc	0m 52s	branch-2.8 passed
Patch Compile Tests			

0	mvndep	0m 16s	Maven dependency ordering for patch
+1	mvninstall	1m 2s	the patch passed
+1	compile	7m 15s	the patch passed
+1	javac	7m 15s	the patch passed
+1	checkstyle	1m 10s	root: The patch generated 0 new + 176 unchanged - 1 fixed = 176 total (was 177)
+1	mvnsite	1m 20s	the patch passed
+1	whitespace	0m 0s	The patch has no whitespace issues.
+1	findbugs	3m 17s	the patch passed
+1	javadoc	0m 52s	the patch passed
Other Tests			
+1	unit	0m 26s	hadoop-yarn-api in the patch passed.
+1	unit	9m 50s	hadoop-yarn-server-nodemanager in the patch passed.
+1	unit	0m 22s	hadoop-mapreduce-client-shuffle in the patch passed.
-1	asflicense	0m 25s	The patch generated 3 ASF License warnings.
		66m 27s	

Subsystem	Report/Notes
Docker	Client=17.05.0-ce Server=17.05.0-ce Image:yetus/hadoop:c2d96dd
JIRA Issue	YARN-7244
JIRA Patch URL	https://issues.apache.org/jira/secure/attachment/12894773/YARN-7244-branch-2.8.002.patch
Optional Tests	asflicense compile javac javadoc mvninstall mvnsite unit shadedclient findbugs checkstyle
uname	Linux 57c1ddadc87d 3.13.0-123-generic #172-Ubuntu SMP Mon Jun 26 18:04:35 UTC 2017 x86_64 x86_64 x86_64 GNU/Linux
Build tool	maven
Personality	/testptch/patchprocess/precommit/personality/provided.sh
git revision	branch-2.8 / 2d4629f
maven	version: Apache Maven 3.0.5
Default Java	1.7.0_151
findbugs	v3.0.0
Test Results	https://builds.apache.org/job/PreCommit-YARN-Build/18241/testReport/
asflicense	https://builds.apache.org/job/PreCommit-YARN-Build/18241/artifact/out/patch-asflicense-problems.txt
modules	C: hadoop-yarn-project/hadoop-yarn/hadoop-yarn-api hadoop-yarn-project/hadoop-yarn/hadoop-yarn-server/hadoop-yarn-server-nodemanager hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-shuffle U: .
Console output	https://builds.apache.org/job/PreCommit-YARN-Build/18241/console
Powered by	Apache Yetus 0.7.0-SNAPSHOT http://yetus.apache.org

This message was automatically generated.


👤 [Kuhu Shukla](#) added a comment - 30/Oct/17 18:19

[jlowe](#), request for comments on the 2.8 version of the patch. Appreciate it!

👤 [Jason Darrell Lowe](#) added a comment - 30/Oct/17 21:06

The ASF warnings are unrelated.


+1 for the branch-2.8 patch. Committing this.

▼  Jason Darrell Lowe added a comment - 30/Oct/17 21:11


Thanks, Kuhu! I committed this to branch-2.8 as well.

▼ People

Assignee:

 Kuhu Shukla

Reporter:

 Kuhu Shukla

Votes:

 0 Vote for this issue

Watchers:

 8 Start watching this issue

▼ Dates

Created:

22/Sep/17 14:13

Updated:

30/Oct/17 21:11

Resolved:

30/Oct/17 21:11